

distortion_classifier

November 23, 2025

```
[1]: import os
from tqdm import tqdm
import torch
import torchvision
import torch.nn as nn
import numpy as np
import torch.nn.functional as F
import torchvision.transforms as T
from torchvision.datasets import ImageFolder
from torch.utils.data import DataLoader
from torch.utils.data import random_split
from torchvision.utils import make_grid
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

matplotlib.rcParams['figure.facecolor'] = '#ffffff'
```

```
[2]: ROOT_DIR = os.path.abspath(os.path.join(os.getcwd(), '..', '..'))
DATA_DIR = os.path.join(ROOT_DIR, 'data', 'cnn')
print(os.listdir(DATA_DIR))
classes = os.listdir(os.path.join(DATA_DIR, 'train')) # Must remove enhanced_
    ↴folder
print(classes)

['test', 'test2', 'test3', 'train', 'valid']
['h_blur', 'low_light', 'low_qual', 'normal', 'v_blur']
```

0.0.1 Data Transformation

```
[4]: # Data transforms (normalization & data augmentation)
stats = ((0.42720765, 0.43359186, 0.44090385), (0.24280364, 0.24590165, 0.
    ↴23916515))
train_tfms = T.Compose([T.Resize([128, 256]),
                        T.RandomPerspective(distortion_scale=0.2, p=0.2),
                        T.ColorJitter(hue=.2),
                        T.ToTensor(),
                        T.Normalize(*stats, inplace=True)])
```

```
valid_tfms = T.Compose([T.Resize([128, 256]), T.ToTensor(), T.  
    Normalize(*stats)])
```

0.0.2 PyTorch Datasets / Data Loaders

```
[6]: train_ds = ImageFolder(os.path.join(DATA_DIR, 'train'), train_tfms)  
valid_ds = ImageFolder(os.path.join(DATA_DIR, 'valid'), valid_tfms)
```

```
[7]: batch_size = 64
```

```
[8]: # PyTorch Data Loaders  
train_dl = DataLoader(train_ds, batch_size, shuffle=True, num_workers=3, □  
    pin_memory=True)  
valid_dl = DataLoader(valid_ds, batch_size*2, num_workers=3, pin_memory=True)
```

```
[9]: def denormalize(images, means, stds):  
    means = torch.tensor(means).reshape(1, 3, 1, 1)  
    stds = torch.tensor(stds).reshape(1, 3, 1, 1)  
    return images * stds + means  
  
def show_batch(dl):  
    for images, labels in dl:  
        fig, ax = plt.subplots(figsize=(12, 12))  
        ax.set_xticks([]); ax.set_yticks([])  
        denorm_images = denormalize(images, *stats)  
        ax.imshow(make_grid(denorm_images[:32], nrow=8).permute(1, 2, 0).  
            clamp(0, 1))  
        break
```

```
[35]: show_batch(train_dl)
```



0.0.3 Using GPU

```
[10]: def get_default_device():
        if torch.cuda.is_available():
            return torch.device('cuda')
        return torch.device('cpu')

def to_device(data, device):
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self.dl)
```

```
[11]: device = get_default_device()
device
```

```
[11]: device(type='cuda')
```

```
[12]: train_dl = DeviceDataLoader(train_dl, device)
valid_dl = DeviceDataLoader(valid_dl, device)
```

0.0.4 Model Development

```
[14]: def accuracy(outputs, labels):
        _, preds = torch.max(outputs, dim=1)
        return torch.tensor(torch.sum(preds == labels).item() / len(preds))

class ImageClassificationBase(nn.Module):
    def training_step(self, batch):
        images, labels = batch
        out = self(images)
        loss = F.cross_entropy(out, labels)
        return loss

    def validation_step(self, batch):
        images, labels = batch
        out = self(images)
```

```

        loss = F.cross_entropy(out, labels)
        acc = accuracy(out, labels)
        return {'val_loss': loss.detach(), 'val_acc': acc}

    def validation_epoch_end(self, outputs):
        batch_losses = [x['val_loss'] for x in outputs]
        epoch_loss = torch.stack(batch_losses).mean()           # Combine losses
        batch_accs = [x['val_acc'] for x in outputs]
        epoch_acc = torch.stack(batch_accs).mean()             # Combine accuracies
        return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}

    def epoch_end(self, epoch, result):
        print(f'Epoch [{epoch}], last_lr: {result['lrs'][-1]:.5f}, train_loss: {result['train_loss']:.4f}, '
              f'val_loss: {result['val_loss']:.4f}, val_acc: {result['val_acc']:.4f}!')

```

```
[15]: def conv_block(in_channels, out_channels, pool=False):
    layers = [nn.Conv2d(in_channels, out_channels, kernel_size=3, padding=1),
              nn.BatchNorm2d(out_channels),
              nn.ReLU(inplace=True)]
    if pool: layers.append(nn.MaxPool2d(2))
    return nn.Sequential(*layers)

class DistortionClassifier(ImageClassificationBase):
    def __init__(self, in_channels, num_classes):
        super().__init__()

        # 3 x 128 x 256
        self.conv1 = conv_block(in_channels, 64, pool=True) # 64 x 64 x 128
        self.conv2 = conv_block(64, 128, pool=True) # 128 x 32 x 64
        self.res1 = nn.Sequential(conv_block(128, 128), conv_block(128, 128)) # 128 x 32 x 64

        self.conv3 = conv_block(128, 256, pool=True) # 256 x 16 x 32
        self.conv4 = conv_block(256, 512, pool=True) # 512 x 8 x 16
        self.res2 = nn.Sequential(conv_block(512, 512), conv_block(512, 512)) # 512 x 8 x 16

        self.classifier = nn.Sequential(nn.MaxPool2d(2), # 512 x 4 x 8
                                         nn.Flatten(),
                                         nn.Dropout(0.2),
                                         nn.Linear(512 * 4 * 8, num_classes))

    def forward(self, xb):
        out = self.conv1(xb)
```

```
        out = self.conv2(out)
        out = self.res1(out) + out
        out = self.conv3(out)
        out = self.conv4(out)
        out = self.res2(out) + out
        out = self.classifier(out)
    return out
```

```
[16]: from torchvision import models
from torchvision.models import resnet18

class DistortionClassifier18(ImageClassificationBase):
    def __init__(self, num_classes):
        super().__init__()
        self.network = models.resnet18()
        self.network.fc = nn.Linear(self.network.fc.in_features, num_classes)

    def forward(self, xb):
        return self.network(xb)
```

```
[17]: model = to_device(DistortionClassifier18(5), device)
model
```

```
[17]: DistortionClassifier18(
    network: ResNet(
        (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
        (layer1): Sequential(
            (0): BasicBlock(
                (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
                (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (relu): ReLU(inplace=True)
                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
                (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            )
            (1): BasicBlock(
                (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
```

```

        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(layer2): Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
            (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (1): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(layer3): Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    )
)

```

```

        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
            (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (1): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(layer4): Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
            (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (1): BasicBlock(
        (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)

```

```

        )
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Linear(in_features=512, out_features=5, bias=True)
)
)

[18]: @torch.no_grad()
def evaluate(model, val_loader):
    model.eval()
    outputs = [model.validation_step(batch) for batch in val_loader]
    return model.validation_epoch_end(outputs)

def get_lr(optimizer):
    for param_group in optimizer.param_groups:
        return param_group['lr']

def fit_one_cycle(epochs, max_lr, model, train_loader, val_loader,
                  weight_decay=0, grad_clip=None, opt_func=torch.optim.SGD):
    torch.cuda.empty_cache()
    history = []

    # Set up custom optimizer with weight decay
    optimizer = opt_func(model.parameters(), max_lr, weight_decay=weight_decay)
    # Set up one-cycle learning rate scheduling
    sched = torch.optim.lr_scheduler.OneCycleLR(optimizer, max_lr,
                                                epochs=epochs,
                                                steps_per_epoch=len(train_loader))

    for epoch in range(epochs):
        # Training phase
        model.train()
        train_losses = []
        lrs = []
        for batch in tqdm(train_loader):
            loss = model.training_step(batch)
            train_losses.append(loss)
            loss.backward()

            # Gradient clipping
            if grad_clip:
                nn.utils.clip_grad_value_(model.parameters(), grad_clip)

            optimizer.step()
            optimizer.zero_grad()

```

```

# Record and update learning rate
lrs.append(get_lr(optimizer))
sched.step()

# Validation phase
result = evaluate(model, val_loader)
result['train_loss'] = torch.stack(train_losses).mean().item()
result['lrs'] = lrs
model.epoch_end(epoch, result)
history.append(result)

return history

```

[21]: history = [evaluate(model, valid_dl)]
history

[21]: {'val_loss': 1.9271801710128784, 'val_acc': 0.18359375}]

[19]: history = []
epochs = 30
max_lr = 0.01
grad_clip = 0.1
weight_decay = 1e-4
opt_func = torch.optim.Adam

[24]: %%time
history += fit_one_cycle(epochs, max_lr, model, train_dl, valid_dl,
grad_clip=grad_clip,
weight_decay=weight_decay,
opt_func=opt_func)

100%|
| 1230/1230 [01:52<00:00, 10.96it/s]

Epoch [0], last_lr: 0.00069, train_loss: 0.0354, val_loss: 0.5561, val_acc:
0.8723

100%|
| 1230/1230 [01:47<00:00, 11.41it/s]

Epoch [1], last_lr: 0.00152, train_loss: 0.0336, val_loss: 0.6916, val_acc:
0.7876

100%|
| 1230/1230 [01:47<00:00, 11.40it/s]

Epoch [2], last_lr: 0.00280, train_loss: 0.0401, val_loss: 0.6363, val_acc:
0.8089

100%|
| 1230/1230 [01:48<00:00, 11.37it/s]

```
Epoch [3] , last_lr: 0.00437, train_loss: 0.0377, val_loss: 0.5536, val_acc: 0.7935
100%|
| 1230/1230 [01:48<00:00, 11.30it/s]

Epoch [4] , last_lr: 0.00603, train_loss: 0.0376, val_loss: 0.4903, val_acc: 0.8391
100%|
| 1230/1230 [01:48<00:00, 11.29it/s]

Epoch [5] , last_lr: 0.00760, train_loss: 0.0350, val_loss: 0.4997, val_acc: 0.8700
100%|
| 1230/1230 [01:48<00:00, 11.37it/s]

Epoch [6] , last_lr: 0.00888, train_loss: 0.0374, val_loss: 0.7302, val_acc: 0.7943
100%|
| 1230/1230 [01:46<00:00, 11.55it/s]

Epoch [7] , last_lr: 0.00971, train_loss: 0.0383, val_loss: 0.4035, val_acc: 0.8584
100%|
| 1230/1230 [01:48<00:00, 11.35it/s]

Epoch [8] , last_lr: 0.01000, train_loss: 0.0353, val_loss: 1.0227, val_acc: 0.7544
100%|
| 1230/1230 [01:47<00:00, 11.43it/s]

Epoch [9] , last_lr: 0.00994, train_loss: 0.0308, val_loss: 0.2939, val_acc: 0.8990
100%|
| 1230/1230 [01:47<00:00, 11.46it/s]

Epoch [10] , last_lr: 0.00978, train_loss: 0.0338, val_loss: 0.8697, val_acc: 0.7615
100%|
| 1230/1230 [02:13<00:00, 9.23it/s]

Epoch [11] , last_lr: 0.00950, train_loss: 0.0315, val_loss: 0.2272, val_acc: 0.9158
100%|
| 1230/1230 [01:48<00:00, 11.31it/s]

Epoch [12] , last_lr: 0.00913, train_loss: 0.0274, val_loss: 0.3721, val_acc: 0.8663
```

```
100%|
| 1230/1230 [01:47<00:00, 11.47it/s]

Epoch [13], last_lr: 0.00867, train_loss: 0.0286, val_loss: 0.4039, val_acc:
0.8756

100%|
| 1230/1230 [01:47<00:00, 11.42it/s]

Epoch [14], last_lr: 0.00812, train_loss: 0.0268, val_loss: 0.3494, val_acc:
0.8897

100%|
| 1230/1230 [01:47<00:00, 11.48it/s]

Epoch [15], last_lr: 0.00750, train_loss: 0.0254, val_loss: 0.9066, val_acc:
0.7996

100%|
| 1230/1230 [01:49<00:00, 11.19it/s]

Epoch [16], last_lr: 0.00683, train_loss: 0.0229, val_loss: 0.3452, val_acc:
0.8859

100%|
| 1230/1230 [01:54<00:00, 10.74it/s]

Epoch [17], last_lr: 0.00611, train_loss: 0.0210, val_loss: 0.8006, val_acc:
0.7866

100%|
| 1230/1230 [01:52<00:00, 10.90it/s]

Epoch [18], last_lr: 0.00537, train_loss: 0.0190, val_loss: 0.6987, val_acc:
0.8074

100%|
| 1230/1230 [01:48<00:00, 11.31it/s]

Epoch [19], last_lr: 0.00463, train_loss: 0.0183, val_loss: 0.3725, val_acc:
0.8914

100%|
| 1230/1230 [01:49<00:00, 11.19it/s]

Epoch [20], last_lr: 0.00389, train_loss: 0.0164, val_loss: 0.5377, val_acc:
0.8873

100%|
| 1230/1230 [01:49<00:00, 11.27it/s]

Epoch [21], last_lr: 0.00317, train_loss: 0.0146, val_loss: 0.2741, val_acc:
0.9070

100%|
| 1230/1230 [01:48<00:00, 11.33it/s]
```

```
Epoch [22], last_lr: 0.00250, train_loss: 0.0131, val_loss: 0.2369, val_acc: 0.9158
100%|
| 1230/1230 [01:47<00:00, 11.40it/s]

Epoch [23], last_lr: 0.00188, train_loss: 0.0108, val_loss: 0.4028, val_acc: 0.8691
100%|
| 1230/1230 [01:47<00:00, 11.44it/s]

Epoch [24], last_lr: 0.00133, train_loss: 0.0086, val_loss: 0.3208, val_acc: 0.9123
100%|
| 1230/1230 [01:47<00:00, 11.44it/s]

Epoch [25], last_lr: 0.00087, train_loss: 0.0079, val_loss: 0.4474, val_acc: 0.8558
100%|
| 1230/1230 [01:47<00:00, 11.41it/s]

Epoch [26], last_lr: 0.00050, train_loss: 0.0064, val_loss: 0.4487, val_acc: 0.8563
100%|
| 1230/1230 [01:47<00:00, 11.43it/s]

Epoch [27], last_lr: 0.00022, train_loss: 0.0056, val_loss: 0.3277, val_acc: 0.8937
100%|
| 1230/1230 [01:47<00:00, 11.48it/s]

Epoch [28], last_lr: 0.00006, train_loss: 0.0049, val_loss: 0.4415, val_acc: 0.8615
100%|
| 1230/1230 [01:48<00:00, 11.38it/s]

Epoch [29], last_lr: 0.00000, train_loss: 0.0044, val_loss: 0.3111, val_acc: 0.9030
CPU times: total: 14min 27s
Wall time: 56min 41s
```

```
[40]: torch.cuda.empty_cache()
```


0.0.5 Save Model

```
[50]: torch.save(model.state_dict(), "distortion_classifier_resnet18-pretrained-90.pt")
```

0.0.6 Performance Graph


```
[46]: def plot_accuracies(history):
    accuracies = [x['val_acc'] for x in history]
    plt.plot(accuracies, '-x')
    plt.xlabel('epoch')
    plt.ylabel('accuracy')
    plt.title('Accuracy vs No. of epochs')

plot_accuracies(history)
```




```
[48]: def plot_losses(history):
    train_losses = [x['train_loss'] for x in history]
    val_losses = [x['val_loss'] for x in history]
    plt.plot(train_losses, '-bx')
    plt.plot(val_losses, '-rx')
    plt.xlabel('epoch')
    plt.ylabel('loss')
    plt.legend(['Training', 'Validation'])
    plt.title('Loss vs No. of epochs')
```

```
plot_losses(history)
```



```
[50]: def plot_lrs(history):
    lrs = np.concatenate([x.get('lrs', []) for x in history])
    plt.plot(lrs)
    plt.xlabel('Batch no.')
    plt.ylabel('Learning rate')
    plt.title('Learning Rate vs Batch no.')

plot_lrs(history)
```



0.0.7 Evaluate Model

```
[45]: from sklearn.metrics import classification_report, confusion_matrix  
import seaborn as sns
```

```
[46]: all_preds = []  
all_labels = []  
  
model.eval()  
with torch.no_grad():  
    for images, labels in valid_dl:  
        images = images.to(device)  
        labels = labels.to(device)  
        outputs = model(images)  
        _, preds = torch.max(outputs, 1)  
  
        all_preds.extend(preds.cpu().numpy())  
        all_labels.extend(labels.cpu().numpy())
```

```
[47]: class_names = train_ds.classes


print("Classification Report:")
print(classification_report(all_labels, all_preds, target_names=class_names))

print("\nConfusion Matrix:")
sns.heatmap(confusion_matrix(all_labels, all_preds), annot=True, fmt='d',
            xticklabels=class_names, yticklabels=class_names)
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
h.blur	0.98	0.91	0.94	214
low_light	1.00	1.00	1.00	146
low_qual	0.91	0.94	0.92	207
normal	0.94	0.99	0.97	188
v.blur	0.97	0.97	0.97	157
accuracy			0.96	912
macro avg	0.96	0.96	0.96	912
weighted avg	0.96	0.96	0.96	912

Confusion Matrix:




0.0.8 Predict Image

```
[36]: def predict_image(img, model):
    # Convert to a batch of 1
    xb = to_device(img.unsqueeze(0), device)
    # Get predictions from the model
    yb = model(xb)
    # Pick index with the highest probability
    _, preds = torch.max(yb, dim=1)
    # Retrieve the class label
    return train_ds.classes[preds[0].item()]

[64]: img, label = valid_ds[0]
denorm_img = denormalize(img, *stats)[0].permute(1, 2, 0)
plt.figure(figsize=(2,20))
plt.imshow(denorm_img)
print(denorm_img.shape)
print(f'Label: {train_ds.classes[label]}, Predicted: {predict_image(img, model)}')

torch.Size([128, 256, 3])
Label: h_blur, Predicted: h_blur
```




```
[88]: img, label = valid_ds[1010]
denorm_img = denormalize(img, *stats)[0].permute(1, 2, 0)
plt.figure(figsize=(2,20))
plt.imshow(denorm_img)
print(denorm_img.shape)
print(f'Label: {train_ds.classes[label]}, Predicted: {predict_image(img, model)}')

torch.Size([128, 256, 3])
Label: normal, Predicted: normal
```



```
[152]: img, label = test_ds[37]
denorm_img = denormalize(img, *stats)[0].permute(1, 2, 0)
plt.figure(figsize=(2,20))
plt.imshow(denorm_img)
print(denorm_img.shape)
print(f'Label: {train_ds.classes[label]}, Predicted: {predict_image(img, model)})')
```

torch.Size([128, 256, 3])
Label: h.blur, Predicted: normal



```
[70]: [evaluate(model, valid_d1)]
```

```
[70]: {'val_loss': 0.3463773727416992, 'val_acc': 0.9001893997192383}]
```

0.0.9 Test with images outside the dataset

```
[42]: from PIL import Image
```

```
[134]: h_blurs = os.listdir(os.path.join(DATA_DIR, 'test3', 'low_qual'))

for image_name in h_blurs:
    external_img = Image.open(os.path.join(DATA_DIR, 'test3', 'low_qual', image_name)).convert('RGB')
    img_tensor = transform(external_img)

    prediction = predict_image(img_tensor, model)
    if prediction != 'low_qual':
```

```
print(image_name)

11.jpg
15.jpg
20.jpg
23.jpg
26.jpg
4.jpg
IMG_3689.PNG
IMG_3690.PNG
IMG_3691.PNG
IMG_3751.PNG
IMG_3752.PNG
```

```
-----  
UnidentifiedImageError                                     Traceback (most recent call last)  
Cell In[134], line 4  
      1 h_blurs = os.listdir(os.path.join(DATA_DIR, 'test3', 'low_qual'))  
      2 for image_name in h_blurs:  
----> 4     external_img = Image.open(os.path.join(DATA_DIR,           ,           ,image_name)).  
      5         convert('RGB')  
      6     img_tensor = transform(external_img)  
      7     prediction = predict_image(img_tensor, model)  
  
File D:\User\Jansen\Self Study\2025 - 05 -  
↳ MAY\LiPAD\lipad-venv\Lib\site-packages\PIL\Image.py:3572, in open(fp, mode,  
↳ formats)  
3570     warnings.warn(message)  
3571 msg = "cannot identify image file %r" % (filename if filename else fp)  
-> 3572 raise UnidentifiedImageError(msg)  
  
UnidentifiedImageError: cannot identify image file 'C:  
↳ \\\\Thesis\\\\LiPAD\\\\data\\\\cnn\\\\test3\\\\low_qual\\IMG_3779(1).HEIC'
```


```
[44]: transform = T.Compose([  
    T.Resize([128, 256]),  
    T.ToTensor(),  
    T.Normalize(*stats)  
)  
  
model.eval()  
img_path = os.path.join(DATA_DIR, 'test', 'h.blur', '67.jpg')  
external_img = Image.open(img_path).convert('RGB')  
img_tensor = transform(external_img)
```

```
[112]: prediction = predict_image(img_tensor, model)
print(f'Predicted distortion: {prediction}')

denorm_img = denormalize(img_tensor, *stats)[0].permute(1, 2, 0)
plt.imshow(denorm_img)
plt.axis('off')
plt.title(f'Prediction: {prediction}')
plt.show()
```

Predicted distortion: low_qual

Prediction: low_qual



0.0.10 Load the model using .pt

```
[38]: model = DistortionClassifier18(5)
model_wts = torch.load('distortion_classifier.pt')
model.load_state_dict(model_wts)
model = to_device(model, device)
```

```
[45]: [evaluate(model, valid_dl)]
```

```
[45]: [{}{'val_loss': 0.010042956098914146, 'val_acc': 0.9973958134651184}]
```

```
[81]: torch.cuda.empty_cache()
```

0.0.11 Test with test dataset

```
[55]: test_ds = ImageFolder(os.path.join(DATA_DIR, 'test'), valid_tfms)
test_dl = DataLoader(test_ds, batch_size*5, num_workers=3, pin_memory=True)
test_dl = DeviceDataLoader(test_dl, device)

[56]: [evaluate(model, test_dl)]

[56]: {'val_loss': 0.25383108854293823, 'val_acc': 0.9370861053466797}

[172]: all_preds = []
all_labels = []

model.eval()
with torch.no_grad():
    for images, labels in test_dl:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, preds = torch.max(outputs, 1)

        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

class_names = train_ds.classes


print("Classification Report:")
print(classification_report(all_labels, all_preds, target_names=class_names))

print("\nConfusion Matrix:")
sns.heatmap(confusion_matrix(all_labels, all_preds), annot=True, fmt='d',
            xticklabels=class_names, yticklabels=class_names)
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
h_blur	0.96	0.92	0.94	114
low_light	1.00	1.00	1.00	46
low_qual	0.87	0.87	0.87	46
normal	0.86	0.97	0.92	39
v_blur	0.95	0.95	0.95	57
accuracy			0.94	302
macro avg	0.93	0.94	0.93	302
weighted avg	0.94	0.94	0.94	302

Confusion Matrix:



1 Transfer Learning

```
[31]: from torchvision import models
from torchvision.models import resnet18, ResNet18_Weights

class DistortionClassifier18(ImageClassificationBase):
    def __init__(self, num_classes):
        super().__init__()
        #weights = ResNet18_Weights.DEFAULT
        self.network = models.resnet18()
        self.network.fc = nn.Linear(self.network.fc.in_features, num_classes)

    def forward(self, xb):
        return self.network(xb)
```

```
[33]: model2 = DistortionClassifier18(10)
to_device(model2, device);
```

```
[247]: epochs = 20
max_lr = 0.01
grad_clip = 0.1
weight_decay = 1e-4
opt_func = torch.optim.Adam
```

```
[251]: %%time
history2 = fit_one_cycle(epochs, max_lr, model2, train_dl, valid_dl,
                         grad_clip=grad_clip,
                         weight_decay=weight_decay,
                         opt_func=opt_func)
```

100%|
| 929/929 [01:42<00:00, 9.09it/s]

Epoch [0], last_lr: 0.00104, train_loss: 0.1258, val_loss: 0.2412, val_acc: 0.9380

100%|
| 929/929 [01:40<00:00, 9.25it/s]

Epoch [1], last_lr: 0.00280, train_loss: 0.1270, val_loss: 0.2264, val_acc: 0.9372

100%|
| 929/929 [01:40<00:00, 9.28it/s]

Epoch [2], last_lr: 0.00520, train_loss: 0.1468, val_loss: 0.3524, val_acc: 0.9017

100%|
| 929/929 [01:40<00:00, 9.22it/s]

Epoch [3], last_lr: 0.00760, train_loss: 0.1496, val_loss: 0.2503, val_acc: 0.9252

100%|
| 929/929 [01:40<00:00, 9.24it/s]

Epoch [4], last_lr: 0.00936, train_loss: 0.1424, val_loss: 0.7712, val_acc: 0.8380

100%|
| 929/929 [01:45<00:00, 8.79it/s]

Epoch [5], last_lr: 0.01000, train_loss: 0.1470, val_loss: 0.2762, val_acc: 0.9431

100%|
| 929/929 [01:45<00:00, 8.82it/s]

Epoch [6], last_lr: 0.00987, train_loss: 0.1359, val_loss: 0.2224, val_acc: 0.9363

```
100%|
| 929/929 [01:44<00:00,  8.91it/s]

Epoch [7], last_lr: 0.00950, train_loss: 0.1272, val_loss: 0.2191, val_acc:
0.9357

100%|
| 929/929 [01:43<00:00,  8.95it/s]

Epoch [8], last_lr: 0.00891, train_loss: 0.1187, val_loss: 0.1375, val_acc:
0.9603

100%|
| 929/929 [01:43<00:00,  8.99it/s]

Epoch [9], last_lr: 0.00812, train_loss: 0.1105, val_loss: 0.1477, val_acc:
0.9572

100%|
| 929/929 [01:45<00:00,  8.80it/s]

Epoch [10], last_lr: 0.00717, train_loss: 0.1007, val_loss: 0.1843, val_acc:
0.9452

100%|
| 929/929 [01:45<00:00,  8.81it/s]

Epoch [11], last_lr: 0.00611, train_loss: 0.0907, val_loss: 0.1898, val_acc:
0.9521

100%|
| 929/929 [01:45<00:00,  8.80it/s]

Epoch [12], last_lr: 0.00500, train_loss: 0.0820, val_loss: 0.2127, val_acc:
0.9503

100%|
| 929/929 [01:45<00:00,  8.77it/s]

Epoch [13], last_lr: 0.00389, train_loss: 0.0754, val_loss: 0.1438, val_acc:
0.9594

100%|
| 929/929 [01:46<00:00,  8.69it/s]

Epoch [14], last_lr: 0.00283, train_loss: 0.0641, val_loss: 0.1024, val_acc:
0.9740

100%|
| 929/929 [01:47<00:00,  8.67it/s]

Epoch [15], last_lr: 0.00188, train_loss: 0.0535, val_loss: 0.0978, val_acc:
0.9736

100%|
| 929/929 [01:47<00:00,  8.66it/s]
```

```
Epoch [16], last_lr: 0.00109, train_loss: 0.0449, val_loss: 0.0994, val_acc: 0.9751
100%|
| 929/929 [01:45<00:00, 8.77it/s]

Epoch [17], last_lr: 0.00050, train_loss: 0.0380, val_loss: 0.0862, val_acc: 0.9796
100%|
| 929/929 [01:45<00:00, 8.81it/s]

Epoch [18], last_lr: 0.00013, train_loss: 0.0326, val_loss: 0.0806, val_acc: 0.9797
100%|
| 929/929 [01:46<00:00, 8.76it/s]

Epoch [19], last_lr: 0.00000, train_loss: 0.0307, val_loss: 0.0807, val_acc: 0.9799
CPU times: total: 8min 43s
Wall time: 39min 56s
```

```
[295]: torch.save(model2.state_dict(), "distortion_classifier_resnet18.pt")
```

```
[1]: !jupyter nbconvert --to pdf distortion_classifier.ipynb
```

```
[NbConvertApp] Converting notebook distortion_classifier.ipynb to pdf
[NbConvertApp] ERROR | Error while converting 'distortion_classifier.ipynb'
Traceback (most recent call last):
  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-packages\nbconvert\nbconvertapp.py", line 487, in export_single_notebook
    output, resources = self.exporter.from_filename(
                        ^^^^^^^^^^^^^^^^^^

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-packages\nbconvert\exporters\templateexporter.py", line 390, in from_filename
    return super().from_filename(filename, resources, **kw)  #
type:ignore[return-value]
                        ^^^^^^^^^^^^^^

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-packages\nbconvert\exporters\exporter.py", line 201, in from_filename
    return self.from_file(f, resources=resources, **kw)
                        ^^^^^^^^^^

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-packages\nbconvert\exporters\templateexporter.py", line 396, in from_file
    return super().from_file(file_stream, resources, **kw)  #
type:ignore[return-value]
                        ^^^^^^^^^^

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-packages\nbconvert\exporters\exporter.py", line 220, in from_file
    return self.from_notebook_node(
```

```

~~~~~
File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\nbconvert\exporters\pdf.py", line 184, in from_notebook_node
    latex, resources = super().from_notebook_node(nb, resources=resources, **kw)
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\nbconvert\exporters\latex.py", line 92, in from_notebook_node
    return super().from_notebook_node(nb, resources, **kw)
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\nbconvert\exporters\templateexporter.py", line 429, in
from_notebook_node
    output = self.template.render(nb=nb_copy, resources=resources)
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\jinja2\environment.py", line 1304, in render
    self.environment.handle_exception()
File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\jinja2\environment.py", line 939, in handle_exception
    raise rewrite_traceback_stack(source=source)
File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\index.tex.j2", line 8, in top-level
template code
    ((* extends cell_style *))
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\style_jupyter.tex.j2", line 176, in
top-level template code
    \prompt{{{{(prompt)}}}}{{{{(prompt_color)}}}}{{{{(execution_count)}}}}{{{{(extra_sp
ace))}}}
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\base.tex.j2", line 7, in top-level
template code
    ((*- extends 'document_contents.tex.j2' -*))
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\document_contents.tex.j2", line 51,
in top-level template code
    ((*- block figure scoped -*))
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\display_priority.j2", line 5, in
top-level template code
    ((*- extends 'null.j2' -*))
~~~~~

File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex>null.j2", line 30, in top-level

```

```

template code
  ((*- block body -*))
  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\base.tex.j2", line 241, in block
'body'
    ((( super() )))
  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\null.j2", line 32, in block 'body'
    ((*- block any_cell scoped -*))
    ~~~~~

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\null.j2", line 85, in block
'any_cell'
    ((*- block markdowncell scoped-*)) ((*- endblock markdowncell -*))
    ~~~~~

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-
venv\share\jupyter\nbconvert\templates\latex\document_contents.tex.j2", line 68,
in block 'markdowncell'
    ((( cell.source | citation2latex | strip_files_prefix |
convert_pandoc('markdown+tex_math_double_backslash', 'json', extra_args=[])
| resolve_references | convert_explicitly_relative_paths |
convert_pandoc('json', 'latex'))))
    ~~~~~

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\nbconvert\filters\pandoc.py", line 36, in convert_pandoc
    return pandoc(source, from_format, to_format, extra_args=extra_args)
    ~~~~~

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\nbconvert\utils\pandoc.py", line 50, in pandoc
    check_pandoc_version()
  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\nbconvert\utils\pandoc.py", line 98, in check_pandoc_version
    v = get_pandoc_version()
    ~~~~~

  File "D:\User\Jansen\Self Study\2025 - 05 - MAY\LiPAD\lipad-venv\Lib\site-
packages\nbconvert\utils\pandoc.py", line 75, in get_pandoc_version
    raise PandocMissing()
nbconvert.utils.pandoc.PandocMissing: Pandoc wasn't found.
Please check that pandoc is installed:
https://pandoc.org/installing.html

```

[]: