

```
In [1]: import os
import sys
from tqdm import tqdm
import numpy as np
import torch
import torch.nn as nn
from torch import optim
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
from torchvision.utils import make_grid
import torchvision.utils as vutils
import torchvision.transforms as T
import torchvision.transforms.functional as TF
import torch.nn.functional as F
from torchvision.utils import save_image
from PIL import Image
import matplotlib.pyplot as plt
%matplotlib inline
from paddleocr import PaddleOCR, draw_ocr
import paddle
from Levenshtein import distance as levenshtein_distance
from pytorch_msssim import ssim as ssim_fn
from lpips import LPIPS
from torchmetrics.image.psnr import PeakSignalNoiseRatio
from torchmetrics.image.ssim import StructuralSimilarityIndexMeasure
from torchmetrics.image.fid import FrechetInceptionDistance
```

```
In [2]: sys.path.append(os.path.abspath('../ ../../src/dataset'))
from paired_image_dataset import PairedImageDataset

DATA_DIR = os.path.join('..', '..', '..', 'data')
print(os.listdir(DATA_DIR))

['h blur', 'low light', 'low qual', 'test', 'v blur']
```

```
In [3]: image_width = 256  
        image_height = 128  
        batch_size = 16  
        stats = (0.5, 0.5, 0.5), (0.5, 0.5, 0.5)
```

```
In [4]: train_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)])  
  
valid_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)
])
```

```

valid_ds = PairedImageDataset(blur_dir=os.path.join(DATA_DIR, 'h.blur', 'v2', 'vali
                           normal_dir=os.path.join(DATA_DIR, 'h.blur', 'v2', 'va
                           transform=valid_transform)

test_ds = ImageFolder(os.path.join(DATA_DIR, 'h.blur', 'v2', 'test v2'), transform=

```

```
In [15]: train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=4)
valid_dl = DataLoader(valid_ds, batch_size=batch_size*2, shuffle=True, num_workers=4)
test_dl = DataLoader(test_ds, batch_size=2, shuffle=False, num_workers=4)
```

Helper functions

```

In [16]: def denorm(img_tensors):
           return img_tensors * stats[1][0] + stats[0][0]

def show_images(input_images, target_images, nmax=16):
    input_images = denorm(input_images.detach()[:nmax])
    target_images = denorm(target_images.detach()[:nmax])

    # Combine into a single tensor for visualization
    combined = torch.cat((input_images, target_images), 0)
    grid = make_grid(combined, nrow=nmax)

    plt.figure(figsize=(nmax, 4))
    plt.imshow(grid.permute(1, 2, 0))
    plt.axis("off")
    plt.title("Top: Horizontal Blur / Blurred | Bottom: Normal / Target")
    plt.show()

def show_batch(dl, nmax=16):
    for input_batch, target_batch in dl:
        show_images(input_batch, target_batch, nmax)
        break

```

```

In [17]: def show_paired_samples(dataset, num_samples=4):
           fig, axes = plt.subplots(num_samples, 2, figsize=(5, 2 * num_samples))

           for i in range(num_samples):
               input_img, target_img = dataset[i]

               input_img = denorm(input_img)
               target_img = denorm(target_img)

               if isinstance(input_img, torch.Tensor):
                   input_img = input_img.permute(1, 2, 0).numpy()
                   target_img = target_img.permute(1, 2, 0).numpy()

               axes[i, 0].imshow(input_img)
               axes[i, 0].set_title("Horizontal Blur (Input)")
               axes[i, 0].axis("off")

               axes[i, 1].imshow(target_img)

```

```
    axes[i, 1].set_title("Normal (Target)")
    axes[i, 1].axis("off")

plt.tight_layout()
plt.show()

def show_generated_images(images, nrow=4):
    images = denorm(images.detach().cpu())

    # Make grid
    grid_img = vutils.make_grid(images, nrow=nrow, normalize=False)

    # Show image
    plt.figure(figsize=(nrow * 2, 2 * (len(images) // nrow)))
    plt.axis('off')
    plt.imshow(grid_img.permute(1, 2, 0))
    plt.show()
```

```
In [18]: show_paired_samples(train_ds)
```



In [19]: `show_batch(valid_dl, nmax=14)`

Top: Horizontal Blur / Blurred | Bottom: Normal / Target



In [20]: `def print_images(image_tensor, num_images):`

```
    images = denorm(image_tensor)
    images = images.detach().cpu()
    image_grid = make_grid(images[:num_images], nrow=5)
```

```
plt.imshow(image_grid.permute(1, 2, 0).squeeze())
plt.show()
```

```
In [21]: def save_samples(generator, test_dl, epoch, sample_dir='./generated_test_outputs'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)

            for b in range(batch_size):
                img = fake_images[b]

                # create folder per image
                folder_path = os.path.join(sample_dir, str(global_idx))
                os.makedirs(folder_path, exist_ok=True)

                file_name = f'{global_idx}_{epoch}epoch.jpg'
                save_path = os.path.join(folder_path, file_name)

                save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
                global_idx += 1
```

```
In [22]: def get_default_device():
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    def __init__(self, dl, device):
        self(dl = dl
        self.device = device

    def __iter__(self):
        for b in self(dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self(dl)
```

```
In [23]: device = get_default_device()
device
```

```
Out[23]: device(type='cuda')
```

```
In [24]: train_dl = DeviceDataLoader(train_dl, device)
valid_dl = DeviceDataLoader(valid_dl, device)
test_dl = DeviceDataLoader(test_dl, device)
```

Building the Model

Generator

```
In [2]: class ResConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3, dropout=0.0, batch_norm=False):
        super().__init__()
        padding = kernel_size // 2

        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size, padding=padding)
        self.bn1 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size, padding=padding)
        self.bn2 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.drop = nn.Dropout(p=dropout) if dropout > 0 else nn.Identity()

        self.shortcut = nn.Conv2d(in_channels, out_channels, kernel_size=1, padding=0)
        self.bn_sc = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

    def forward(self, x):
        residual = x

        out = self.conv1(x)
        out = self.bn1(out)
        out = F.relu(out)

        out = self.conv2(out)
        out = self.bn2(out)
        out = self.drop(out)

        shortcut = self.shortcut(residual)
        shortcut = self.bn_sc(shortcut)

        out += shortcut
        out = F.relu(out)
        return out

class GatingSignal(nn.Module):
    def __init__(self, in_channels, out_channels, batch_norm=False):
        super().__init__()
```

```

        layers = [nn.Conv2d(in_channels, out_channels, kernel_size=1, padding=0)]
        if batch_norm:
            layers.append(nn.BatchNorm2d(out_channels))
        layers.append(nn.ReLU(inplace=True))
        self.gate = nn.Sequential(*layers)

    def forward(self, x):
        return self.gate(x)

class AttentionBlock(nn.Module):
    def __init__(self, in_channels, gating_channels, inter_channels):
        super().__init__()

        self.theta_x = nn.Conv2d(in_channels, inter_channels, kernel_size=2, stride=2)
        self.phi_g = nn.Conv2d(gating_channels, inter_channels, kernel_size=1, padding=0)

        self.upsample_g = nn.ConvTranspose2d(inter_channels, inter_channels, kernel_size=2, stride=2)
        self.combine = nn.Sequential(
            nn.ReLU(inplace=True),
            nn.Conv2d(inter_channels, 1, kernel_size=1),
            nn.Sigmoid())

        self.final_conv = nn.Sequential(
            nn.Conv2d(in_channels, in_channels, kernel_size=1),
            nn.BatchNorm2d(in_channels))

    def forward(self, x, g):
        # x: skip connection feature map
        # g: gating signal (decoder feature map)
        theta_x = self.theta_x(x)
        phi_g = self.phi_g(g)

        # Upsample gating to match theta_x
        upsample_g = self.upsample_g(phi_g)

        if upsample_g.shape != theta_x.shape:
            upsample_g = F.interpolate(upsample_g, size=theta_x.shape[2:], mode='bilinear')

        psi = self.combine(theta_x + upsample_g)
        psi = F.interpolate(psi, size=x.shape[2:], mode='bilinear', align_corners=True)
        psi = psi.expand(-1, x.shape[1], -1, -1)

        y = x * psi
        return self.final_conv(y)

class AttentionResUNetGenerator(nn.Module):
    def __init__(self, in_channels=3, out_channels=3, base_filters=64, dropout=0.0,
                 super().__init__()

        # Encoder
        self.down1 = ResConvBlock(in_channels, base_filters, dropout=dropout, batch_norm=True)
        self.pool1 = nn.MaxPool2d(2)
        self.down2 = ResConvBlock(base_filters, base_filters * 2, dropout=dropout, batch_norm=True)
        self.pool2 = nn.MaxPool2d(2)
        self.down3 = ResConvBlock(base_filters * 2, base_filters * 4, dropout=dropout, batch_norm=True)
        self.pool3 = nn.MaxPool2d(2)
        self.up1 = ResConvBlock(base_filters * 4, base_filters * 2, dropout=dropout, batch_norm=True)
        self.up2 = ResConvBlock(base_filters * 2, base_filters, dropout=dropout, batch_norm=True)
        self.up3 = ResConvBlock(base_filters, base_filters, dropout=dropout, batch_norm=True)
        self.out = nn.Conv2d(base_filters, out_channels, kernel_size=1)

```

```

        self.pool3 = nn.MaxPool2d(2)
        self.down4 = ResConvBlock(base_filters * 4, base_filters * 8, dropout=dropout)
        self.pool4 = nn.MaxPool2d(2)

    # Bottleneck
    self.bottleneck = ResConvBlock(base_filters * 8, base_filters * 16, dropout=dropout)

    # Gating and Attention
    self.gate4 = GatingSignal(base_filters * 16, base_filters * 8, batch_norm)
    self.attn4 = AttentionBlock(base_filters * 8, base_filters * 8, base_filters * 16)

    self.gate3 = GatingSignal(base_filters * 8, base_filters * 4, batch_norm) #
    self.attn3 = AttentionBlock(base_filters * 4, base_filters * 4, base_filters * 8)

    self.gate2 = GatingSignal(base_filters * 4, base_filters * 2, batch_norm) #
    self.attn2 = AttentionBlock(base_filters * 2, base_filters * 2, base_filters * 4)

    self.gate1 = GatingSignal(base_filters * 2, base_filters, batch_norm) # 128
    self.attn1 = AttentionBlock(base_filters, base_filters, base_filters) # 64

    # Decoder
    self.up4 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec4 = ResConvBlock(base_filters * 16 + base_filters * 8, base_filters * 8)

    self.up3 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec3 = ResConvBlock(base_filters * 8 + base_filters * 4, base_filters * 4)

    self.up2 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec2 = ResConvBlock(base_filters * 4 + base_filters * 2, base_filters * 2)

    self.up1 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec1 = ResConvBlock(base_filters * 2 + base_filters, base_filters, dropout=dropout)

    # Output Layer
    self.final_conv = nn.Sequential(
        nn.Conv2d(base_filters, out_channels, kernel_size=1),
        nn.Tanh())

def forward(self, x):
    d1 = self.down1(x)
    p1 = self.pool1(d1)
    d2 = self.down2(p1)
    p2 = self.pool2(d2)
    d3 = self.down3(p2)
    p3 = self.pool3(d3)
    d4 = self.down4(p3)
    p4 = self.pool4(d4)

    bn = self.bottleneck(p4)

    g4 = self.gate4(bn)
    a4 = self.attn4(d4, g4)
    u4 = self.up4(bn)
    u4 = torch.cat([u4, a4], dim=1)
    d5 = self.dec4(u4)

```

```

        g3 = self.gate3(d5)
        a3 = self.attn3(d3, g3)
        u3 = self.up3(d5)
        u3 = torch.cat([u3, a3], dim=1)
        d6 = self.dec3(u3)

        g2 = self.gate2(d6)
        a2 = self.attn2(d2, g2)
        u2 = self.up2(d6)
        u2 = torch.cat([u2, a2], dim=1)
        d7 = self.dec2(u2)

        g1 = self.gate1(d7)
        a1 = self.attn1(d1, g1)
        u1 = self.up1(d7)
        u1 = torch.cat([u1, a1], dim=1)
        d8 = self.dec1(u1)

    return self.final_conv(d8)

```

In [3]:

```

generator = AttentionResUNetGenerator()
x = torch.randn(4, 3, 128, 256) # Batch of horizontal motion blurred images
y = generator(x) # Deblurred output
print(y.shape) # Should be (4, 3, 128, 256)

```

torch.Size([4, 3, 128, 256])

Discriminator

In [27]:

```

class PatchDiscriminator(nn.Module):
    def __init__(self, in_channels=3, base_filters=64):
        super().__init__()
        layers = [
            nn.Conv2d(in_channels, base_filters, kernel_size=4, stride=2, padding=1,
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters, base_filters * 2, kernel_size=4, stride=2, padding=1,
                     nn.BatchNorm2d(base_filters * 2),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 2, base_filters * 4, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 4),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 4, base_filters * 8, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 8),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 8, 1, kernel_size=4, stride=2, padding=1),
            nn.Sigmoid()
        ]
        self.model = nn.Sequential(*layers)

```

```
def forward(self, x):
    return self.model(x)
```

```
In [28]: image1 = torch.rand((1, 3, 128, 256))

discriminator = PatchDiscriminator()
output = discriminator(image1)
print(output.shape)

torch.Size([1, 1, 4, 8])
```

Model Parameters

```
In [4]: total_params = sum(p.numel() for p in generator.parameters())
total_params
```

```
Out[4]: 39068999
```

```
In [30]: discriminator = to_device(discriminator, device)
generator = to_device(generator, device)
```

Training

Loss Functions

```
In [31]: comparison_loss = nn.BCEWithLogitsLoss()
L1_loss_fn = nn.L1Loss()
```

PaddleOCR

```
In [32]: ocr_model = PaddleOCR(use_angle_cls=False, lang='en', det=False, rec=True, use_gpu=False)

def tensor_to_bgr_numpy(tensor_img, mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5)):
    if tensor_img.dim() == 4:
        tensor_img = tensor_img[0]

    if tensor_img.shape[0] != 3:
        raise ValueError(f"Expected 3 channels, got shape {tensor_img.shape}")

    unnorm = torch.zeros_like(tensor_img)
    for c in range(3):
        unnorm[c] = tensor_img[c] * float(std[c]) + float(mean[c])
    tensor_img = unnorm.clamp(0, 1)

    pil_img = TF.to_pil_image(tensor_img.cpu())
    np_img = np.array(pil_img)[:, :, ::-1]

    return np_img
```

```
[2025/09/01 06:55:05] ppocr DEBUG: Namespace(help='==SUPPRESS==', use_gpu=True, use_xpu=False, use_npu=False, use_mlu=False, use_gcu=False, ir_optim=True, use_tensorrt=False, min_subgraph_size=15, precision='fp32', gpu_mem=500, gpu_id=0, image_dir=None, page_num=0, det_algorithm='DB', det_model_dir='C:\\Users\\ADMIN/.paddleocr/whl\\det\\en\\PP-OCRv3_det_infer', det_limit_side_len=960, det_limit_type='max', det_box_type='quad', det_db_thresh=0.3, det_db_box_thresh=0.6, det_db_unclip_ratio=1.5, max_batch_size=10, use_dilation=False, det_db_score_mode='fast', det_east_score_thresh=0.8, det_east_cover_thresh=0.1, det_east_nms_thresh=0.2, det_sast_score_thresh=0.5, det_sast_nms_thresh=0.2, det_pse_thresh=0, det_pse_box_thresh=0.85, det_pse_min_area=16, det_pse_scale=1, scales=[8, 16, 32], alpha=1.0, beta=1.0, fourier_degree=5, rec_algorithm='SVTR_LCNet', rec_model_dir='C:\\Users\\ADMIN/.paddleocr/whl\\rec\\en\\PP-OCRv4_rec_infer', rec_image_inverse=True, rec_image_shape='3, 48, 320', rec_batch_num=6, max_text_length=25, rec_char_dict_path='C:\\\\Thesis\\\\LiPAD with Paddle\\\\lipadenv3.9\\\\lib\\\\site-packages\\\\paddleocr\\\\ppocr\\\\utils\\\\en_dict.txt', use_space_char=True, vis_font_path='./doc/fonts/simfang.ttf', drop_score=0.5, e2e_algorithm='PGNet', e2e_model_dir=None, e2e_limit_side_len=768, e2e_limit_type='max', e2e_pgnet_score_thresh=0.5, e2e_char_dict_path='./ppocr/utils/ic15_dict.txt', e2e_pgnet_valid_set='totaltext', e2e_pgnet_mode='fast', use_angle_cls=False, cls_model_dir='C:\\\\Users\\ADMIN/.paddleocr/whl\\cls\\ch_ppocr_mobile_v2.0_cls_infer', cls_image_shape='3, 48, 192', label_list=['0', '180'], cls_batch_num=6, cls_thresh=0.9, enable_mkldnn=False, cpu_threads=10, use_pd-serving=False, warmup=False, sr_model_dir=None, sr_image_shape='3, 32, 128', sr_batch_num=1, draw_img_save_dir='./inference_results', save_crop_res=False, crop_res_save_dir='./output', use_mp=False, total_process_num=1, process_id=0, benchmark=False, save_log_path='./log_output/', show_log=True, use_onnx=False, onnx_providers=False, onnx_sess_options=False, return_word_box=False, output='./output', table_max_len=488, table_algorithm='TableAttn', table_model_dir=None, merge_no_span_structure=True, table_char_dict_path=None, formula_algorithm='LaTeXOCR', formula_model_dir=None, formula_char_dict_path=None, formula_batch_num=1, layout_model_dir=None, layout_dict_path=None, layout_score_threshold=0.5, layout_nms_threshold=0.5, kie_algorithm='LayoutXLM', ser_model_dir=None, re_model_dir=None, use_visual Backbone=True, ser_dict_path='../train_data/XFUND/class_list_xfun.txt', ocr_order_method=None, mode='structure', image_orientation=False, layout=True, table=True, formula=False, ocr=True, recovery=False, recovery_to_markdown=False, use_pdf2docx_api=False, invert=False, binarize=False, alphacolor=(255, 255, 255), lang='en', det=False, rec=True, type='ocr', savefile=False, ocr_version='PP-OCRv4', structure_version='PP-StructureV2')
```

```
[2025/09/01 06:55:05] ppocr WARNING: The first GPU is used for inference by default, GPU ID: 0
```

```
[2025/09/01 06:55:07] ppocr WARNING: The first GPU is used for inference by default, GPU ID: 0
```

Discriminator Training Function

```
In [33]: def train_discriminator(discriminator, generator, inputs, targets, d_opt):
    discriminator.train()

    real_images = targets
    fake_images = generator(inputs).detach()

    pred_real = discriminator(real_images)
    pred_fake = discriminator(fake_images)

    real_labels = torch.ones_like(pred_real)
    fake_labels = torch.zeros_like(pred_fake)
```

```

real_score = torch.sigmoid(pred_real).mean().item()
fake_score = torch.sigmoid(pred_fake).mean().item()

loss_real = comparison_loss(pred_real, real_labels)
loss_fake = comparison_loss(pred_fake, fake_labels)

d_loss = (loss_real + loss_fake) / 2

d_opt.zero_grad()
d_loss.backward()
d_opt.step()

return d_loss.item(), real_score, fake_score

```

Generator Training Function

```

In [34]: def ssim(pred, target):
    return 1 - ssim_fn(pred, target, data_range=1.0, size_average=True)

lpips_loss = LPIPS(net='vgg').to(device)
lpips_loss.eval()
for param in lpips_loss.parameters():
    param.requires_grad = False

def train_generator(generator, discriminator, inputs, targets, g_opt, adv_lambda=1.
generator.train()

fake_images = generator(inputs)
pred_fake = discriminator(fake_images)

real_labels = torch.ones_like(pred_fake)

adv_loss = comparison_loss(pred_fake, real_labels)
l1_loss = L1_loss_fn(fake_images, targets)
perceptual_loss = lpips_loss(fake_images.clamp(-1, 1), targets.clamp(-1, 1)).me
ssim_loss = ssim(fake_images, targets)

# Text Loss
fake_np = tensor_to_bgr_numpy(fake_images)
real_np = tensor_to_bgr_numpy(targets)

with torch.no_grad():
    fake_text = ocr_model.ocr(fake_np, cls=False, det=False)[0]
    real_text = ocr_model.ocr(real_np, cls=False, det=False)[0]

fake_str = fake_text[0][0] if fake_text else ''
real_str = real_text[0][0] if real_text else ''

fake_str = fake_str.strip().replace(' ', '')
real_str = real_str.strip().replace(' ', '')

# Use normalized edit distance as loss
edit_dist = levenshtein_distance(fake_str, real_str)
norm_edit_dist = edit_dist / max(len(real_str), 1)
text_loss = torch.tensor(norm_edit_dist, device=device)

```

```

g_loss = (
    adv_loss * adv_lambda +
    l1_loss * l1_lambda +
    ssim_loss * ssim_lambda +
    perceptual_loss * perceptual_lambda +
    text_loss * text_lambda
)

g_opt.zero_grad()
g_loss.backward()
g_opt.step()

return g_loss.item(), fake_images, {
    "total_loss": g_loss.item(),
    "adv": adv_loss.item(),
    "l1": l1_loss.item(),
    "ssim": ssim_loss.item(),
    "perceptual": perceptual_loss.item(),
    "text": text_loss.item()
}

```

Setting up [LPIPS] perceptual loss: trunk [vgg], v[0.1], spatial [off]
Loading model from: C:\Thesis\LiPAD with Paddle\lipadenv3.9\lib\site-packages\lpips\weights\v0.1\vgg.pth

```

In [35]: psnr_metric = PeakSignalNoiseRatio(data_range=1.0).to(device)
ssim_metric = StructuralSimilarityIndexMeasure(data_range=1.0).to(device)
fid_metric = FrechetInceptionDistance(feature=2048).to(device)

def to_uint8(tensor):
    # Clamp to [0, 1], scale to [0, 255], then convert to uint8
    tensor = torch.clamp(tensor, 0.0, 1.0)
    tensor = (tensor * 255.0).to(torch.uint8)
    return tensor

def evaluate_test(generator):
    raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl,
                                                                fig, axes = plt.subplots(1, 2, figsize=(10, 8)) # 1 row, 2 columns

    # Raw Test Set
    axes[0].imshow(raw_grid_img)
    axes[0].axis("off")
    axes[0].set_title("Raw Test Set")

    # Generated Test Set
    axes[1].imshow(gen_grid_img)
    axes[1].axis("off")
    axes[1].set_title("Generated Test Set")

    plt.tight_layout()
    plt.show()

```

```

In [36]: def create_image_grid_from_loader(generator, dataloader, device, num_images=16, image_size=256):

```

```

generator.to(device)

raw_images = []
gen_images = []
gen_count = 0
raw_count = 0

for batch in dataloader:
    # Assume batch is either just images or (images, labels)
    if isinstance(batch, (tuple, list)):
        inputs = batch[0]
    else:
        inputs = batch

    inputs = inputs.to(device)
    outputs = generator(inputs)

    for img in inputs:
        if denormalize:
            img = denorm(img)

        img_resized = TF.resize(img, image_size)
        raw_images.append(img_resized)

        raw_count += 1
        if raw_count >= num_images:
            break

    for out_img in outputs:
        if denormalize:
            out_img = denorm(out_img)

        out_img_resized = TF.resize(out_img, image_size)
        gen_images.append(out_img_resized)

        gen_count += 1
        if gen_count >= num_images:
            break

    if raw_count >= num_images:
        break

raw_grid = make_grid(raw_images, nrow=4)
gen_grid = make_grid(gen_images, nrow=4)

raw_ndarr = raw_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()
gen_ndarr = gen_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()

raw_grid_image = Image.fromarray(raw_ndarr)
gen_grid_image = Image.fromarray(gen_ndarr)

if save_path:
    gen_grid_image.save(save_path)
    print(f"Saved image grid to {save_path}")

return raw_grid_image, gen_grid_image

```

```
In [37]: @torch.no_grad()
def evaluate_generator(generator, valid_dl, epoch):
    generator.eval()
    psnr_vals = []
    ssim_vals = []
    fid_metric.reset()
    shown = False

    for val_inputs, val_targets in valid_dl:
        val_inputs = val_inputs.to(device)
        val_targets = val_targets.to(device)

        val_outputs = generator(val_inputs)

        if not shown:
            print('Validation Data')
            print_images(val_inputs, 5)
            print_images(val_outputs, 5)
            print_images(val_targets, 5)
            shown = True

        # Convert to uint8 for FID
        gen_uint8 = to_uint8(val_outputs)
        target_uint8 = to_uint8(val_targets)

        fid_metric.update(gen_uint8, real=False)
        fid_metric.update(target_uint8, real=True)

        # Compute PSNR/SSIM
        psnr = psnr_metric(val_outputs, val_targets).item()
        ssim = ssim_metric(val_outputs, val_targets).item()
        psnr_vals.append(psnr)
        ssim_vals.append(ssim)

        mean_psnr = sum(psnr_vals) / len(psnr_vals)
        mean_ssime = sum(ssim_vals) / len(ssim_vals)
        fid = fid_metric.compute().item()

        print(f"Epoch [{epoch+1}/{epochs}] - Mean PSNR: {mean_psnr:.4f}, Mean SSIM: {me

    evaluate_test(generator)
    return mean_psnr, mean_ssime, fid

def fit(generator, discriminator, epochs, lr, adv_lambda=1.0, l1_lambda=100.0, ssim_perceptual_lambda=5.0, text_lambda=7.5, g_opt=None, d_opt=None, checkpoint=None):
    torch.cuda.empty_cache()

    losses_g = []
    losses_d = []
    real_scores = []
    fake_scores = []
    psnrs = []
    ssims = []
    fids = []
```

```

g_losses_trace = []

# Create optimizers
if d_opt is None:
    d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
if g_opt is None:
    g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))

# Load checkpoint state if resuming
if checkpoint is not None:
    print("Resuming training from checkpoint...")
    d_opt.load_state_dict(checkpoint['opt_d'])
    g_opt.load_state_dict(checkpoint['opt_g'])

# Train
for epoch in range(80, epochs):
    torch.cuda.empty_cache()
    generator.train()
    discriminator.train()
    for inputs, targets in tqdm(train_dl):
        inputs = inputs.to(device)
        targets = targets.to(device)

        # Train discriminator
        d_loss, real_score, fake_score = train_discriminator(discriminator, gen)

        # Train generator
        g_loss, fake_images, g_loss_trace = train_generator(generator, discriminator,
                                                             l1_lambda, ssim_lambda, perceptual_lambda)

    print('Training Data')
    print_images(inputs, 5)
    print_images(fake_images, 5)
    print_images(targets, 5)

    if (epoch + 1) % 5 == 0:
        torch.save(generator.state_dict(), f'[RAU-NET]generator-Datasetv2-{epoch}.pt')
        torch.save(discriminator.state_dict(), f'[RAU-NET]discriminator-Datasetv2-{epoch}.pt')
        torch.save({
            'generator': generator.state_dict(),
            'discriminator': discriminator.state_dict(),
            'opt_g': g_opt.state_dict(),
            'opt_d': d_opt.state_dict(),
            'losses_g': losses_g,
            'losses_d': losses_d,
            'real_scores': real_scores,
            'fake_scores': fake_scores,
            'psnrs': psnrs,
            'ssims': ssims,
            'fids': fids,
            'epoch': epoch + 1
        }, f'[RAU-NET]Checkpoint-Logs-Datasetv2-{epoch + 1} epoch.pt')

# Print progress
try:
    print('Epoch [{}/{}], loss_g: {:.4f}, loss_d: {:.4f}, real_score: {:.4f}, fake_score: {:.4f}, psnr: {:.4f}, ssim: {:.4f}, fid: {:.4f}'.format(epoch, epochs, losses_g, losses_d, real_scores, fake_scores, psnrs, ssims, fids))

```

```

        epoch+1, epochs, g_loss, d_loss, real_score, fake_score))
except Exception as err:
    print('Error:', str(err))

mean_psnr, mean_ssim, fid = evaluate_generator(generator, valid_dl, epoch)

# Record losses and scores
losses_g.append(g_loss)
losses_d.append(d_loss)
real_scores.append(real_score)
fake_scores.append(fake_score)
psnrs.append(mean_psnr)
ssims.append(mean_ssim)
fids.append(fid)
g_losses_trace.append(g_loss_trace)

save_samples(generator, test_dl, epoch + 1, sample_dir='./visualization-2')

return losses_g, losses_d, real_scores, fake_scores, psnrs, ssims, fids, g_loss

```

Hyperparameter Configurations

```
In [38]: adv_lambda = 0.25 # Change to 0.25
l1_lambda = 100
ssim_lambda = 5 # Change to 5
perceptual_lambda = 5.0
text_lambda = 10 # Change to 10

lr = 0.00005
beta1 = 0.5
beta2 = 0.999

epochs = 100
```

```
In [30]: g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))
d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
```

```
In [31]: history = []
```

```
In [32]: %time
history += fit(generator,
                discriminator,
                epochs,
                lr,
                adv_lambda=adv_lambda,
                l1_lambda=l1_lambda,
                ssim_lambda=ssim_lambda,
                perceptual_lambda=perceptual_lambda,
                text_lambda=text_lambda,
                g_opt=g_opt,
                d_opt=d_opt,
                checkpoint=checkpoint)
```

```
NameError Traceback (most recent call last)
File <timed exec>:12
      NameError: name 'checkpoint' is not defined
```

Post Training

```
In [34]: torch.cuda.empty_cache()
```

Save Model

```
In [34]: torch.save(generator.state_dict(), "[RAU-NET]generator-Datasetv2-5epoch.pt")
torch.save(discriminator.state_dict(), "[RAU-NET]discriminator-Datasetv2-5epoch.pt"
```

```
In [35]: torch.save({
    'generator_state_dict': generator.state_dict(),
    'discriminator_state_dict': discriminator.state_dict(),
    'opt_g_state_dict': g_opt.state_dict(),
    'opt_d_state_dict': d_opt.state_dict(),
    'history': history
}, f'[RAU-NET]Checkpoint-Datasetv2-30epoch.pt')
```

Model Evaluation

```
In [ ]: psnr_vals = []
ssim_vals = []

fid_metric = FrechetInceptionDistance(feature=2048).to(device)
fid_metric.reset()

generator.eval()
for val_input, val_target in valid_dl:
    val_input = val_input.to(device)
    val_target = val_target.to(device)

    with torch.no_grad():
        gen_output = generator(val_input)

    # Convert to uint8 for FID
    gen_uint8 = to_uint8(gen_output)
    target_uint8 = to_uint8(val_target)

    fid_metric.update(gen_uint8, real=False)
    fid_metric.update(target_uint8, real=True)

    show_generated_images(gen_output)

# Compute PSNR/SSIM
```

```

psnr = psnr_metric(gen_output, val_target).item()
ssim = ssim_metric(gen_output, val_target).item()
psnr_vals.append(psnr)
ssim_vals.append(ssim)

print("Mean PSNR:", sum(psnr_vals) / len(psnr_vals))
print("Mean SSIM:", sum(ssim_vals) / len(ssim_vals))
print("FID:", fid_metric.compute().item())

```











```
In [41]: def deblur_image(image_path, generator, device):
    image = Image.open(image_path).convert('RGB')
    input_tensor = valid_transform(image).unsqueeze(0).to(device)

    generator.eval()
    with torch.no_grad():
        gen_image = generator(input_tensor)

    output = denorm(gen_image.squeeze(0).cpu())
```

```
plt.imshow(output.permute(1, 2, 0)) # CHW → HWC
plt.axis('off')
plt.title("Generated Image")
plt.show()
```

In [43]: deblur_image(os.path.join(DATA_DIR, 'h.blur', 'v2', 'test', 'h.blur', '135.jpg'), g

Generated Image



```
In [42]: raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl, device)

fig, axes = plt.subplots(1, 2, figsize=(15, 15)) # 1 row, 2 columns

# Raw Test Set
axes[0].imshow(raw_grid_img)
axes[0].axis("off")
axes[0].set_title("Raw Test Set")

# Generated Test Set
axes[1].imshow(gen_grid_img)
axes[1].axis("off")
axes[1].set_title("Generated Test Set")

plt.tight_layout()
plt.show()
```



```
In [31]: torch.cuda.empty_cache()
```

Load Model

```
In [39]: MODELS_DIR = os.path.join('..', '..', '..', 'models', 'h_blur')
```

```
In [40]: generator = AttentionResUNetGenerator()
loaded_g_wts = torch.load(os.path.join(MODELS_DIR, 'BEST MODELS', '0.25 Adv Loss Model.pt'))
generator.load_state_dict(loaded_g_wts)
generator = to_device(generator, device)
```

```
discriminator = PatchDiscriminator()
loaded_d_wts = torch.load(os.path.join(MODELS_DIR, 'BEST MODELS', '0.25 Adv Loss Model.pt'))
discriminator.load_state_dict(loaded_d_wts)
discriminator = to_device(discriminator, device)
```

```
In [30]: checkpoint = torch.load(os.path.join(MODELS_DIR, 'BEST MODELS', '0.25 Adv Loss Model.pt'))
#g_opt.load_state_dict(checkpoint['opt_g_state_dict'])
#d_opt.load_state_dict(checkpoint['opt_d_state_dict'])
#history = checkpoint['losses_g']
```

```
In [39]: def are_models_identical(model1, model2):
    # Check if state dictionaries have the same keys
    if model1.state_dict().keys() != model2.state_dict().keys():
        return False

    # Compare each parameter's values
    for key in model1.state_dict():
        if not torch.equal(model1.state_dict()[key], model2.state_dict()[key]):
            return False
    return True
```

```
In [45]: if are_models_identical(generator1, generator2):
    print("Models have identical weights and biases")
else:
    print("Models have different weights or biases")
```

Models have different weights or biases

```
In [9]: len(history)
```

```
Out[9]: 79
```

Save Reconstructed Test Set

```
In [12]: def save_generated_tests(generator, test_dl, sample_dir='./results'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)
```

```
        for b in range(batch_size):
            img = fake_images[b]

            file_name = f"{global_idx}.jpg"
            save_path = os.path.join(sample_dir, file_name)

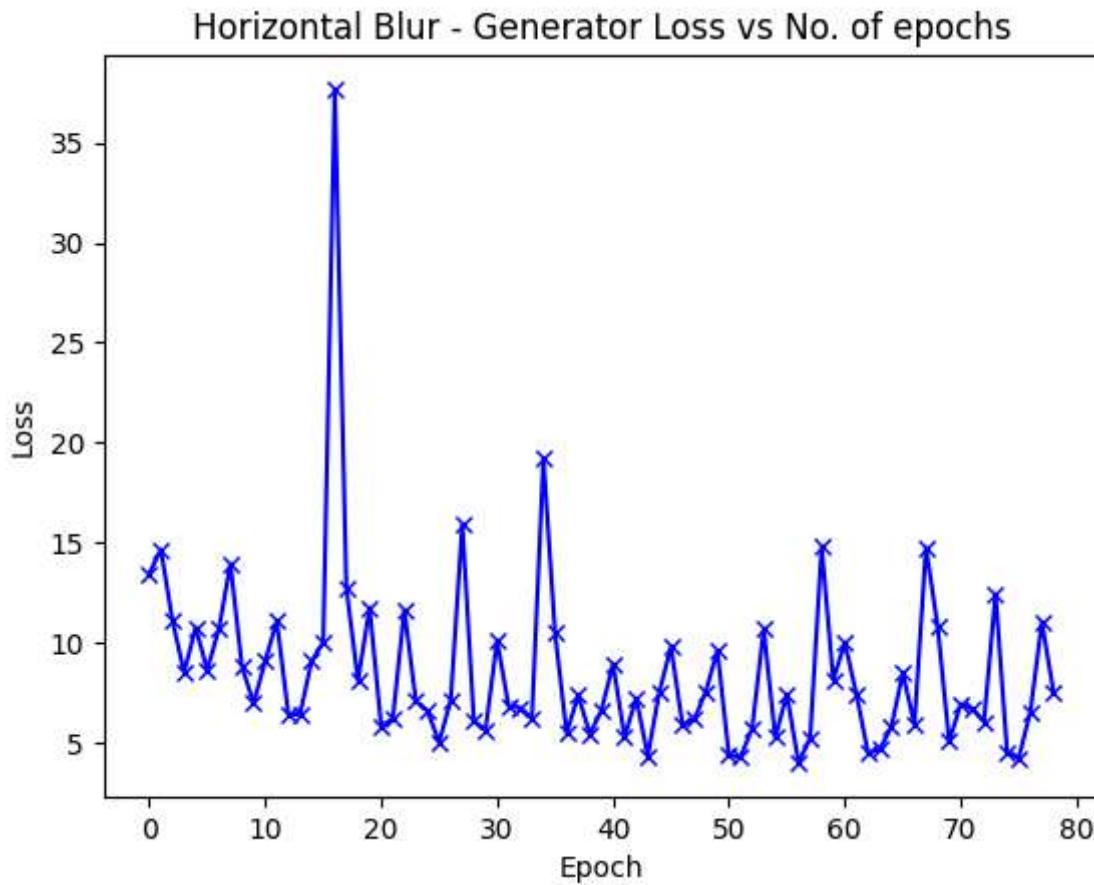
            save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
            global_idx += 1
```

```
In [40]: save_generated_tests(generator, test_dl)
```

Plots

```
In [68]: def plot_g_losses(history):
    losses_g = [x for x in history['losses_g']]
    plt.plot(losses_g, '-bx')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Horizontal Blur - Generator Loss vs No. of epochs')

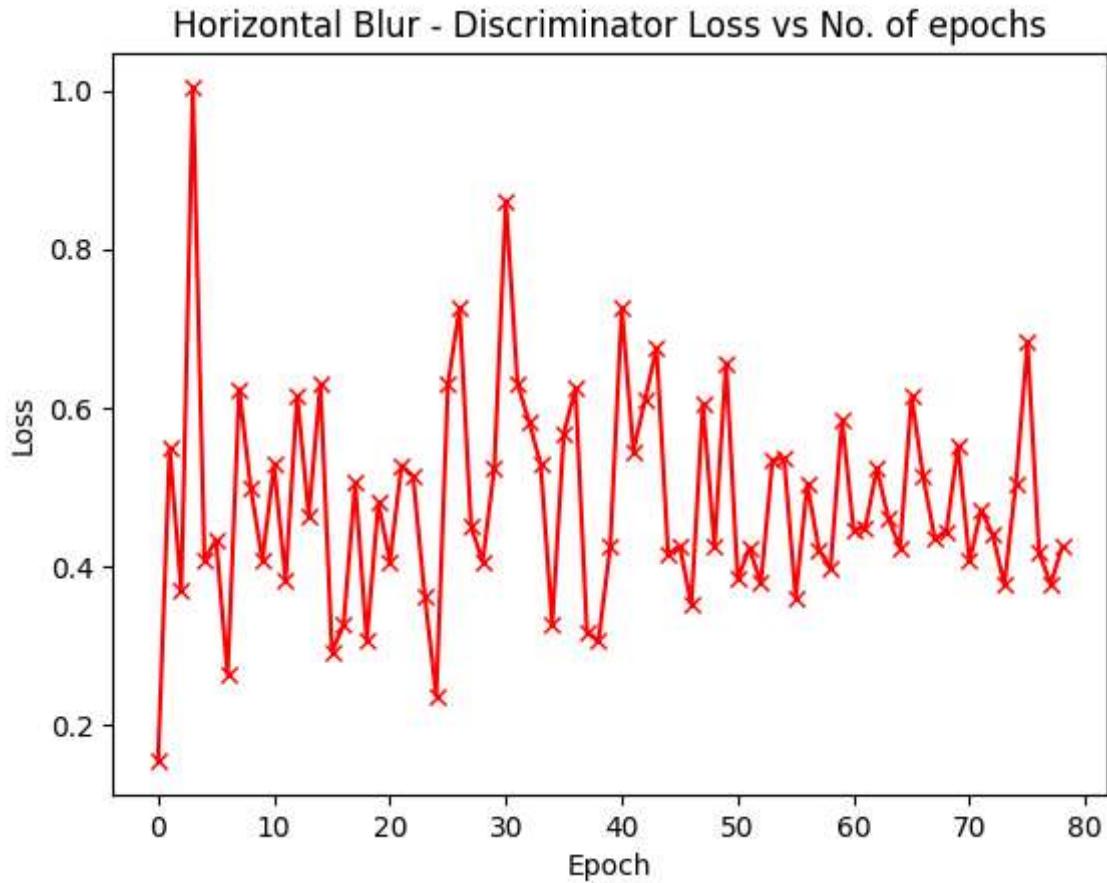
plot_g_losses(checkpoint)
```



```
In [71]: def plot_d_losses(history):
    losses_d = [x for x in history['losses_d']]
    plt.plot(losses_d, '-rx')
```

```
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Horizontal Blur - Discriminator Loss vs No. of epochs')

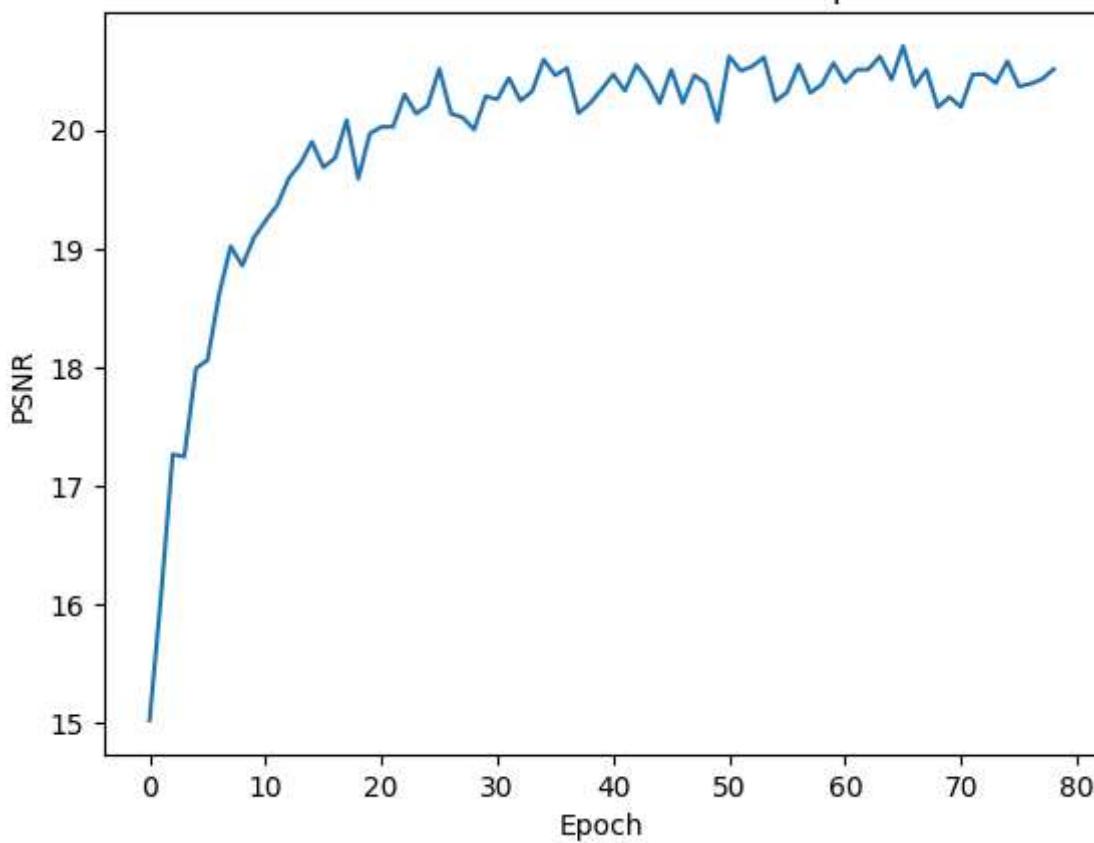
plot_d_losses(checkpoint)
```



```
In [62]: def plot_psnrs(history):
    psnrs = [x for x in history['psnrs']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('PSNR')
    plt.title('Horizontal Blur - PSNR vs No. of epochs')

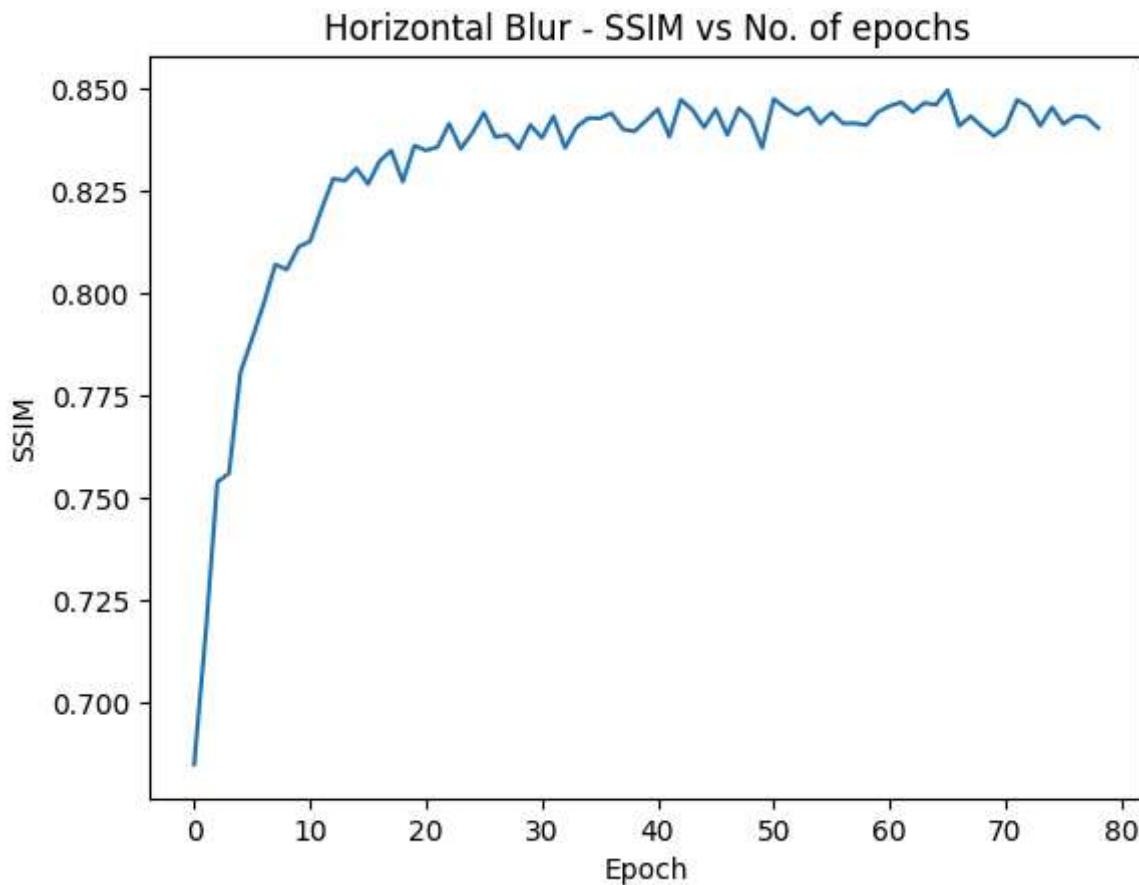
plot_psnrs(checkpoint)
```

Horizontal Blur - PSNR vs No. of epochs



```
In [65]: def plot_ssims(history):
    psnrs = [x for x in history['ssims']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('SSIM')
    plt.title('Horizontal Blur - SSIM vs No. of epochs')

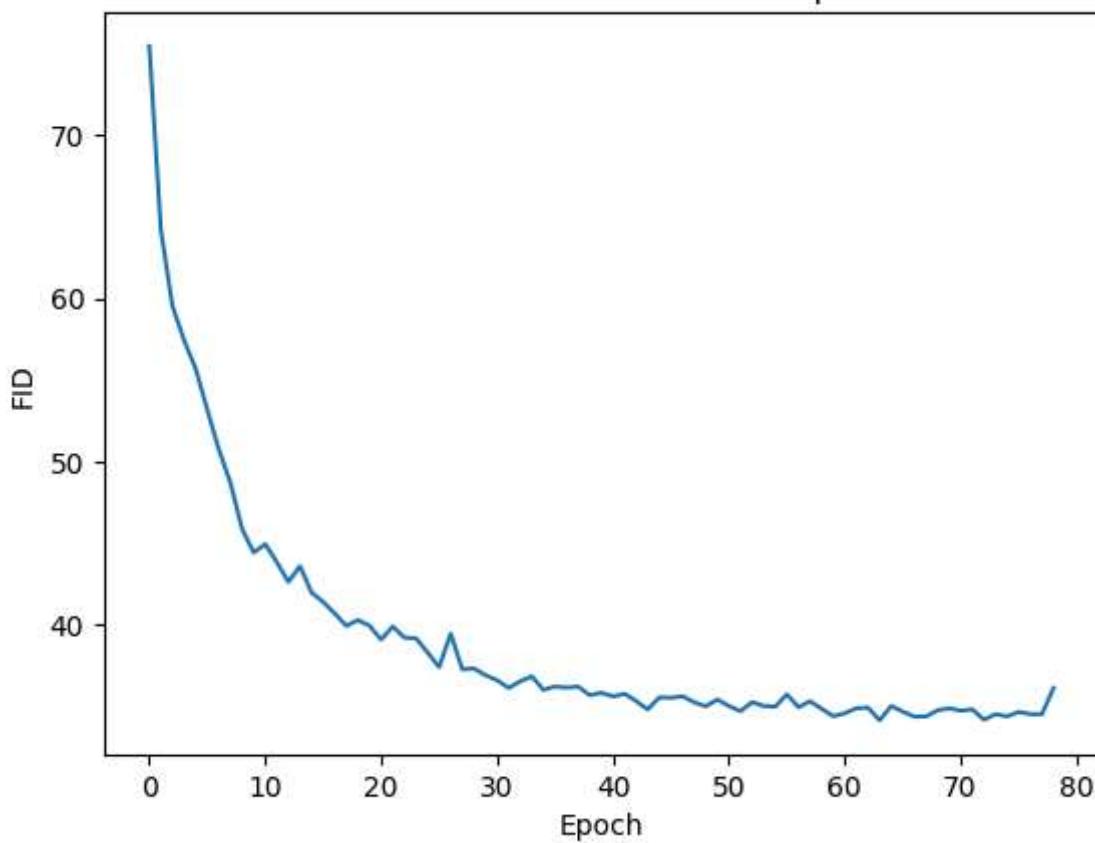
plot_ssims(checkpoint)
```



```
In [67]: def plot_fids(history):
    psnrs = [x for x in history['fids']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('FID')
    plt.title('Horizontal Blur - FID vs No. of epochs')

plot_fids(checkpoint)
```

Horizontal Blur - FID vs No. of epochs



In []:

In []: