

```
In [1]: import os
import sys
from tqdm import tqdm
import numpy as np
import torch
import torch.nn as nn
from torch.optim.lr_scheduler import ReduceLROnPlateau
from torch import optim
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
from torchvision.utils import make_grid
import torchvision.utils as vutils
import torchvision.transforms as T
import torchvision.transforms.functional as TF
import torch.nn.functional as F
from torchvision.utils import save_image
from PIL import Image
import matplotlib.pyplot as plt
%matplotlib inline
from paddleocr import PaddleOCR, draw_ocr
import paddle
from Levenshtein import distance as levenshtein_distance
from pytorch_msssim import ssim as ssim_fn
from lpips import LPIPS
from torchmetrics.image.psnr import PeakSignalNoiseRatio
from torchmetrics.image.ssim import StructuralSimilarityIndexMeasure
from torchmetrics.image.fid import FrechetInceptionDistance
```

```
In [2]: sys.path.append(os.path.abspath('../..../src/dataset'))
from paired_image_dataset import PairedImageDataset

DATA_DIR = os.path.join('..', '..', '..', 'data')
print(os.listdir(DATA_DIR))

['h_blur', 'low_light', 'low_qual', 'test', 'v_blur']
```

```
In [3]: image_width = 256
image_height = 128
batch_size = 16
stats = (0.5, 0.5, 0.5), (0.5, 0.5, 0.5)
```

```
In [5]: train_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)])

valid_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)
])
```

```
In [6]: train_ds = PairedImageDataset(blur_dir=os.path.join(DATA_DIR, 'low_light', 'train',
normal_dir=os.path.join(DATA_DIR, 'low_light', 'train'))
```

```

        transform=train_transform)

valid_ds = PairedImageDataset(blur_dir=os.path.join(DATA_DIR, 'low_light', 'valid',
                                                    normal_dir=os.path.join(DATA_DIR, 'low_light', 'valid',
                                                    transform=valid_transform))

test_ds = ImageFolder(os.path.join(DATA_DIR, 'low_light', 'test v2'), transform=val

```

```
In [7]: train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=4)
valid_dl = DataLoader(valid_ds, batch_size=batch_size*2, shuffle=True, num_workers=4)
test_dl = DataLoader(test_ds, batch_size=batch_size//2, shuffle=False, num_workers=4)
```

## Helper functions

```
In [8]: def denorm(img_tensors):
    return img_tensors * stats[1][0] + stats[0][0]

def show_images(input_images, target_images, nmax=16):
    input_images = denorm(input_images.detach()[:nmax])
    target_images = denorm(target_images.detach()[:nmax])

    # Combine into a single tensor for visualization
    combined = torch.cat((input_images, target_images), 0)
    grid = make_grid(combined, nrow=nmax)

    plt.figure(figsize=(nmax, 4))
    plt.imshow(grid.permute(1, 2, 0))
    plt.axis("off")
    plt.title("Top: Horizontal Blur / Blurred | Bottom: Normal / Target")
    plt.show()

def show_batch(dl, nmax=16):
    for input_batch, target_batch in dl:
        show_images(input_batch, target_batch, nmax)
        break
```

```
In [9]: def show_paired_samples(dataset, num_samples=4):
    fig, axes = plt.subplots(num_samples, 2, figsize=(5, 2 * num_samples))

    for i in range(num_samples):
        input_img, target_img = dataset[i]

        input_img = denorm(input_img)
        target_img = denorm(target_img)

        if isinstance(input_img, torch.Tensor):
            input_img = input_img.permute(1, 2, 0).numpy()
            target_img = target_img.permute(1, 2, 0).numpy()

        axes[i, 0].imshow(input_img)
        axes[i, 0].set_title("Horizontal Blur (Input)")
        axes[i, 0].axis("off")
```

```
axes[i, 1].imshow(target_img)
axes[i, 1].set_title("Normal (Target)")
axes[i, 1].axis("off")

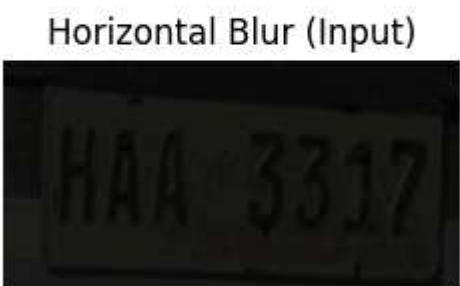
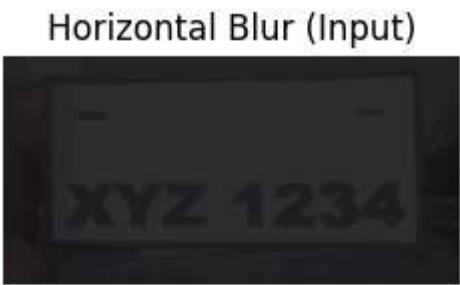
plt.tight_layout()
plt.show()

def show_generated_images(images, nrow=4):
    images = denorm(images.detach().cpu())

    # Make grid
    grid_img = vutils.make_grid(images, nrow=nrow, normalize=False)

    # Show image
    plt.figure(figsize=(nrow * 2, 2 * (len(images) // nrow)))
    plt.axis('off')
    plt.imshow(grid_img.permute(1, 2, 0))
    plt.show()
```

```
In [10]: show_paired_samples(train_ds)
```



```
In [11]: show_batch(valid_dl, nmax=14)
```

Top: Horizontal Blur / Blurred | Bottom: Normal / Target



```
In [12]: def print_images(image_tensor, num_images):
```

```
    images = denorm(image_tensor)
    images = images.detach().cpu()
    image_grid = make_grid(images[:num_images], nrow=5)
```

```
plt.imshow(image_grid.permute(1, 2, 0).squeeze())
plt.show()
```

```
In [13]: def save_samples(generator, test_dl, epoch, sample_dir='./generated_test_outputs'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)

            for b in range(batch_size):
                img = fake_images[b]

                # create folder per image
                folder_path = os.path.join(sample_dir, str(global_idx))
                os.makedirs(folder_path, exist_ok=True)

                file_name = f'{global_idx}_{epoch}epoch.jpg'
                save_path = os.path.join(folder_path, file_name)

                save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
                global_idx += 1
```

```
In [14]: def get_default_device():
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    def __init__(self, dl, device):
        self(dl = dl
        self.device = device

    def __iter__(self):
        for b in self(dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self(dl)
```

```
In [15]: device = get_default_device()
device
```

```
Out[15]: device(type='cuda')
```

```
In [16]: train_dl = DeviceDataLoader(train_dl, device)
valid_dl = DeviceDataLoader(valid_dl, device)
test_dl = DeviceDataLoader(test_dl, device)
```

# Building the Model

## Generator

```
In [17]: class ResConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3, dropout=0.0, batch_norm=False):
        super().__init__()
        padding = kernel_size // 2

        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size, padding=padding)
        self.bn1 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size, padding=padding)
        self.bn2 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.drop = nn.Dropout(p=dropout) if dropout > 0 else nn.Identity()

        self.shortcut = nn.Conv2d(in_channels, out_channels, kernel_size=1, padding=0)
        self.bn_sc = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

    def forward(self, x):
        residual = x

        out = self.conv1(x)
        out = self.bn1(out)
        out = F.relu(out)

        out = self.conv2(out)
        out = self.bn2(out)
        out = self.drop(out)

        shortcut = self.shortcut(residual)
        shortcut = self.bn_sc(shortcut)

        out += shortcut
        out = F.relu(out)
        return out

class GatingSignal(nn.Module):
    def __init__(self, in_channels, out_channels, batch_norm=False):
        super().__init__()
```



```

        self.pool3 = nn.MaxPool2d(2)
        self.down4 = ResConvBlock(base_filters * 4, base_filters * 8, dropout=dropout)
        self.pool4 = nn.MaxPool2d(2)

    # Bottleneck
    self.bottleneck = ResConvBlock(base_filters * 8, base_filters * 16, dropout=dropout)

    # Gating and Attention
    self.gate4 = GatingSignal(base_filters * 16, base_filters * 8, batch_norm)
    self.attn4 = AttentionBlock(base_filters * 8, base_filters * 8, base_filters * 16)

    self.gate3 = GatingSignal(base_filters * 8, base_filters * 4, batch_norm) #
    self.attn3 = AttentionBlock(base_filters * 4, base_filters * 4, base_filters * 8)

    self.gate2 = GatingSignal(base_filters * 4, base_filters * 2, batch_norm) #
    self.attn2 = AttentionBlock(base_filters * 2, base_filters * 2, base_filters * 4)

    self.gate1 = GatingSignal(base_filters * 2, base_filters, batch_norm) # 128
    self.attn1 = AttentionBlock(base_filters, base_filters, base_filters) # 64

    # Decoder
    self.up4 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec4 = ResConvBlock(base_filters * 16 + base_filters * 8, base_filters * 8)

    self.up3 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec3 = ResConvBlock(base_filters * 8 + base_filters * 4, base_filters * 4)

    self.up2 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec2 = ResConvBlock(base_filters * 4 + base_filters * 2, base_filters * 2)

    self.up1 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec1 = ResConvBlock(base_filters * 2 + base_filters, base_filters, dropout=dropout)

    # Output Layer
    self.final_conv = nn.Sequential(
        nn.Conv2d(base_filters, out_channels, kernel_size=1),
        nn.Tanh())

def forward(self, x):
    d1 = self.down1(x)
    p1 = self.pool1(d1)
    d2 = self.down2(p1)
    p2 = self.pool2(d2)
    d3 = self.down3(p2)
    p3 = self.pool3(d3)
    d4 = self.down4(p3)
    p4 = self.pool4(d4)

    bn = self.bottleneck(p4)

    g4 = self.gate4(bn)
    a4 = self.attn4(d4, g4)
    u4 = self.up4(bn)
    u4 = torch.cat([u4, a4], dim=1)
    d5 = self.dec4(u4)

```

```

        g3 = self.gate3(d5)
        a3 = self.attn3(d3, g3)
        u3 = self.up3(d5)
        u3 = torch.cat([u3, a3], dim=1)
        d6 = self.dec3(u3)

        g2 = self.gate2(d6)
        a2 = self.attn2(d2, g2)
        u2 = self.up2(d6)
        u2 = torch.cat([u2, a2], dim=1)
        d7 = self.dec2(u2)

        g1 = self.gate1(d7)
        a1 = self.attn1(d1, g1)
        u1 = self.up1(d7)
        u1 = torch.cat([u1, a1], dim=1)
        d8 = self.dec1(u1)

    return self.final_conv(d8)

```

In [18]:

```

generator = AttentionResUNetGenerator()
x = torch.randn(4, 3, 128, 256) # Batch of horizontal motion blurred images
y = generator(x) # Deblurred output
print(y.shape) # Should be (4, 3, 128, 256)

```

torch.Size([4, 3, 128, 256])

## Discriminator

In [19]:

```

class PatchDiscriminator(nn.Module):
    def __init__(self, in_channels=3, base_filters=64):
        super().__init__()
        layers = [
            nn.Conv2d(in_channels, base_filters, kernel_size=4, stride=2, padding=1,
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters, base_filters * 2, kernel_size=4, stride=2, padding=1,
                     nn.BatchNorm2d(base_filters * 2),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 2, base_filters * 4, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 4),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 4, base_filters * 8, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 8),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 8, 1, kernel_size=4, stride=2, padding=1),
        ]
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)

```

```
In [20]: image1 = torch.rand((1, 3, 128, 256))

discriminator = PatchDiscriminator()
output = discriminator(image1)
print(output.shape)

torch.Size([1, 1, 4, 8])
```

## Model Parameters

```
In [21]: total_params = sum(p.numel() for p in discriminator.parameters())
total_params
```

```
Out[21]: 2766529
```

```
In [22]: discriminator = to_device(discriminator, device)
generator = to_device(generator, device)
```

## Training

### Loss Functions

```
In [23]: comparison_loss = nn.BCEWithLogitsLoss()
L1_loss_fn = nn.L1Loss()
```

## PaddleOCR

```
In [24]: ocr_model = PaddleOCR(use_angle_cls=True, lang='en', use_gpu=True, show_log=False)

def tensor_to_bgr_numpy(tensor_img, mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5)):
    if tensor_img.dim() == 4:
        tensor_img = tensor_img[0]

    if tensor_img.shape[0] != 3:
        raise ValueError(f"Expected 3 channels, got shape {tensor_img.shape}")

    unnorm = torch.zeros_like(tensor_img)
    for c in range(3):
        unnorm[c] = tensor_img[c] * float(std[c]) + float(mean[c])
    tensor_img = unnorm.clamp(0, 1)

    pil_img = TF.to_pil_image(tensor_img.cpu())
    np_img = np.array(pil_img)[:, :, ::-1]

    return np_img
```

```
[2025/09/01 06:57:59] ppocr WARNING: The first GPU is used for inference by default,  
GPU ID: 0  
[2025/09/01 06:58:01] ppocr WARNING: The first GPU is used for inference by default,  
GPU ID: 0  
[2025/09/01 06:58:03] ppocr WARNING: The first GPU is used for inference by default,  
GPU ID: 0
```

## Discriminator Training Function

```
In [25]: def train_discriminator(discriminator, generator, inputs, targets, d_opt):  
    discriminator.train()  
  
    real_images = targets  
    fake_images = generator(inputs).detach()  
  
    pred_real = discriminator(real_images)  
    pred_fake = discriminator(fake_images)  
  
    real_labels = torch.ones_like(pred_real)  
    fake_labels = torch.zeros_like(pred_fake)  
  
    real_score = torch.sigmoid(pred_real).mean().item()  
    fake_score = torch.sigmoid(pred_fake).mean().item()  
  
    loss_real = comparison_loss(pred_real, real_labels)  
    loss_fake = comparison_loss(pred_fake, fake_labels)  
  
    d_loss = (loss_real + loss_fake) / 2  
  
    d_opt.zero_grad()  
    d_loss.backward()  
    d_opt.step()  
  
    return d_loss.item(), real_score, fake_score
```

## Generator Training Function

```
In [26]: def ssim(pred, target):  
    return 1 - ssim_fn(pred, target, data_range=1.0, size_average=True)  
  
lpips_loss = LPIPS(net='vgg').to(device)  
lpips_loss.eval()  
for param in lpips_loss.parameters():  
    param.requires_grad = False  
  
def train_generator(generator, discriminator, inputs, targets, g_opt, adv_lambda=1.  
    generator.train()  
  
    fake_images = generator(inputs)  
    pred_fake = discriminator(fake_images)  
  
    real_labels = torch.ones_like(pred_fake)  
  
    adv_loss = comparison_loss(pred_fake, real_labels)
```

```

l1_loss = L1_loss_fn(fake_images, targets)
perceptual_loss = lpips_loss(fake_images.clamp(-1, 1), targets.clamp(-1, 1)).mse
ssim_loss = ssim(fake_images, targets)

# Text Loss
fake_np = tensor_to_bgr_numpy(fake_images)
real_np = tensor_to_bgr_numpy(targets)

with torch.no_grad():
    fake_text = ocr_model.ocr(fake_np, cls=False)[0]
    real_text = ocr_model.ocr(real_np, cls=False)[0]

fake_str = fake_text[0][1][0] if fake_text else ''
real_str = real_text[0][1][0] if real_text else ''

fake_str = fake_str.strip().replace(' ', '')
real_str = real_str.strip().replace(' ', '')

# Use normalized edit distance as loss
edit_dist = levenshtein_distance(fake_str, real_str)
norm_edit_dist = edit_dist / max(len(real_str), 1)
text_loss = torch.tensor(norm_edit_dist, device=device)

g_loss = (
    adv_loss * adv_lambda +
    l1_loss * l1_lambda +
    ssim_loss * ssim_lambda +
    perceptual_loss * perceptual_lambda +
    text_loss * text_lambda
)

g_opt.zero_grad()
g_loss.backward()
g_opt.step()

return g_loss.item(), fake_images, {
    "total_loss": g_loss.item(),
    "adv": adv_loss.item(),
    "l1": l1_loss.item(),
    "ssim": ssim_loss.item(),
    "perceptual": perceptual_loss.item(),
    "text": text_loss.item()
}

```

Setting up [LPIPS] perceptual loss: trunk [vgg], v[0.1], spatial [off]  
Loading model from: C:\Thesis\LiPAD with Paddle\lipadvenv3.9\lib\site-packages\lpips\weights\v0.1\vgg.pth

```

In [27]: psnr_metric = PeakSignalNoiseRatio(data_range=1.0).to(device)
ssim_metric = StructuralSimilarityIndexMeasure(data_range=1.0).to(device)
fid_metric = FrechetInceptionDistance(feature=2048).to(device)

def to_uint8(tensor):
    # Clamp to [0, 1], scale to [0, 255], then convert to uint8
    tensor = torch.clamp(tensor, 0.0, 1.0)
    tensor = (tensor * 255.0).to(torch.uint8)

```

```

    return tensor

def evaluate_test(generator):
    raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl,
                                                               fig, axes = plt.subplots(1, 2, figsize=(10, 8)) # 1 row, 2 columns

    # Raw Test Set
    axes[0].imshow(raw_grid_img)
    axes[0].axis("off")
    axes[0].set_title("Raw Test Set")

    # Generated Test Set
    axes[1].imshow(gen_grid_img)
    axes[1].axis("off")
    axes[1].set_title("Generated Test Set")

    plt.tight_layout()
    plt.show()

```

In [28]:

```

def create_image_grid_from_loader(generator, dataloader, device, num_images=16, image_size=256):
    generator.eval()
    generator.to(device)

    raw_images = []
    gen_images = []
    gen_count = 0
    raw_count = 0

    for batch in dataloader:
        # Assume batch is either just images or (images, labels)
        if isinstance(batch, (tuple, list)):
            inputs = batch[0]
        else:
            inputs = batch

        inputs = inputs.to(device)
        outputs = generator(inputs)

        for img in inputs:
            if denormalize:
                img = denorm(img)

            img_resized = TF.resize(img, image_size)
            raw_images.append(img_resized)

            raw_count += 1
            if raw_count >= num_images:
                break

        for out_img in outputs:
            if denormalize:
                out_img = denorm(out_img)

            out_img_resized = TF.resize(out_img, image_size)
            gen_images.append(out_img_resized)

```

```

        gen_count += 1
        if gen_count >= num_images:
            break

    if raw_count >= num_images:
        break

raw_grid = make_grid(raw_images, nrow=4)
gen_grid = make_grid(gen_images, nrow=4)

raw_ndarr = raw_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()
gen_ndarr = gen_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()

raw_grid_image = Image.fromarray(raw_ndarr)
gen_grid_image = Image.fromarray(gen_ndarr)

if save_path:
    gen_grid_image.save(save_path)
    print(f"Saved image grid to {save_path}")

return raw_grid_image, gen_grid_image

```

In [29]:

```

@torch.no_grad()
def evaluate_generator(generator, valid_dl, epoch):
    generator.eval()
    psnr_vals = []
    ssim_vals = []
    fid_metric.reset()
    shown = False

    for val_inputs, val_targets in valid_dl:
        val_inputs = val_inputs.to(device)
        val_targets = val_targets.to(device)

        val_outputs = generator(val_inputs)

        if not shown:
            print('Validation Data')
            print_images(val_inputs, 5)
            print_images(val_outputs, 5)
            print_images(val_targets, 5)
            shown = True

        # Convert to uint8 for FID
        gen_uint8 = to_uint8(val_outputs)
        target_uint8 = to_uint8(val_targets)

        fid_metric.update(gen_uint8, real=False)
        fid_metric.update(target_uint8, real=True)

        # Compute PSNR/SSIM
        psnr = psnr_metric(val_outputs, val_targets).item()
        ssim = ssim_metric(val_outputs, val_targets).item()
        psnr_vals.append(psnr)
        ssim_vals.append(ssim)

```

```

        mean_psnr = sum(psnr_vals) / len(psnr_vals)
        mean_ssim = sum(ssim_vals) / len(ssim_vals)
        fid = fid_metric.compute().item()

        print(f"Epoch [{epoch+1}/{epochs}] - Mean PSNR: {mean_psnr:.4f}, Mean SSIM: {mean_ssim:.4f}, FID: {fid:.4f}")

    evaluate_test(generator)
    return mean_psnr, mean_ssim, fid

def get_lr(optimizer):
    return optimizer.param_groups[0]['lr']

def fit(generator, discriminator, epochs, lr, adv_lambda=1.0, l1_lambda=100.0, ssim_lambda=1.0,
        perceptual_lambda=5.0, text_lambda=7.5, g_opt=None, d_opt=None, checkpoint=None):
    torch.cuda.empty_cache()

    losses_g = []
    losses_d = []
    real_scores = []
    fake_scores = []
    psnrs = []
    ssims = []
    fids = []
    g_losses_trace = []
    g_lrs = []
    d_lrs = []

    # Create optimizers
    if d_opt is None:
        d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
    if g_opt is None:
        g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))

    # Load checkpoint state if resuming
    if checkpoint is not None:
        print("Resuming training from checkpoint...")
        d_opt.load_state_dict(checkpoint['opt_d'])
        g_opt.load_state_dict(checkpoint['opt_g'])

    g_scheduler = ReduceLROnPlateau(g_opt, mode='min', factor=0.5, patience=10, verbose=False)
    d_scheduler = ReduceLROnPlateau(d_opt, mode='min', factor=0.5, patience=12, verbose=False)

    # Train
    for epoch in range(epochs):
        torch.cuda.empty_cache()
        generator.train()
        discriminator.train()
        for inputs, targets in tqdm(train_dl):
            inputs = inputs.to(device)
            targets = targets.to(device)

            # Train discriminator
            d_loss, real_score, fake_score = train_discriminator(discriminator, generator, inputs, targets, device, g_opt, d_opt, adv_lambda, l1_lambda, ssim_lambda, perceptual_lambda, text_lambda)
            losses_d.append(d_loss.item())
            real_scores.append(real_score.item())
            fake_scores.append(fake_score.item())
            psnrs.append(psnr_fn(inputs, targets).item())
            ssims.append(ssim_fn(inputs, targets).item())
            fids.append(fid_fn(generator, discriminator, inputs, device).item())
            g_losses_trace.append(g_opt.state_dict()['state'][0]['step'])

        # Train generator
        g_opt.zero_grad()
        generator.train()
        for inputs, targets in tqdm(train_dl):
            inputs = inputs.to(device)
            targets = targets.to(device)

            # Train generator
            g_loss = train_generator(generator, discriminator, inputs, targets, device, g_opt, d_opt, adv_lambda, l1_lambda, ssim_lambda, perceptual_lambda, text_lambda)
            losses_g.append(g_loss.item())
            g_lrs.append(g_opt.state_dict()['state'][0]['step'])

        # Save checkpoint
        if (epoch + 1) % 5 == 0:
            checkpoint = {
                'epoch': epoch + 1,
                'opt_d': d_opt.state_dict(),
                'opt_g': g_opt.state_dict(),
                'losses_d': losses_d,
                'real_scores': real_scores,
                'fake_scores': fake_scores,
                'psnrs': psnrs,
                'ssims': ssims,
                'fids': fids,
                'g_losses_trace': g_losses_trace,
                'g_lrs': g_lrs,
                'd_lrs': d_lrs
            }
            torch.save(checkpoint, f'checkpoint_{epoch+1}.pt')

    # Evaluate test set
    mean_psnr, mean_ssim, fid = evaluate_test(generator)
    print(f"Epoch [{epoch+1}/{epochs}] - Mean PSNR: {mean_psnr:.4f}, Mean SSIM: {mean_ssim:.4f}, FID: {fid:.4f}")

    evaluate_test(generator)
    return mean_psnr, mean_ssim, fid

```

```

# Train generator
g_loss, fake_images, g_loss_trace = train_generator(generator, discriminator,
                                                    kl_lambda, ssim_lambda, perceptron)

print('Training Data')
print_images(inputs, 5)
print_images(fake_images, 5)
print_images(targets, 5)

if (epoch + 1) % 5 == 0:
    torch.save(generator.state_dict(), f'[RAU-NET OCRDET LR]generator-Datas{epoch}.pt')
    torch.save(discriminator.state_dict(), f'[RAU-NET OCRDET LR]discriminator-{epoch}.pt')
    torch.save({
        'generator': generator.state_dict(),
        'discriminator': discriminator.state_dict(),
        'opt_g': g_opt.state_dict(),
        'opt_d': d_opt.state_dict(),
        'losses_g': losses_g,
        'losses_d': losses_d,
        'real_score': real_scores,
        'fake_score': fake_scores,
        'psnrs': psnrs,
        'ssims': ssims,
        'fids': fids,
        'g_lrs': g_lrs,
        'd_lrs': d_lrs,
        'epoch': epoch + 1
    }, f'[RAU-NET OCRDET LR]Checkpoint-Logs-Datasetv2-{epoch + 1} epoch.pt')

# Print progress
try:
    curr_glr = get_lr(g_opt)
    curr_dlr = get_lr(d_opt)

    print(f'Epoch [{epoch}/{epochs}], loss_g: {g_loss:.4f}, loss_d: {d_loss:.4f}, real_score: {real_scores[-1]:.4f}')
    print(f"Current LRs - Generator: {curr_glr:.6f}, Discriminator: {curr_dlr:.6f}")
except Exception as err:
    print('Error:', str(err))

mean_psnr, mean_ssims, fid = evaluate_generator(generator, valid_dl, epoch)

g_scheduler.step(g_loss)
d_scheduler.step(d_loss)

# Record Losses and scores
losses_g.append(g_loss)
losses_d.append(d_loss)
real_scores.append(real_scores[-1])
fake_scores.append(fake_scores[-1])
psnrs.append(mean_psnr)
ssims.append(mean_ssims)
fids.append(fid)
g_losses_trace.append(g_loss_trace)
g_lrs.append(curr_glr)
d_lrs.append(curr_dlr)

```

```
    save_samples(generator, test_dl, epoch + 1, sample_dir='./visualization-ocr')

    return losses_g, losses_d, real_scores, fake_scores, psnrs, ssims, fids, g_loss
```

## Hyperparameter Configurations

```
In [30]: adv_lambda = 0.25 # Change to 0.25
        l1_lambda = 100
        ssim_lambda = 5 # Change to 5
        perceptual_lambda = 5.0
        text_lambda = 10 # Change to 10

        lr = 0.0002
        beta1 = 0.5
        beta2 = 0.999

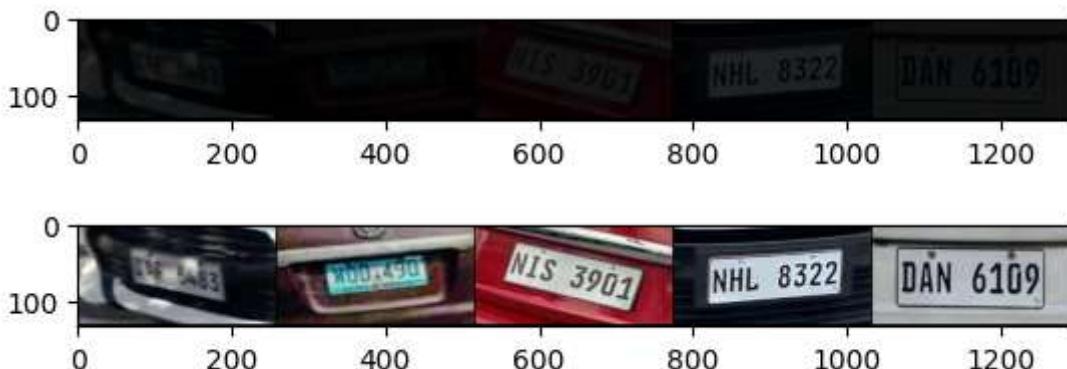
        epochs = 200
```

```
In [30]: g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))
        d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
```

```
In [31]: history = []
```

```
In [32]: %time
history += fit(generator,
                discriminator,
                epochs,
                lr,
                adv_lambda=adv_lambda,
                l1_lambda=l1_lambda,
                ssim_lambda=ssim_lambda,
                perceptual_lambda=perceptual_lambda,
                text_lambda=text_lambda,
                g_opt=g_opt,
                d_opt=d_opt)
```

100% |  
1250/1250 [19:52<00:00, 1.05it/s]  
Training Data





Epoch [1/200], loss\_g: 18.2146, loss\_d: 0.6682, real\_score: 0.3407, fake\_score: 0.2039

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [1/200] - Mean PSNR: 16.8681, Mean SSIM: 0.7625, FID: 54.599266052246094

Raw Test Set

Generated Test Set



100% |

1250/1250 [19:37<00:00, 1.06it/s]

Training Data



Epoch [2/200], loss\_g: 10.7149, loss\_d: 0.4601, real\_score: 0.9112, fake\_score: 0.54  
72

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [2/200] - Mean PSNR: 17.5805, Mean SSIM: 0.7811, FID: 51.86952209472656

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:46<00:00, 1.05it/s]

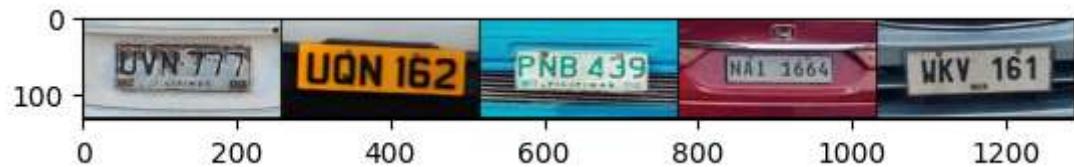
Training Data



Epoch [3/200], loss\_g: 10.5099, loss\_d: 0.5808, real\_score: 0.6860, fake\_score: 0.53  
17

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



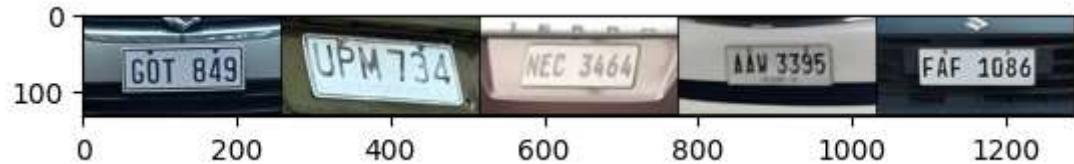
Epoch [3/200] - Mean PSNR: 17.1474, Mean SSIM: 0.7727, FID: 48.428184509277344

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:16<00:00, 1.08it/s]  
Training Data



Epoch [4/200], loss\_g: 10.5624, loss\_d: 0.4266, real\_score: 0.5125, fake\_score: 0.14  
28

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [4/200] - Mean PSNR: 17.0121, Mean SSIM: 0.7601, FID: 49.03854751586914

Raw Test Set

Generated Test Set



100%

1250/1250 [19:04<00:00, 1.09it/s]

Training Data



Epoch [5/200], loss\_g: 12.2944, loss\_d: 0.9898, real\_score: 0.9777, fake\_score: 0.84

44

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [5/200] - Mean PSNR: 18.3938, Mean SSIM: 0.7984, FID: 46.98902130126953

Raw Test Set

Generated Test Set



100% |

1250/1250 [19:17<00:00, 1.08it/s]

Training Data

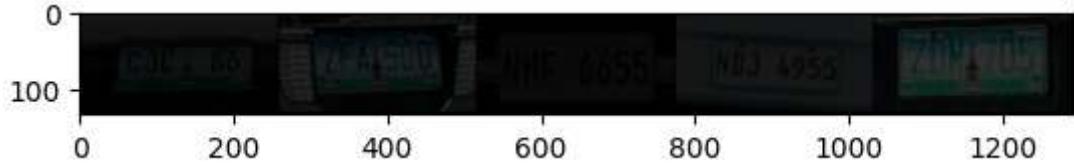


Epoch [6/200], loss\_g: 9.6396, loss\_d: 1.1877, real\_score: 0.1133, fake\_score: 0.110

2

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [6/200] - Mean PSNR: 18.2961, Mean SSIM: 0.8042, FID: 46.424007415771484



100% |  
1250/1250 [19:42<00:00, 1.06it/s]  
Training Data



Epoch [7/200], loss\_g: 8.1887, loss\_d: 0.4803, real\_score: 0.6531, fake\_score: 0.386  
4

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [7/200] - Mean PSNR: 18.4826, Mean SSIM: 0.8010, FID: 44.67730712890625



100% |  
1250/1250 [19:40<00:00, 1.06it/s]  
Training Data



Epoch [8/200], loss\_g: 13.6996, loss\_d: 0.7281, real\_score: 0.5370, fake\_score: 0.5254

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [8/200] - Mean PSNR: 18.7647, Mean SSIM: 0.8087, FID: 43.99444580078125

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:20<00:00, 1.08it/s]

Training Data





Epoch [9/200], loss\_g: 8.4032, loss\_d: 0.4975, real\_score: 0.6807, fake\_score: 0.4400

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [9/200] - Mean PSNR: 19.0534, Mean SSIM: 0.8129, FID: 42.22925567626953

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:06<00:00, 1.09it/s]  
Training Data





Epoch [10/200], loss\_g: 9.9892, loss\_d: 0.5280, real\_score: 0.9569, fake\_score: 0.6030

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [10/200] - Mean PSNR: 18.9131, Mean SSIM: 0.8105, FID: 43.680694580078125

Raw Test Set

Generated Test Set



100% |

1250/1250 [23:27<00:00, 1.13s/it]

Training Data



Epoch [11/200], loss\_g: 10.6698, loss\_d: 0.3221, real\_score: 0.8033, fake\_score: 0.3  
358

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [11/200] - Mean PSNR: 18.8334, Mean SSIM: 0.8119, FID: 43.634002685546875

Raw Test Set

Generated Test Set



100% |  
1250/1250 [23:12<00:00, 1.11s/it]

Training Data



Epoch [12/200], loss\_g: 9.5788, loss\_d: 0.7207, real\_score: 0.5507, fake\_score: 0.54  
29

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [12/200] - Mean PSNR: 19.4375, Mean SSIM: 0.8178, FID: 42.24984359741211

Raw Test Set

Generated Test Set



100% |  
1250/1250 [21:25<00:00, 1.03s/it]

Training Data



Epoch [13/200], loss\_g: 7.0039, loss\_d: 0.7567, real\_score: 0.2811, fake\_score: 0.15  
34

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [13/200] - Mean PSNR: 18.8323, Mean SSIM: 0.8165, FID: 41.2432861328125

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:27<00:00, 1.07it/s]

Training Data



Epoch [14/200], loss\_g: 8.4413, loss\_d: 0.5858, real\_score: 0.4340, fake\_score: 0.2354

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [14/200] - Mean PSNR: 19.4585, Mean SSIM: 0.8214, FID: 41.35641860961914

Raw Test Set

Generated Test Set



100% |

1250/1250 [19:31<00:00, 1.07it/s]

Training Data



Epoch [15/200], loss\_g: 7.9358, loss\_d: 0.6933, real\_score: 0.8897, fake\_score: 0.6886

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [15/200] - Mean PSNR: 18.6810, Mean SSIM: 0.8018, FID: 43.930599212646484



100% |  
1250/1250 [18:53<00:00, 1.10it/s]  
Training Data



Epoch [16/200], loss\_g: 9.5477, loss\_d: 0.4396, real\_score: 0.5152, fake\_score: 0.1606

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [16/200] - Mean PSNR: 18.7868, Mean SSIM: 0.8106, FID: 41.41378402709961



100% |  
1250/1250 [19:01<00:00, 1.10it/s]  
Training Data



Epoch [17/200], loss\_g: 9.0014, loss\_d: 0.3868, real\_score: 0.6185, fake\_score: 0.2206

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [17/200] - Mean PSNR: 16.7515, Mean SSIM: 0.7782, FID: 43.40441131591797

Raw Test Set

Generated Test Set



100% |  
1250/1250 [22:56<00:00, 1.10s/it]

Training Data





Epoch [18/200], loss\_g: 7.5288, loss\_d: 0.4232, real\_score: 0.7459, fake\_score: 0.3966

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [18/200] - Mean PSNR: 19.5137, Mean SSIM: 0.8267, FID: 40.65325164794922

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:52<00:00, 1.10it/s]

Training Data

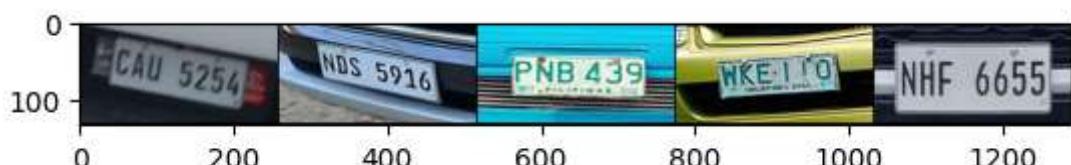
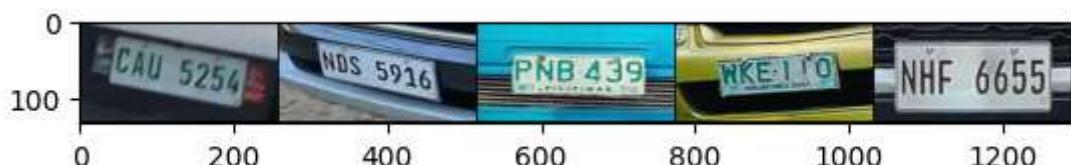
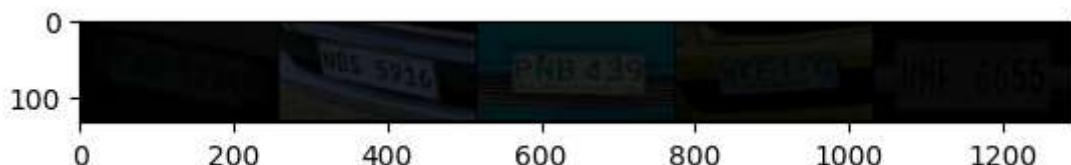




Epoch [19/200], loss\_g: 9.3909, loss\_d: 0.3903, real\_score: 0.6573, fake\_score: 0.27  
24

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [19/200] - Mean PSNR: 18.5926, Mean SSIM: 0.8067, FID: 41.4069938659668

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [20/200], loss\_g: 18.4405, loss\_d: 0.5489, real\_score: 0.6026, fake\_score: 0.4075

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [20/200] - Mean PSNR: 19.0711, Mean SSIM: 0.8102, FID: 43.037315368652344

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:31<00:00, 1.07it/s]

Training Data



Epoch [21/200], loss\_g: 8.5958, loss\_d: 0.5092, real\_score: 0.5487, fake\_score: 0.3047

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [21/200] - Mean PSNR: 18.8582, Mean SSIM: 0.8164, FID: 38.688232421875

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:58<00:00, 1.10it/s]

Training Data



Epoch [22/200], loss\_g: 6.5751, loss\_d: 0.6358, real\_score: 0.5311, fake\_score: 0.4137

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [22/200] - Mean PSNR: 19.3957, Mean SSIM: 0.8226, FID: 38.86845016479492  
Raw Test Set Generated Test Set



100% |  
1250/1250 [18:52<00:00, 1.10it/s]  
Training Data



Epoch [23/200], loss\_g: 9.2164, loss\_d: 0.4331, real\_score: 0.5851, fake\_score: 0.23  
46

Current LRs – Generator: 0.000200, Discriminator: 0.000200  
Validation Data





Epoch [23/200] - Mean PSNR: 19.2563, Mean SSIM: 0.8223, FID: 39.82571792602539

Raw Test Set

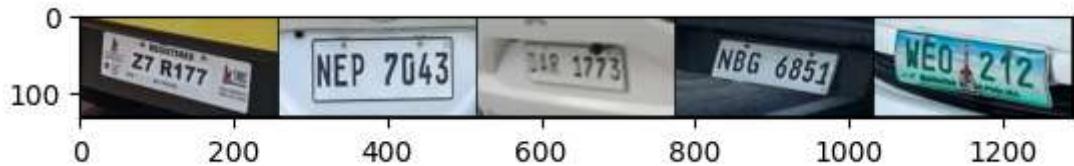
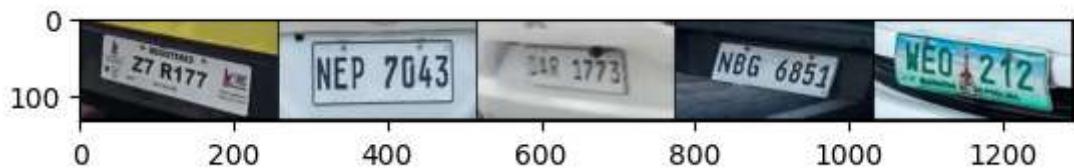
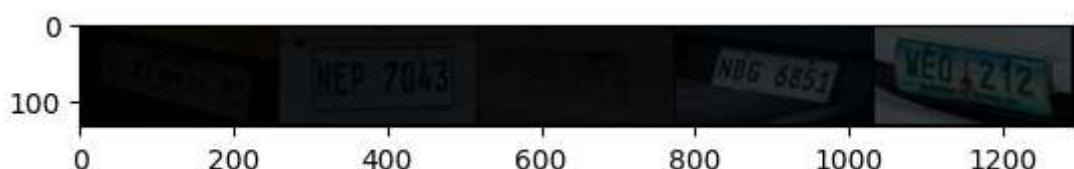
Generated Test Set



100%

1250/1250 [18:52<00:00, 1.10it/s]

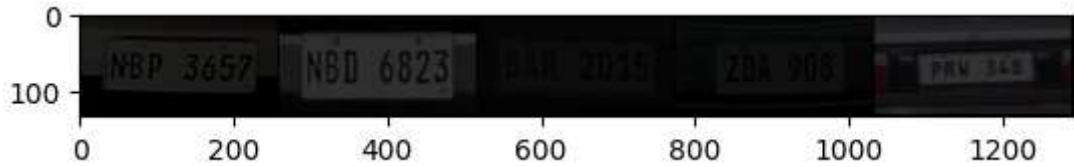
Training Data



Epoch [24/200], loss\_g: 6.6367, loss\_d: 0.4406, real\_score: 0.6854, fake\_score: 0.35  
96

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



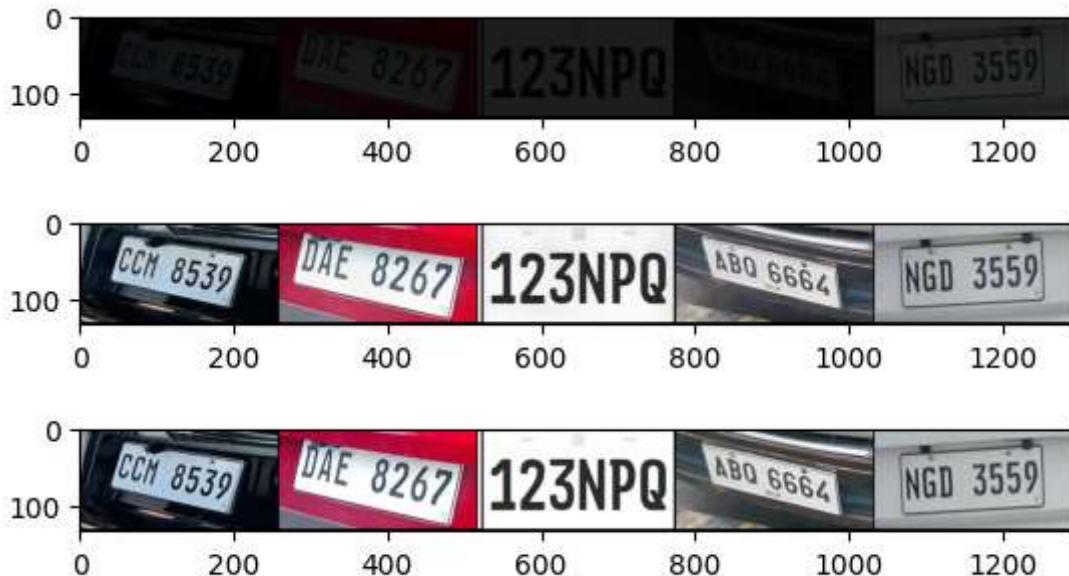
Epoch [24/200] - Mean PSNR: 19.4466, Mean SSIM: 0.8276, FID: 38.446407318115234



Epoch 00024: reducing learning rate of group 0 to 1.0000e-04.

100% |  
1250/1250 [18:51<00:00, 1.10it/s]

Training Data



Epoch [25/200], loss\_g: 6.8812, loss\_d: 0.3451, real\_score: 0.9111, fake\_score: 0.42  
12

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [25/200] - Mean PSNR: 18.6253, Mean SSIM: 0.8169, FID: 40.5386848449707



100% |  
1250/1250 [18:52<00:00, 1.10it/s]

### Training Data



Epoch [26/200], loss\_g: 7.4148, loss\_d: 0.5868, real\_score: 0.6321, fake\_score: 0.4575

Current LRs - Generator: 0.000200, Discriminator: 0.000100

### Validation Data



Epoch [26/200] - Mean PSNR: 19.1911, Mean SSIM: 0.8207, FID: 39.65391159057617

### Raw Test Set

### Generated Test Set



100% | 1250/1250 [18:52<00:00, 1.10it/s]

### Training Data





Epoch [27/200], loss\_g: 7.5348, loss\_d: 0.5608, real\_score: 0.8693, fake\_score: 0.5935

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [27/200] - Mean PSNR: 19.5470, Mean SSIM: 0.8248, FID: 40.53878402709961

Raw Test Set

Generated Test Set



100%

1250/1250 [18:52<00:00, 1.10it/s]

Training Data





Epoch [28/200], loss\_g: 8.2881, loss\_d: 0.3993, real\_score: 0.7811, fake\_score: 0.3907

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [28/200] - Mean PSNR: 19.2068, Mean SSIM: 0.8188, FID: 37.874725341796875

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:51<00:00, 1.10it/s]

Training Data



Epoch [29/200], loss\_g: 6.1950, loss\_d: 0.3684, real\_score: 0.6133, fake\_score: 0.18  
16

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [29/200] - Mean PSNR: 19.7472, Mean SSIM: 0.8326, FID: 38.3443489074707

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [30/200], loss\_g: 8.0763, loss\_d: 0.4904, real\_score: 0.6919, fake\_score: 0.41  
79

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [30/200] - Mean PSNR: 18.9249, Mean SSIM: 0.8193, FID: 37.139339447021484

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:51<00:00, 1.10it/s]  
Training Data



Epoch [31/200], loss\_g: 6.3003, loss\_d: 0.3672, real\_score: 0.7716, fake\_score: 0.3429

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



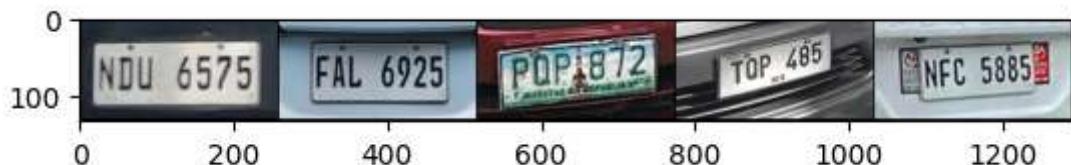


Epoch [31/200] - Mean PSNR: 19.5960, Mean SSIM: 0.8316, FID: 37.91668701171875  
 Raw Test Set                                  Generated Test Set



100% |  
 1250/1250 [18:51<00:00, 1.10it/s]

Training Data



Epoch [32/200], loss\_g: 7.3745, loss\_d: 0.7945, real\_score: 0.2545, fake\_score: 0.09  
 33

Current LRs - Generator: 0.000200, Discriminator: 0.000100

Validation Data





Epoch [32/200] - Mean PSNR: 18.4271, Mean SSIM: 0.8119, FID: 37.850799560546875

Raw Test Set

Generated Test Set



100%

1250/1250 [18:50<00:00, 1.11it/s]

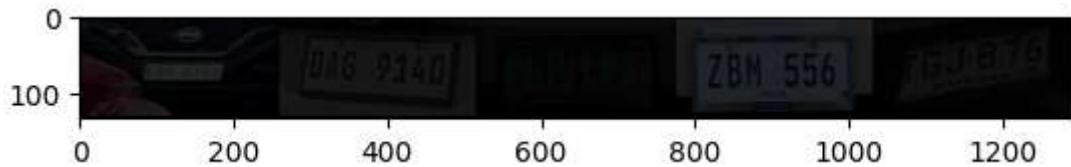
Training Data



Epoch [33/200], loss\_g: 5.9734, loss\_d: 0.5630, real\_score: 0.4399, fake\_score: 0.18  
64

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [33/200] - Mean PSNR: 19.3256, Mean SSIM: 0.8245, FID: 36.93244552612305



100% | 1250/1250 [18:49<00:00, 1.11it/s]  
Training Data



Epoch [34/200], loss\_g: 6.5473, loss\_d: 0.3547, real\_score: 0.7066, fake\_score: 0.26  
22

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [34/200] - Mean PSNR: 19.2308, Mean SSIM: 0.8264, FID: 36.49135208129883



100% | 1250/1250 [18:49<00:00, 1.11it/s]  
Training Data



Epoch [35/200], loss\_g: 7.0147, loss\_d: 0.4041, real\_score: 0.7281, fake\_score: 0.3457

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [35/200] - Mean PSNR: 19.2603, Mean SSIM: 0.8260, FID: 35.86093521118164

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:51<00:00, 1.10it/s]

Training Data





Epoch [36/200], loss\_g: 6.3441, loss\_d: 0.2667, real\_score: 0.6959, fake\_score: 0.1315

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [36/200] - Mean PSNR: 19.6823, Mean SSIM: 0.8325, FID: 36.370670318603516

Raw Test Set

Generated Test Set



100%

1250/1250 [18:52<00:00, 1.10it/s]

Training Data





Epoch [37/200], loss\_g: 5.7839, loss\_d: 0.4171, real\_score: 0.8979, fake\_score: 0.4830

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [37/200] - Mean PSNR: 19.3616, Mean SSIM: 0.8301, FID: 35.28145217895508

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:51<00:00, 1.10it/s]

Training Data



Epoch [38/200], loss\_g: 5.5461, loss\_d: 0.3104, real\_score: 0.8567, fake\_score: 0.34  
53

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [38/200] - Mean PSNR: 19.6853, Mean SSIM: 0.8278, FID: 36.04048156738281

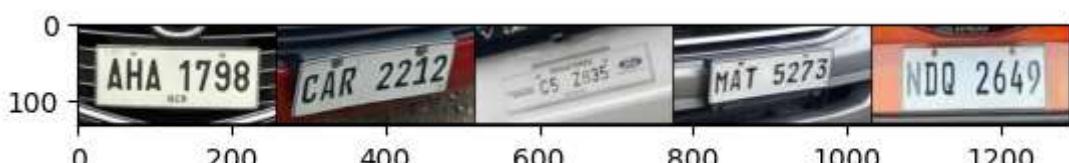
Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:51<00:00, 1.10it/s]

Training Data



Epoch [39/200], loss\_g: 5.6925, loss\_d: 0.5041, real\_score: 0.9068, fake\_score: 0.56  
11

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [39/200] - Mean PSNR: 19.6486, Mean SSIM: 0.8346, FID: 37.27739334106445

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:43<00:00, 1.06it/s]

Training Data



Epoch [40/200], loss\_g: 5.1701, loss\_d: 0.4868, real\_score: 0.6584, fake\_score: 0.3881

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data





100% |  
1250/1250 [18:53<00:00, 1.10it/s]

Training Data



Epoch [41/200], loss\_g: 5.1171, loss\_d: 0.6099, real\_score: 0.7514, fake\_score: 0.5397

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data





Epoch [41/200] - Mean PSNR: 18.8694, Mean SSIM: 0.8104, FID: 41.353050231933594

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [42/200], loss\_g: 5.8439, loss\_d: 0.4059, real\_score: 0.5288, fake\_score: 0.1113

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [42/200] - Mean PSNR: 19.6525, Mean SSIM: 0.8296, FID: 35.89630889892578



100% |  
1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [43/200], loss\_g: 6.4200, loss\_d: 0.6108, real\_score: 0.7863, fake\_score: 0.58  
13

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [43/200] - Mean PSNR: 19.7165, Mean SSIM: 0.8333, FID: 35.90210723876953



100% |  
1250/1250 [18:53<00:00, 1.10it/s]

Training Data



Epoch [44/200], loss\_g: 6.4869, loss\_d: 0.5800, real\_score: 0.4240, fake\_score: 0.1801

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [44/200] - Mean PSNR: 19.3867, Mean SSIM: 0.8283, FID: 35.04948806762695

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:11<00:00, 1.09it/s]

Training Data





Epoch [45/200], loss\_g: 8.3734, loss\_d: 0.3464, real\_score: 0.9302, fake\_score: 0.4298

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [45/200] - Mean PSNR: 19.8657, Mean SSIM: 0.8340, FID: 35.597843170166016

Raw Test Set

Generated Test Set



100%

1250/1250 [18:53<00:00, 1.10it/s]

Training Data





Epoch [46/200], loss\_g: 7.0787, loss\_d: 0.4467, real\_score: 0.6669, fake\_score: 0.34  
33

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [46/200] - Mean PSNR: 19.8077, Mean SSIM: 0.8324, FID: 36.3724250793457

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:55<00:00, 1.10it/s]

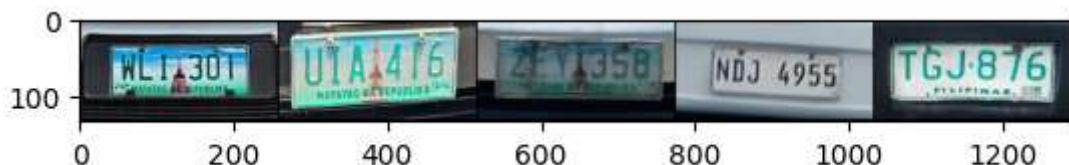
Training Data



Epoch [47/200], loss\_g: 4.9231, loss\_d: 0.7286, real\_score: 0.3368, fake\_score: 0.20  
44

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [47/200] - Mean PSNR: 19.9029, Mean SSIM: 0.8369, FID: 35.2142219543457

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [48/200], loss\_g: 4.8662, loss\_d: 0.4509, real\_score: 0.6252, fake\_score: 0.28  
98

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [48/200] - Mean PSNR: 19.1885, Mean SSIM: 0.8145, FID: 39.32527160644531

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:52<00:00, 1.10it/s]  
Training Data



Epoch [49/200], loss\_g: 4.9082, loss\_d: 0.4044, real\_score: 0.7409, fake\_score: 0.3698

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data





Epoch [49/200] - Mean PSNR: 19.6040, Mean SSIM: 0.8319, FID: 34.38228225708008  
Raw Test Set Generated Test Set



Epoch 00049: reducing learning rate of group 0 to 5.0000e-05.

1250/1250 [18:52<00:00, 1.10it/s]

## Training Data



Epoch [50/200], loss\_g: 5.6646, loss\_d: 0.4466, real\_score: 0.5714, fake\_score: 0.2401

Current LRs = Generator: 0.000200, Discriminator: 0.000050

## Validation Data





Epoch [50/200] - Mean PSNR: 19.4780, Mean SSIM: 0.8273, FID: 35.92219543457031

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [51/200], loss\_g: 5.9107, loss\_d: 0.4719, real\_score: 0.5146, fake\_score: 0.19  
00

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [51/200] - Mean PSNR: 19.3230, Mean SSIM: 0.8229, FID: 39.22008514404297



100% |  
1250/1250 [18:54<00:00, 1.10it/s]

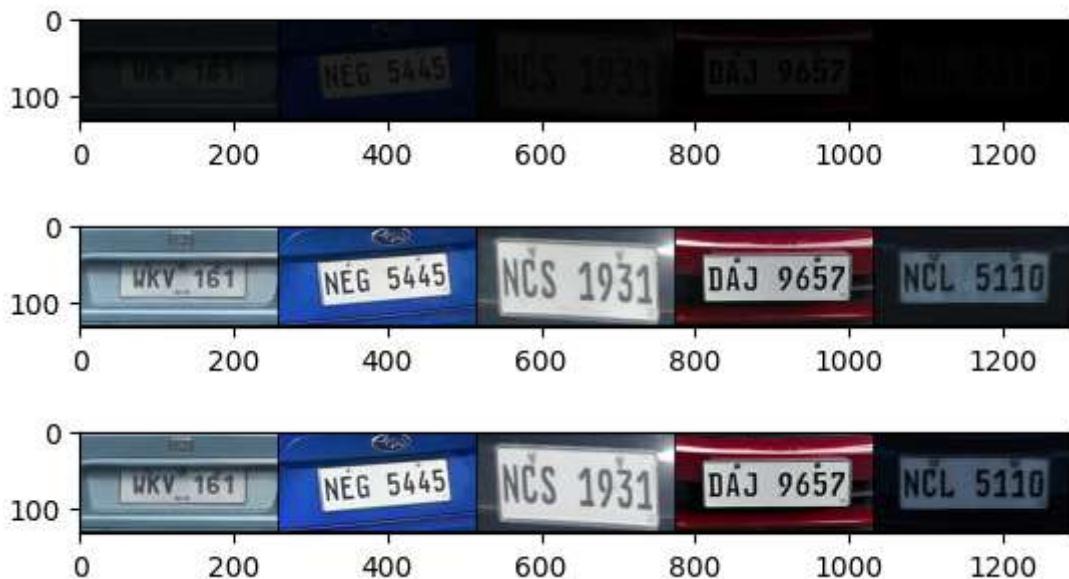
Training Data



Epoch [52/200], loss\_g: 5.7576, loss\_d: 0.4770, real\_score: 0.7740, fake\_score: 0.4648

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [52/200] - Mean PSNR: 19.5790, Mean SSIM: 0.8318, FID: 35.6588134765625



100% |  
1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [53/200], loss\_g: 5.3253, loss\_d: 0.5075, real\_score: 0.5630, fake\_score: 0.3145

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [53/200] - Mean PSNR: 19.1471, Mean SSIM: 0.8299, FID: 33.63098907470703

Raw Test Set

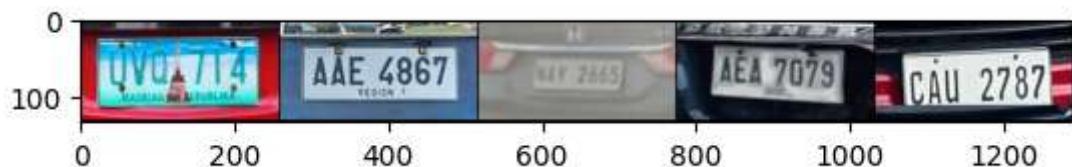
Generated Test Set



100% |  
1250/1250 [18:54<00:00, 1.10it/s]

Training Data





Epoch [54/200], loss\_g: 5.8746, loss\_d: 0.6818, real\_score: 0.3782, fake\_score: 0.2284

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [54/200] - Mean PSNR: 18.9682, Mean SSIM: 0.8186, FID: 40.606109619140625

Raw Test Set

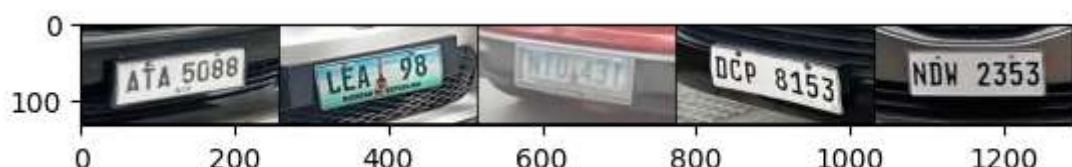
Generated Test Set



100%

1250/1250 [18:55<00:00, 1.10it/s]

Training Data





Epoch [55/200], loss\_g: 5.4337, loss\_d: 0.2782, real\_score: 0.6853, fake\_score: 0.13  
66

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [55/200] - Mean PSNR: 19.5003, Mean SSIM: 0.8295, FID: 36.19438552856445

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [56/200], loss\_g: 5.4361, loss\_d: 0.3521, real\_score: 0.7003, fake\_score: 0.26  
34

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [56/200] - Mean PSNR: 19.9919, Mean SSIM: 0.8382, FID: 33.99836730957031

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [57/200], loss\_g: 3.9459, loss\_d: 0.6200, real\_score: 0.6484, fake\_score: 0.50  
01

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [57/200] - Mean PSNR: 19.5011, Mean SSIM: 0.8229, FID: 37.963592529296875

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:53<00:00, 1.10it/s]  
Training Data



Epoch [58/200], loss\_g: 4.8085, loss\_d: 0.3983, real\_score: 0.6430, fake\_score: 0.2594

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data





Epoch [58/200] - Mean PSNR: 19.8338, Mean SSIM: 0.8371, FID: 33.86127853393555  
Raw Test Set Generated Test Set



100% |████████████████████████████████████████████████████████████████████████|  
1250/1250 [18:55<00:00, 1.10it/s]  
Termination Date:

## Training Data



Epoch [59/200], loss\_g: 5.5164, loss\_d: 0.4907, real\_score: 0.5346, fake\_score: 0.2369

Current LRs – Generator: 0.000200, Discriminator: 0.000050

#### Validation Data





Epoch [59/200] - Mean PSNR: 19.6117, Mean SSIM: 0.8320, FID: 36.521175384521484

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [60/200], loss\_g: 6.4611, loss\_d: 0.6633, real\_score: 0.9854, fake\_score: 0.6857

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [60/200] - Mean PSNR: 18.7329, Mean SSIM: 0.8182, FID: 37.373294830322266



100% |  
1250/1250 [18:54<00:00, 1.10it/s]  
Training Data



Epoch [61/200], loss\_g: 5.5181, loss\_d: 0.3960, real\_score: 0.7753, fake\_score: 0.38  
19

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [61/200] - Mean PSNR: 19.4703, Mean SSIM: 0.8283, FID: 34.49689483642578



100% |  
1250/1250 [18:55<00:00, 1.10it/s]  
Training Data



Epoch [62/200], loss\_g: 7.1770, loss\_d: 0.3835, real\_score: 0.5847, fake\_score: 0.1590

Current LRs – Generator: 0.000200, Discriminator: 0.000050

Validation Data



Epoch [62/200] - Mean PSNR: 19.6412, Mean SSIM: 0.8327, FID: 33.97242736816406

Raw Test Set

Generated Test Set



Epoch 00062: reducing learning rate of group 0 to 2.5000e-05.

100% |

1250/1250 [18:55<00:00, 1.10it/s]

Training Data

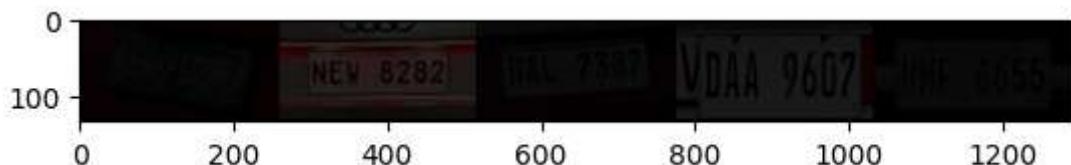




Epoch [63/200], loss\_g: 4.7057, loss\_d: 0.5082, real\_score: 0.5016, fake\_score: 0.2238

Current LRs – Generator: 0.000200, Discriminator: 0.000025

Validation Data



Epoch [63/200] - Mean PSNR: 20.0471, Mean SSIM: 0.8383, FID: 33.27448654174805

Raw Test Set

Generated Test Set



100%

1250/1250 [18:55<00:00, 1.10it/s]

Training Data





Epoch [64/200], loss\_g: 4.4399, loss\_d: 0.5381, real\_score: 0.5633, fake\_score: 0.3456

Current LRs – Generator: 0.000200, Discriminator: 0.000025

Validation Data



Epoch [64/200] - Mean PSNR: 19.8891, Mean SSIM: 0.8381, FID: 32.707763671875

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [65/200], loss\_g: 4.9765, loss\_d: 0.3520, real\_score: 0.8566, fake\_score: 0.39  
65

Current LRs – Generator: 0.000200, Discriminator: 0.000025

Validation Data



Epoch [65/200] - Mean PSNR: 19.6465, Mean SSIM: 0.8330, FID: 33.009307861328125

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [66/200], loss\_g: 4.0273, loss\_d: 0.5859, real\_score: 0.5196, fake\_score: 0.33  
89

Current LRs – Generator: 0.000200, Discriminator: 0.000025

Validation Data



Epoch [66/200] - Mean PSNR: 19.9328, Mean SSIM: 0.8391, FID: 32.64397048950195

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:55<00:00, 1.10it/s]

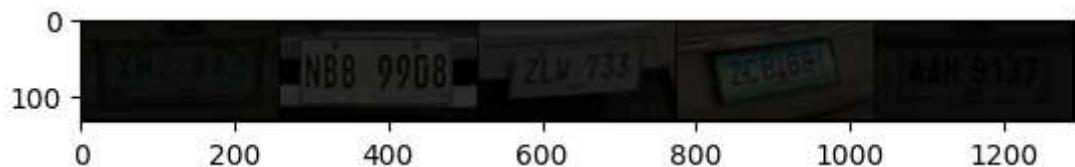
Training Data



Epoch [67/200], loss\_g: 5.0199, loss\_d: 0.3486, real\_score: 0.6669, fake\_score: 0.2154

Current LRs – Generator: 0.000200, Discriminator: 0.000025

Validation Data





Epoch [67/200] - Mean PSNR: 19.6101, Mean SSIM: 0.8333, FID: 34.50214767456055  
Raw Test Set Generated Test Set



100% |  
1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [68/200], loss\_g: 4.4843, loss\_d: 0.5318, real\_score: 0.7383, fake\_score: 0.4897

Current LRs – Generator: 0.000200, Discriminator: 0.000025

Validation Data





Epoch [68/200] - Mean PSNR: 19.5516, Mean SSIM: 0.8278, FID: 36.05794143676758

Raw Test Set

Generated Test Set



Epoch 00068: reducing learning rate of group 0 to 1.0000e-04.

100% |  
1250/1250 [18:53<00:00, 1.10it/s]  
Training Data



Epoch [69/200], loss\_g: 5.2764, loss\_d: 0.7444, real\_score: 0.6191, fake\_score: 0.5829

Current LRs – Generator: 0.000100, Discriminator: 0.000025

Validation Data



Epoch [69/200] - Mean PSNR: 19.9033, Mean SSIM: 0.8383, FID: 32.36399841308594

Raw Test Set

Generated Test Set



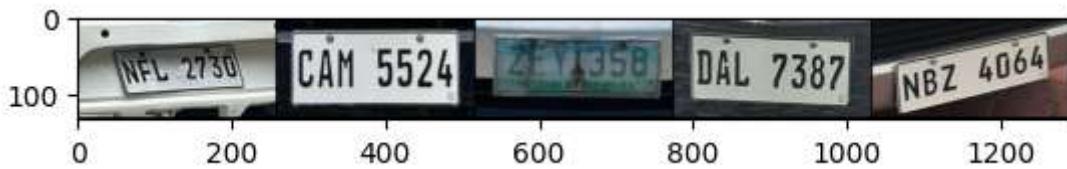
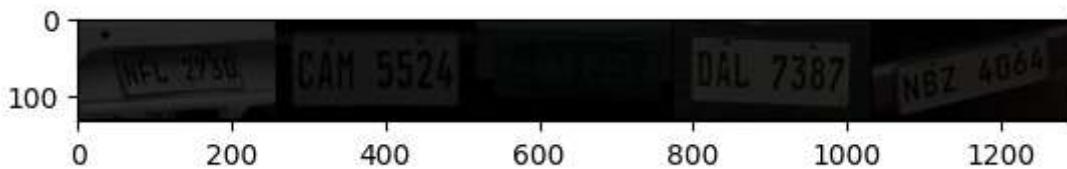
100% |  
1250/1250 [18:53<00:00, 1.10it/s]  
Training Data



Epoch [70/200], loss\_g: 5.2822, loss\_d: 0.3369, real\_score: 0.6463, fake\_score: 0.1805

Current LRs – Generator: 0.000100, Discriminator: 0.000025

Validation Data



Epoch [70/200] - Mean PSNR: 19.9513, Mean SSIM: 0.8362, FID: 33.87950134277344

Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:53<00:00, 1.10it/s]

### Training Data



Epoch [71/200], loss\_g: 4.5609, loss\_d: 0.4064, real\_score: 0.6947, fake\_score: 0.3235

Current LRs - Generator: 0.000100, Discriminator: 0.000025

### Validation Data



Epoch [71/200] - Mean PSNR: 20.0985, Mean SSIM: 0.8401, FID: 31.13563346862793

### Raw Test Set

### Generated Test Set



100% |  
1250/1250 [18:53<00:00, 1.10it/s]

### Training Data

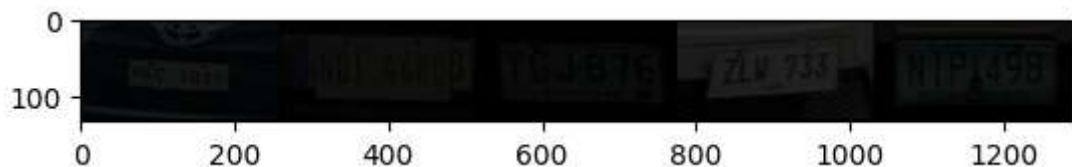




Epoch [72/200], loss\_g: 4.3425, loss\_d: 0.4502, real\_score: 0.6470, fake\_score: 0.3281

Current LRs – Generator: 0.000100, Discriminator: 0.000025

Validation Data



Epoch [72/200] - Mean PSNR: 19.9807, Mean SSIM: 0.8428, FID: 32.067867279052734

Raw Test Set

Generated Test Set



100%

1250/1250 [18:54<00:00, 1.10it/s]

Training Data





Current LRs – Generator: 0.000100, Discriminator: 0.000025

Validation Data



Epoch [73/200] - Mean PSNR: 19.6563, Mean SSIM: 0.8357, FID: 32.406192779541016

Raw Test Set

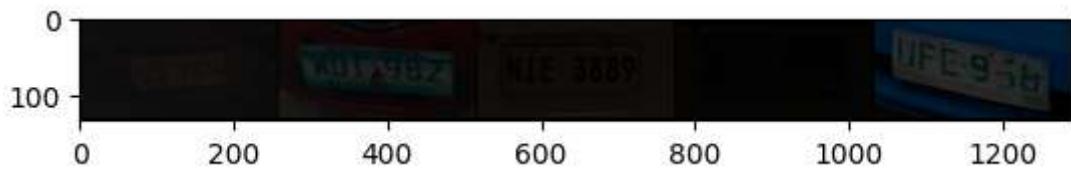
Generated Test Set



100% |

1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [74/200], loss\_g: 4.4238, loss\_d: 0.6718, real\_score: 0.7015, fake\_score: 0.58  
08

Current LRs – Generator: 0.000100, Discriminator: 0.000025

Validation Data



Epoch [74/200] - Mean PSNR: 19.6914, Mean SSIM: 0.8361, FID: 33.11274719238281

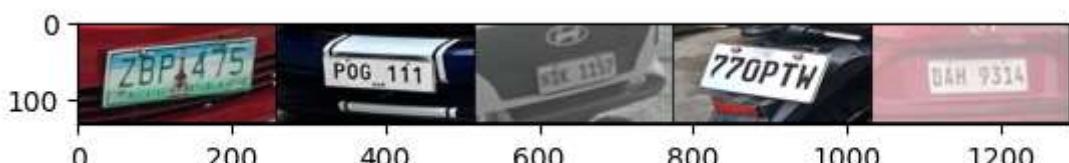
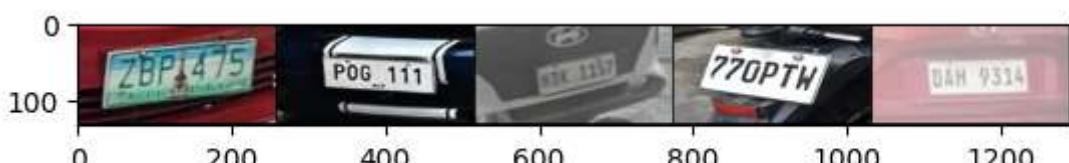
Raw Test Set

Generated Test Set



100% |  
1250/1250 [18:53<00:00, 1.10it/s]

Training Data



Epoch [75/200], loss\_g: 4.3721, loss\_d: 0.6064, real\_score: 0.4503, fake\_score: 0.27  
76

Current LRs – Generator: 0.000100, Discriminator: 0.000025

Validation Data



Epoch [75/200] - Mean PSNR: 19.7120, Mean SSIM: 0.8353, FID: 34.76650619506836

Raw Test Set

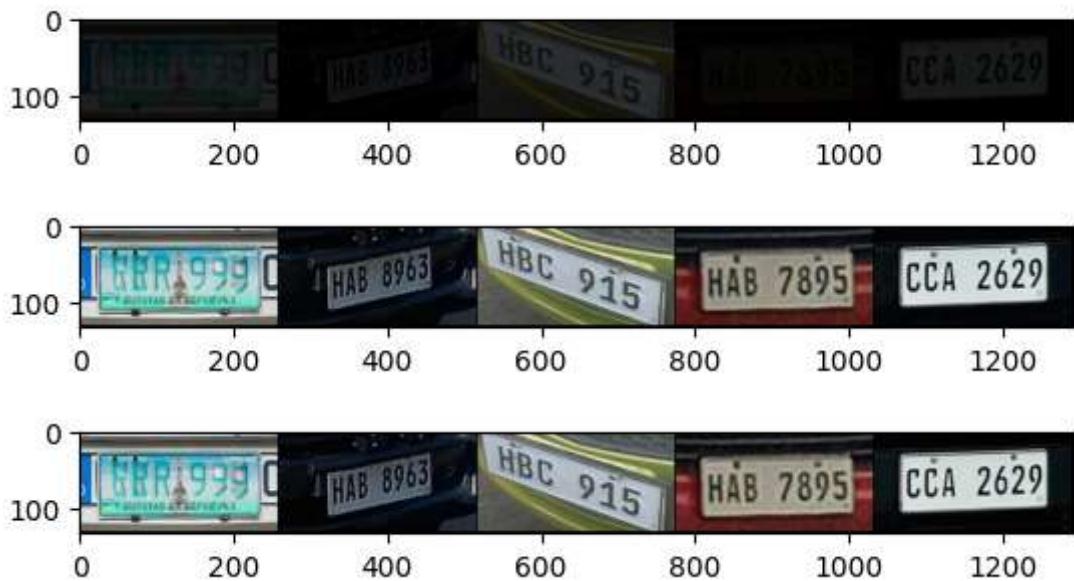
Generated Test Set



Epoch 00075: reducing learning rate of group 0 to 1.2500e-05.

100% |  
1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [76/200], loss\_g: 5.2036, loss\_d: 0.6257, real\_score: 0.4503, fake\_score: 0.2957

Current LRs – Generator: 0.000100, Discriminator: 0.000013

Validation Data





Epoch [76/200] - Mean PSNR: 19.8029, Mean SSIM: 0.8357, FID: 32.24737548828125  
Raw Test Set Generated Test Set



100% |████████████████████████████████████████████████████████████████████████|  
1250/1250 [18:54<00:00, 1.10it/s]  
Transfert : 0.0%

## Training Data



Epoch [77/200], loss\_g: 3.9547, loss\_d: 0.6082, real\_score: 0.5193, fake\_score: 0.3750

Current LRs = Generator: 0.000100, Discriminator: 0.000013

### Validation Data





Epoch [77/200] - Mean PSNR: 20.3417, Mean SSIM: 0.8443, FID: 32.42433166503906

Raw Test Set

Generated Test Set



100% |  
1250/1250 [19:12<00:00, 1.08it/s]

Training Data



Epoch [78/200], loss\_g: 3.8956, loss\_d: 0.4747, real\_score: 0.6029, fake\_score: 0.31  
49

Current LRs – Generator: 0.000100, Discriminator: 0.000013

Validation Data



Epoch [78/200] - Mean PSNR: 20.1911, Mean SSIM: 0.8436, FID: 31.16459846496582



100% |  
1250/1250 [18:56<00:00, 1.10it/s]  
Training Data



Epoch [79/200], loss\_g: 4.6817, loss\_d: 0.3201, real\_score: 0.7232, fake\_score: 0.24  
55

Current LRs – Generator: 0.000100, Discriminator: 0.000013

Validation Data



Epoch [79/200] - Mean PSNR: 19.8689, Mean SSIM: 0.8388, FID: 32.296390533447266



100% |  
1250/1250 [19:43<00:00, 1.06it/s]  
Training Data



Epoch [80/200], loss\_g: 4.3185, loss\_d: 0.5358, real\_score: 0.6889, fake\_score: 0.46  
16

Current LRs – Generator: 0.000100, Discriminator: 0.000013

Validation Data



Epoch [80/200] - Mean PSNR: 19.7793, Mean SSIM: 0.8401, FID: 31.546554565429688

Raw Test Set

Generated Test Set



0%||  
| 3/1250 [00:11<1:18:03, 3.76s/it]

```
-----  
KeyboardInterrupt  
File <timed exec>:1  
  
Cell In[28], line 91, in fit(generator, discriminator, epochs, lr, adv_lambda, l1_lambda, ssim_lambda, perceptual_lambda, text_lambda, g_opt, d_opt, checkpoint)  
    88         d_loss, real_score, fake_score = train_discriminator(discriminator, generator, inputs, targets, d_opt)  
    89         # Train generator  
---> 90         g_loss, fake_images, g_loss_trace = train_generator(generator, discriminator, inputs, targets, g_opt, adv_lambda,  
    91                                         l1_lambda, ssim_lambda, perceptual_lambda, text_lambda)  
    92  
    93     print('Training Data')  
    94     print_images(inputs, 5)  
  
Cell In[25], line 54, in train_generator(generator, discriminator, inputs, targets, g_opt, adv_lambda, l1_lambda, ssim_lambda, perceptual_lambda, text_lambda)  
    50     g_loss.backward()  
    51     g_opt.step()  
    52     return g_loss.item(), fake_images, {  
---> 53         "total_loss": g_loss.item(),  
    54         "adv": adv_loss.item(),  
    55         "l1": l1_loss.item(),  
    56         "ssim": ssim_loss.item(),  
    57         "perceptual": perceptual_loss.item(),  
    58         "text": text_loss.item()  
    59     }  
    60 }  
  
KeyboardInterrupt:
```

## Post Training

```
In [34]: torch.cuda.empty_cache()
```

## Save Model

```
In [34]: torch.save(generator.state_dict(), "[RAU-NET]generator-Datasetv2-5epoch.pt")  
torch.save(discriminator.state_dict(), "[RAU-NET]discriminator-Datasetv2-5epoch.pt")
```

```
In [35]: torch.save({  
    'generator_state_dict': generator.state_dict(),  
    'discriminator_state_dict': discriminator.state_dict(),  
    'opt_g_state_dict': g_opt.state_dict(),  
    'opt_d_state_dict': d_opt.state_dict(),  
    'history': history  
}, f'[RAU-NET]Checkpoint-Datasetv2-30epoch.pt')
```

## Model Evaluation

```
In [ ]: psnr_vals = []
ssim_vals = []

fid_metric = FrechetInceptionDistance(feature=2048).to(device)
fid_metric.reset()

generator.eval()
for val_input, val_target in valid_dl:
    val_input = val_input.to(device)
    val_target = val_target.to(device)

    with torch.no_grad():
        gen_output = generator(val_input)

    # Convert to uint8 for FID
    gen_uint8 = to_uint8(gen_output)
    target_uint8 = to_uint8(val_target)

    fid_metric.update(gen_uint8, real=False)
    fid_metric.update(target_uint8, real=True)

    show_generated_images(gen_output)

    # Compute PSNR/SSIM
    psnr = psnr_metric(gen_output, val_target).item()
    ssim = ssim_metric(gen_output, val_target).item()
    psnr_vals.append(psnr)
    ssim_vals.append(ssim)

print("Mean PSNR:", sum(psnr_vals) / len(psnr_vals))
print("Mean SSIM:", sum(ssim_vals) / len(ssim_vals))
print("FID:", fid_metric.compute().item())
```











```
In [32]: def deblur_image(image_path, generator, device):
    image = Image.open(image_path).convert('RGB')
    input_tensor = valid_transform(image).unsqueeze(0).to(device)

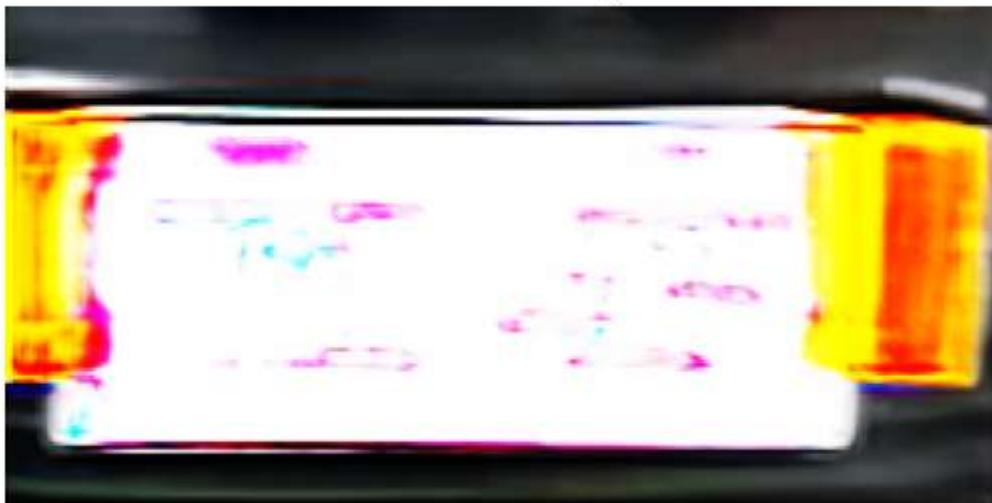
    generator.eval()
    with torch.no_grad():
        gen_image = generator(input_tensor)

    output = denorm(gen_image.squeeze(0).cpu())
```

```
plt.imshow(output.permute(1, 2, 0)) # CHW → HWC
plt.axis('off')
plt.title("Generated Image")
plt.show()
```

In [31]: deblur\_image(os.path.join(DATA\_DIR, 'h.blur', 'v2', 'test', 'h.blur', '129.jpg'), g

Generated Image



```
raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl, device)

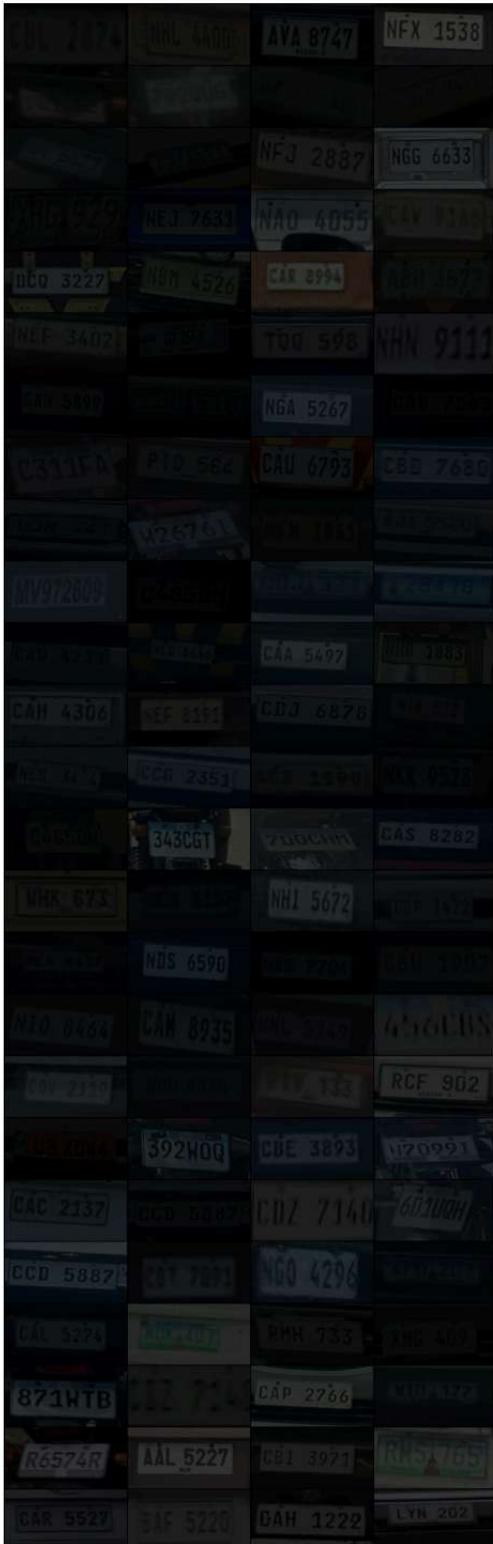
fig, axes = plt.subplots(1, 2, figsize=(15, 15)) # 1 row, 2 columns

# Raw Test Set
axes[0].imshow(raw_grid_img)
axes[0].axis("off")
axes[0].set_title("Raw Test Set")

# Generated Test Set
axes[1].imshow(gen_grid_img)
axes[1].axis("off")
axes[1].set_title("Generated Test Set")

plt.tight_layout()
plt.show()
```

Raw Test Set



Generated Test Set



```
In [33]: torch.cuda.empty_cache()
```

## Load Model

```
In [31]: MODELS_DIR = os.path.join('..', '..', '..', 'models', 'h_blur')
```

```
In [31]: generator = AttentionResUNetGenerator()
loaded_g_wts = torch.load('[RAU-NET OCRDET LR]generator-Datasetv2-80 epoch.pt')
generator.load_state_dict(loaded_g_wts)
generator = to_device(generator, device)

discriminator = PatchDiscriminator()
loaded_d_wts = torch.load('[RAU-NET OCRDET LR]discriminator-Datasetv2-80 epoch.pt')
discriminator.load_state_dict(loaded_d_wts)
discriminator = to_device(discriminator, device)
```

```
In [29]: checkpoint = torch.load('[RAU-NET OCRDET LR]Checkpoint-Logs-Datasetv2-80 epoch.pt')
#g_opt.load_state_dict(checkpoint['opt_g_state_dict'])
#d_opt.load_state_dict(checkpoint['opt_d_state_dict'])
#history = checkpoint['losses_g']
```

```
In [39]: def are_models_identical(model1, model2):
    # Check if state dictionaries have the same keys
    if model1.state_dict().keys() != model2.state_dict().keys():
        return False

    # Compare each parameter's values
    for key in model1.state_dict():
        if not torch.equal(model1.state_dict()[key], model2.state_dict()[key]):
            return False
    return True
```

```
In [45]: if are_models_identical(generator1, generator2):
    print("Models have identical weights and biases")
else:
    print("Models have different weights or biases")
```

Models have different weights or biases

```
In [9]: len(history)
```

```
Out[9]: 79
```

## Save Reconstructed Test Set

```
In [39]: def save_generated_tests(generator, test_dl, sample_dir='./results'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)
```

```

        for b in range(batch_size):
            img = fake_images[b]

            file_name = f"{global_idx}.jpg"
            save_path = os.path.join(sample_dir, file_name)

            save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
            global_idx += 1

```

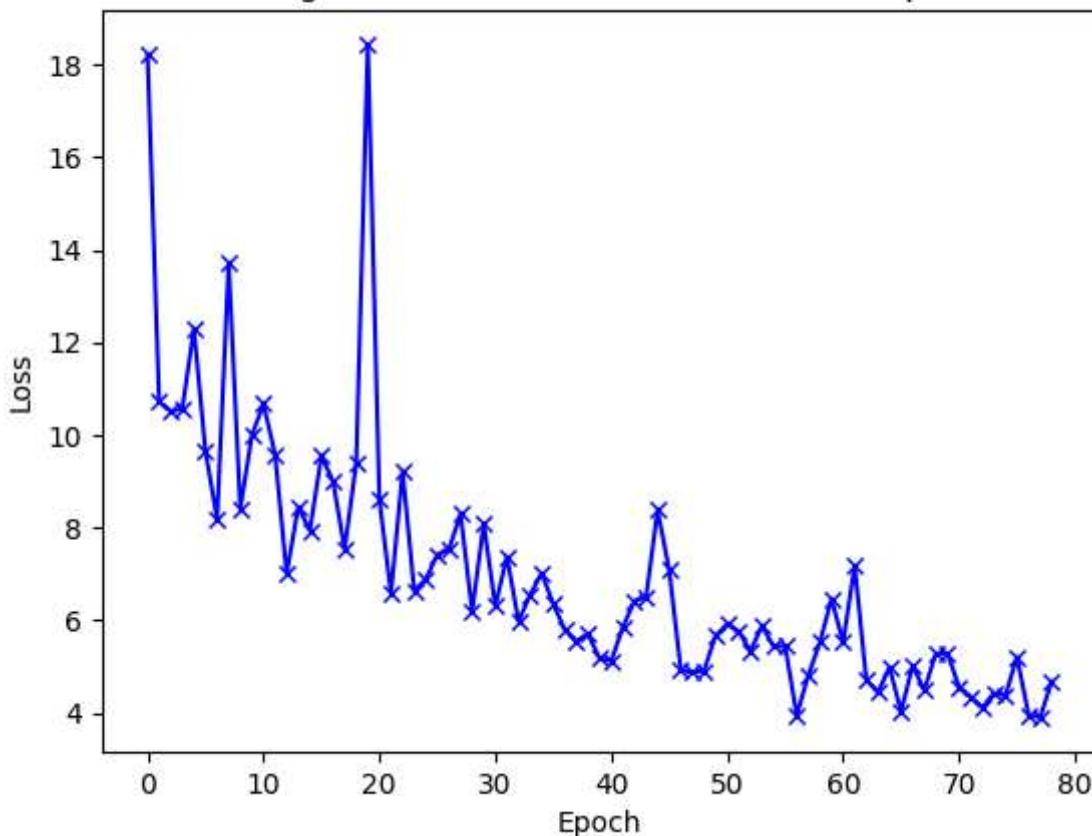
In [40]: `save_generated_tests(generator, test_dl)`

In [ ]:

In [38]: `def plot_g_losses(history):
 losses_g = [x for x in history['losses_g']]
 plt.plot(losses_g, '-bx')
 plt.xlabel('Epoch')
 plt.ylabel('Loss')
 plt.title('Low Light Blur - Generator Loss vs No. of epochs')

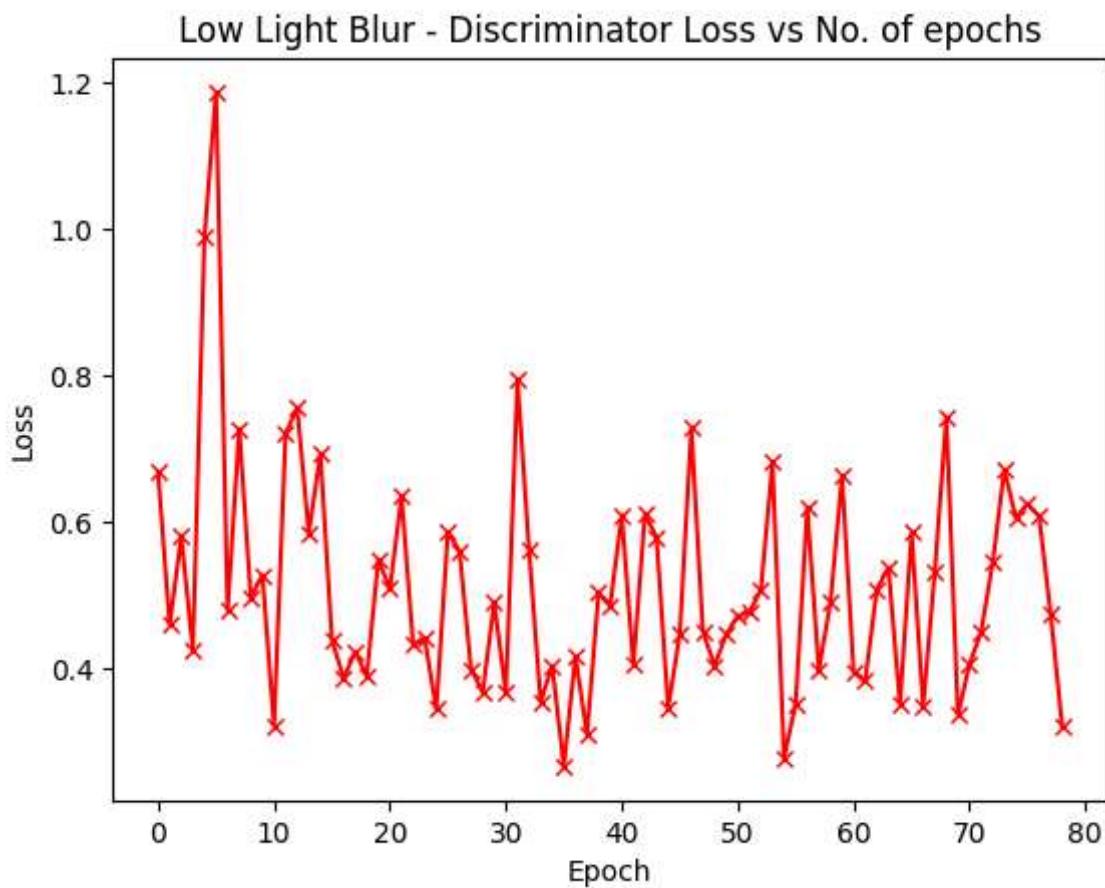
plot_g_losses(checkpoint)`

Low Light Blur - Generator Loss vs No. of epochs

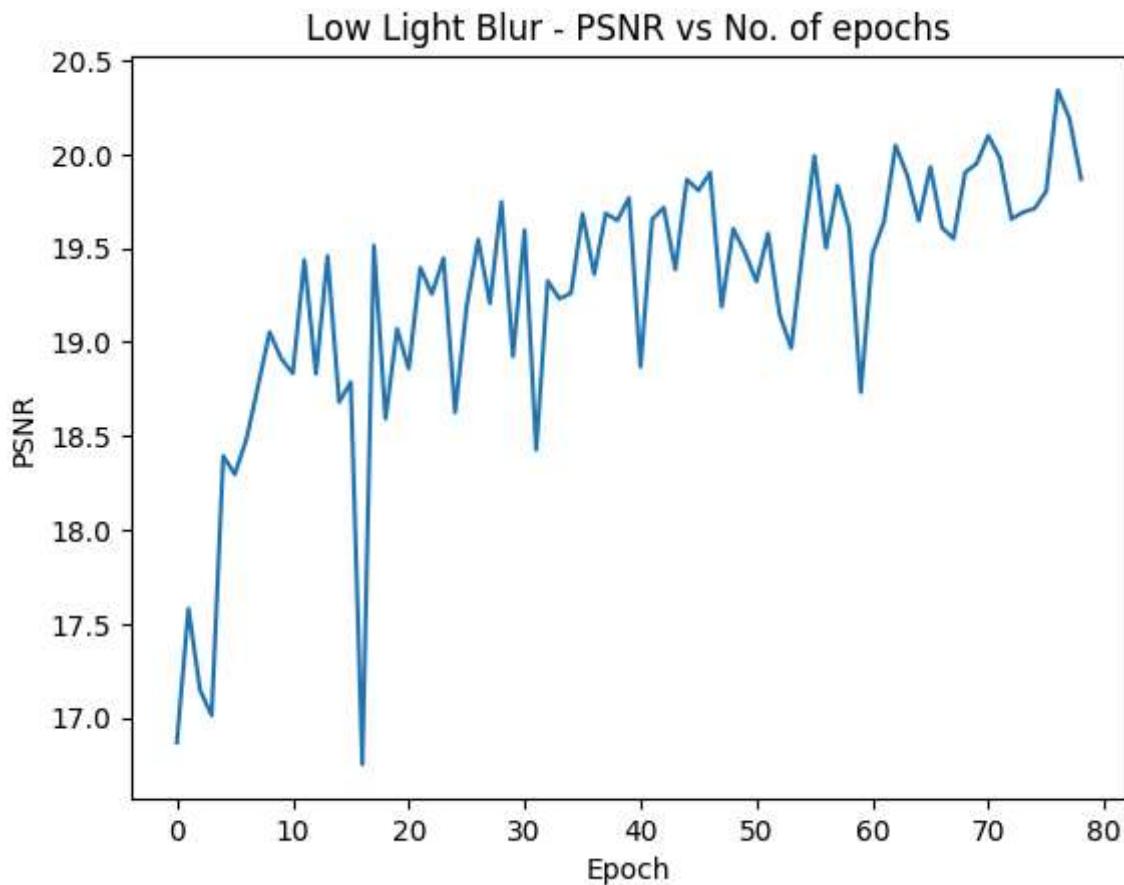


In [39]: `def plot_d_losses(history):
 losses_d = [x for x in history['losses_d']]
 plt.plot(losses_d, '-rx')
 plt.xlabel('Epoch')
 plt.ylabel('Loss')`

```
plt.title('Low Light Blur - Discriminator Loss vs No. of epochs')  
plot_d_losses(checkpoint)
```



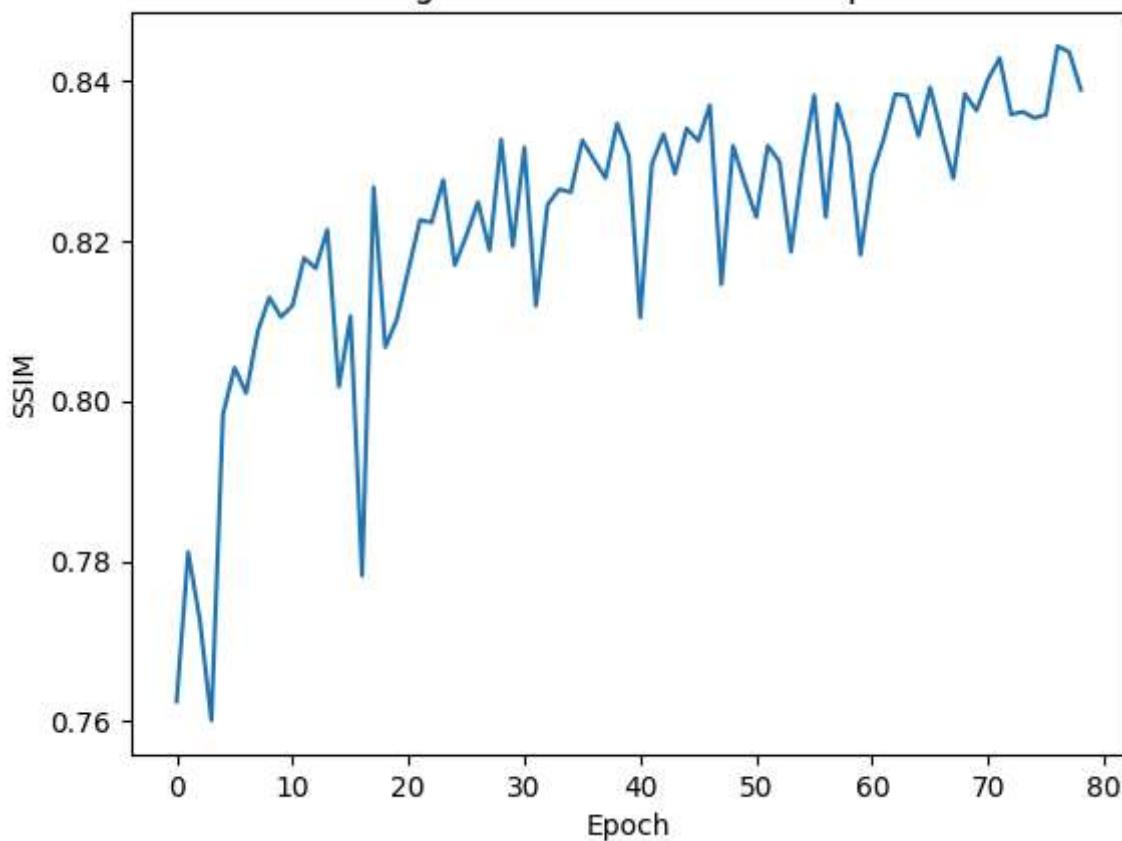
```
In [40]: def plot_psnrs(history):  
    psnrs = [x for x in history['psnrs']]  
    plt.plot(psnrs)  
    plt.xlabel('Epoch')  
    plt.ylabel('PSNR')  
    plt.title('Low Light Blur - PSNR vs No. of epochs')  
  
plot_psnrs(checkpoint)
```



```
In [41]: def plot_ssims(history):
    psnrs = [x for x in history['ssims']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('SSIM')
    plt.title('Low Light Blur - SSIM vs No. of epochs')

plot_ssims(checkpoint)
```

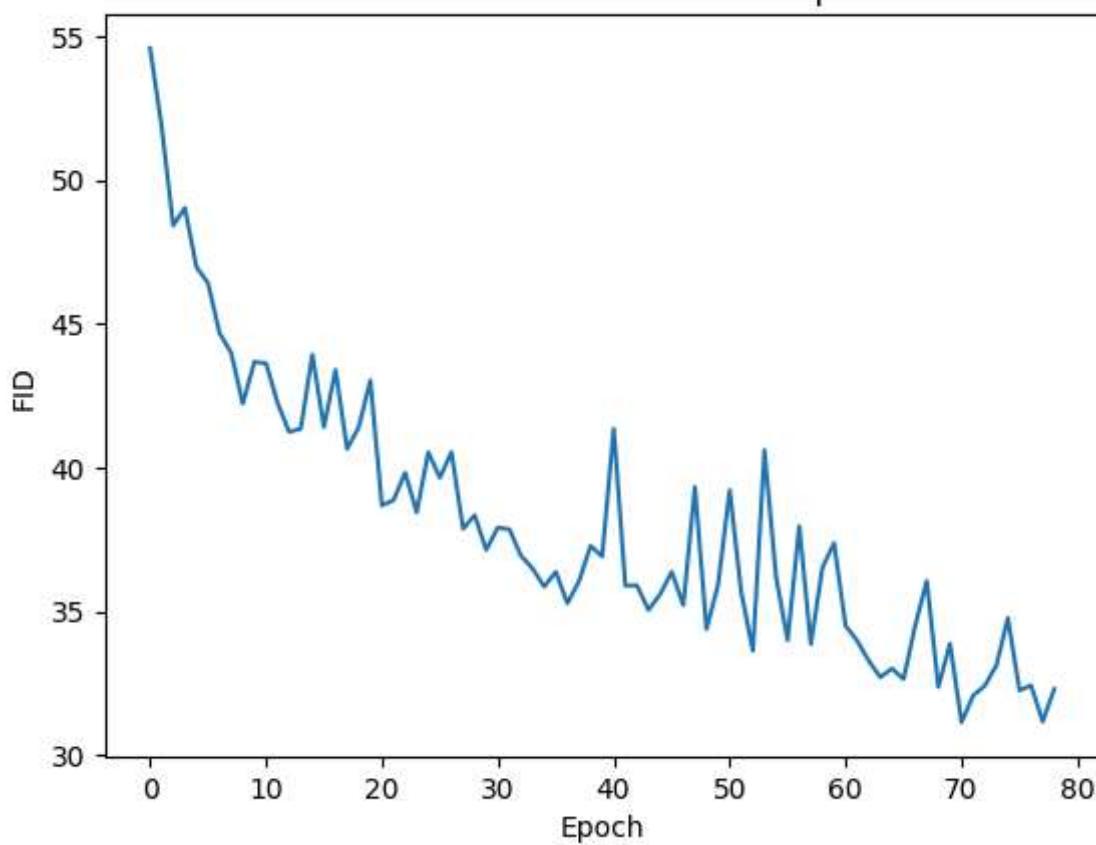
Low Light Blur - SSIM vs No. of epochs



```
In [30]: def plot_fids(history):
    psnrs = [x for x in history['fids']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('FID')
    plt.title('Horizontal Blur - FID vs No. of epochs')

plot_fids(checkpoint)
```

Horizontal Blur - FID vs No. of epochs



In [ ]: