

```
In [1]: import os
import sys
from tqdm import tqdm
import numpy as np
import torch
import torch.nn as nn
from torch import optim
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
from torchvision.utils import make_grid
import torchvision.utils as vutils
import torchvision.transforms as T
import torchvision.transforms.functional as TF
import torch.nn.functional as F
from torchvision.utils import save_image
from PIL import Image
import matplotlib.pyplot as plt
%matplotlib inline
from paddleocr import PaddleOCR, draw_ocr
import paddle
from Levenshtein import distance as levenshtein_distance
from pytorch_msssim import ssim as ssim_fn
from lpips import LPIPS
from torchmetrics.image.psnr import PeakSignalNoiseRatio
from torchmetrics.image.ssim import StructuralSimilarityIndexMeasure
from torchmetrics.image.fid import FrechetInceptionDistance
```

```
In [2]: sys.path.append(os.path.abspath('../../../../../src/dataset'))
from paired_image_dataset import PairedImageDataset

DATA_DIR = os.path.join('..', '..', '..', 'data')
print(os.listdir(DATA_DIR))

['h.blur', 'low.light', 'low.qual', 'test', 'v.blur']
```

```
In [3]: image_width = 256  
        image_height = 128  
        batch_size = 16  
        stats = (0.5, 0.5, 0.5), (0.5, 0.5, 0.5)
```

```
In [4]: train_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)]))

valid_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)
])
```

```

valid_ds = PairedImageDataset(blur_dir=os.path.join(DATA_DIR, 'low_qual', 'valid',
                                                    normal_dir=os.path.join(DATA_DIR, 'low_qual', 'valid',
                                                                transform=valid_transform))

test_ds = ImageFolder(os.path.join(DATA_DIR, 'low_qual', 'test'), transform=valid_t

```

```
In [36]: train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=4)
valid_dl = DataLoader(valid_ds, batch_size=batch_size*2, shuffle=True, num_workers=4)
test_dl = DataLoader(test_ds, batch_size=batch_size//2, shuffle=False, num_workers=4)
```

Helper functions

```

In [7]: def denorm(img_tensors):
    return img_tensors * stats[1][0] + stats[0][0]

def show_images(input_images, target_images, nmax=16):
    input_images = denorm(input_images.detach()[:nmax])
    target_images = denorm(target_images.detach()[:nmax])

    # Combine into a single tensor for visualization
    combined = torch.cat((input_images, target_images), 0)
    grid = make_grid(combined, nrow=nmax)

    plt.figure(figsize=(nmax, 4))
    plt.imshow(grid.permute(1, 2, 0))
    plt.axis("off")
    plt.title("Top: Horizontal Blur / Blurred | Bottom: Normal / Target")
    plt.show()

def show_batch(dl, nmax=16):
    for input_batch, target_batch in dl:
        show_images(input_batch, target_batch, nmax)
        break

```

```

In [8]: def show_paired_samples(dataset, num_samples=4):
    fig, axes = plt.subplots(num_samples, 2, figsize=(5, 2 * num_samples))

    for i in range(num_samples):
        input_img, target_img = dataset[i]

        input_img = denorm(input_img)
        target_img = denorm(target_img)

        if isinstance(input_img, torch.Tensor):
            input_img = input_img.permute(1, 2, 0).numpy()
            target_img = target_img.permute(1, 2, 0).numpy()

        axes[i, 0].imshow(input_img)
        axes[i, 0].set_title("Horizontal Blur (Input)")
        axes[i, 0].axis("off")

        axes[i, 1].imshow(target_img)

```

```
axes[i, 1].set_title("Normal (Target)")
axes[i, 1].axis("off")

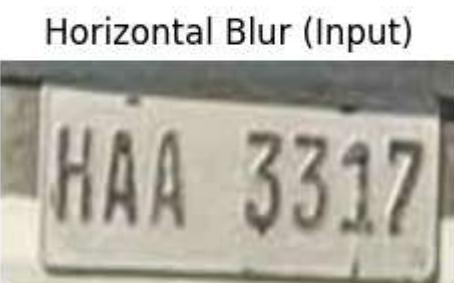
plt.tight_layout()
plt.show()

def show_generated_images(images, nrow=4):
    images = denorm(images.detach().cpu())

    # Make grid
    grid_img = vutils.make_grid(images, nrow=nrow, normalize=False)

    # Show image
    plt.figure(figsize=(nrow * 2, 2 * (len(images) // nrow)))
    plt.axis('off')
    plt.imshow(grid_img.permute(1, 2, 0))
    plt.show()
```

```
In [9]: show_paired_samples(train_ds)
```



```
In [10]: show_batch(valid_dl, nmax=14)
```

Top: Horizontal Blur / Blurred | Bottom: Normal / Target



```
In [11]: def print_images(image_tensor, num_images):
```

```
    images = denorm(image_tensor)
    images = images.detach().cpu()
    image_grid = make_grid(images[:num_images], nrow=5)
```

```
plt.imshow(image_grid.permute(1, 2, 0).squeeze())
plt.show()
```

```
In [12]: def save_samples(generator, test_dl, epoch, sample_dir='./generated_test_outputs'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)

            for b in range(batch_size):
                img = fake_images[b]

                # create folder per image
                folder_path = os.path.join(sample_dir, str(global_idx))
                os.makedirs(folder_path, exist_ok=True)

                file_name = f'{global_idx}_{epoch}epoch.jpg'
                save_path = os.path.join(folder_path, file_name)

                save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
                global_idx += 1
```

```
In [13]: def get_default_device():
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    def __init__(self, dl, device):
        self(dl = dl
        self.device = device

    def __iter__(self):
        for b in self(dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self(dl)
```

```
In [14]: device = get_default_device()
device
```

```
Out[14]: device(type='cuda')
```

```
In [15]: train_dl = DeviceDataLoader(train_dl, device)
valid_dl = DeviceDataLoader(valid_dl, device)
test_dl = DeviceDataLoader(test_dl, device)
```

Building the Model

Generator

```
In [16]: class ResConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3, dropout=0.0, batch_norm=False):
        super().__init__()
        padding = kernel_size // 2

        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size, padding=padding)
        self.bn1 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size, padding=padding)
        self.bn2 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.drop = nn.Dropout(p=dropout) if dropout > 0 else nn.Identity()

        self.shortcut = nn.Conv2d(in_channels, out_channels, kernel_size=1, padding=0)
        self.bn_sc = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

    def forward(self, x):
        residual = x

        out = self.conv1(x)
        out = self.bn1(out)
        out = F.relu(out)

        out = self.conv2(out)
        out = self.bn2(out)
        out = self.drop(out)

        shortcut = self.shortcut(residual)
        shortcut = self.bn_sc(shortcut)

        out += shortcut
        out = F.relu(out)
        return out

class GatingSignal(nn.Module):
    def __init__(self, in_channels, out_channels, batch_norm=False):
        super().__init__()
```

```

        layers = [nn.Conv2d(in_channels, out_channels, kernel_size=1, padding=0)]
        if batch_norm:
            layers.append(nn.BatchNorm2d(out_channels))
        layers.append(nn.ReLU(inplace=True))
        self.gate = nn.Sequential(*layers)

    def forward(self, x):
        return self.gate(x)

class AttentionBlock(nn.Module):
    def __init__(self, in_channels, gating_channels, inter_channels):
        super().__init__()

        self.theta_x = nn.Conv2d(in_channels, inter_channels, kernel_size=2, stride=2)
        self.phi_g = nn.Conv2d(gating_channels, inter_channels, kernel_size=1, padding=0)

        self.upsample_g = nn.ConvTranspose2d(inter_channels, inter_channels, kernel_size=2, stride=2)
        self.combine = nn.Sequential(
            nn.ReLU(inplace=True),
            nn.Conv2d(inter_channels, 1, kernel_size=1),
            nn.Sigmoid())

        self.final_conv = nn.Sequential(
            nn.Conv2d(in_channels, in_channels, kernel_size=1),
            nn.BatchNorm2d(in_channels))

    def forward(self, x, g):
        # x: skip connection feature map
        # g: gating signal (decoder feature map)
        theta_x = self.theta_x(x)
        phi_g = self.phi_g(g)

        # Upsample gating to match theta_x
        upsample_g = self.upsample_g(phi_g)

        if upsample_g.shape != theta_x.shape:
            upsample_g = F.interpolate(upsample_g, size=theta_x.shape[2:], mode='bilinear')

        psi = self.combine(theta_x + upsample_g)
        psi = F.interpolate(psi, size=x.shape[2:], mode='bilinear', align_corners=True)
        psi = psi.expand(-1, x.shape[1], -1, -1)

        y = x * psi
        return self.final_conv(y)

class AttentionResUNetGenerator(nn.Module):
    def __init__(self, in_channels=3, out_channels=3, base_filters=64, dropout=0.0,
                 super().__init__()

        # Encoder
        self.down1 = ResConvBlock(in_channels, base_filters, dropout=dropout, batch_norm=True)
        self.pool1 = nn.MaxPool2d(2)
        self.down2 = ResConvBlock(base_filters, base_filters * 2, dropout=dropout, batch_norm=True)
        self.pool2 = nn.MaxPool2d(2)
        self.down3 = ResConvBlock(base_filters * 2, base_filters * 4, dropout=dropout, batch_norm=True)
        self.pool3 = nn.MaxPool2d(2)
        self.up1 = ResConvBlock(base_filters * 4, base_filters * 2, dropout=dropout, batch_norm=True)
        self.up2 = ResConvBlock(base_filters * 2, base_filters, dropout=dropout, batch_norm=True)
        self.up3 = ResConvBlock(base_filters, base_filters, dropout=dropout, batch_norm=True)
        self.out = nn.Conv2d(base_filters, out_channels, kernel_size=1)

```

```

        self.pool3 = nn.MaxPool2d(2)
        self.down4 = ResConvBlock(base_filters * 4, base_filters * 8, dropout=dropout)
        self.pool4 = nn.MaxPool2d(2)

    # Bottleneck
    self.bottleneck = ResConvBlock(base_filters * 8, base_filters * 16, dropout=dropout)

    # Gating and Attention
    self.gate4 = GatingSignal(base_filters * 16, base_filters * 8, batch_norm)
    self.attn4 = AttentionBlock(base_filters * 8, base_filters * 8, base_filters * 16)

    self.gate3 = GatingSignal(base_filters * 8, base_filters * 4, batch_norm) #
    self.attn3 = AttentionBlock(base_filters * 4, base_filters * 4, base_filters * 8)

    self.gate2 = GatingSignal(base_filters * 4, base_filters * 2, batch_norm) #
    self.attn2 = AttentionBlock(base_filters * 2, base_filters * 2, base_filters * 4)

    self.gate1 = GatingSignal(base_filters * 2, base_filters, batch_norm) # 128
    self.attn1 = AttentionBlock(base_filters, base_filters, base_filters) # 64

    # Decoder
    self.up4 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec4 = ResConvBlock(base_filters * 16 + base_filters * 8, base_filters * 8)

    self.up3 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec3 = ResConvBlock(base_filters * 8 + base_filters * 4, base_filters * 4)

    self.up2 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec2 = ResConvBlock(base_filters * 4 + base_filters * 2, base_filters * 2)

    self.up1 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec1 = ResConvBlock(base_filters * 2 + base_filters, base_filters, dropout=dropout)

    # Output Layer
    self.final_conv = nn.Sequential(
        nn.Conv2d(base_filters, out_channels, kernel_size=1),
        nn.Tanh())

def forward(self, x):
    d1 = self.down1(x)
    p1 = self.pool1(d1)
    d2 = self.down2(p1)
    p2 = self.pool2(d2)
    d3 = self.down3(p2)
    p3 = self.pool3(d3)
    d4 = self.down4(p3)
    p4 = self.pool4(d4)

    bn = self.bottleneck(p4)

    g4 = self.gate4(bn)
    a4 = self.attn4(d4, g4)
    u4 = self.up4(bn)
    u4 = torch.cat([u4, a4], dim=1)
    d5 = self.dec4(u4)

```

```

        g3 = self.gate3(d5)
        a3 = self.attn3(d3, g3)
        u3 = self.up3(d5)
        u3 = torch.cat([u3, a3], dim=1)
        d6 = self.dec3(u3)

        g2 = self.gate2(d6)
        a2 = self.attn2(d2, g2)
        u2 = self.up2(d6)
        u2 = torch.cat([u2, a2], dim=1)
        d7 = self.dec2(u2)

        g1 = self.gate1(d7)
        a1 = self.attn1(d1, g1)
        u1 = self.up1(d7)
        u1 = torch.cat([u1, a1], dim=1)
        d8 = self.dec1(u1)

    return self.final_conv(d8)

```

In [17]:

```

generator = AttentionResUNetGenerator()
x = torch.randn(4, 3, 128, 256) # Batch of horizontal motion blurred images
y = generator(x) # Deblurred output
print(y.shape) # Should be (4, 3, 128, 256)

```

torch.Size([4, 3, 128, 256])

Discriminator

In [18]:

```

class PatchDiscriminator(nn.Module):
    def __init__(self, in_channels=3, base_filters=64):
        super().__init__()
        layers = [
            nn.Conv2d(in_channels, base_filters, kernel_size=4, stride=2, padding=1,
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters, base_filters * 2, kernel_size=4, stride=2, padding=1,
                     nn.BatchNorm2d(base_filters * 2),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 2, base_filters * 4, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 4),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 4, base_filters * 8, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 8),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 8, 1, kernel_size=4, stride=2, padding=1),
        ]
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)

```

```
In [19]: image1 = torch.rand((1, 3, 128, 256))

discriminator = PatchDiscriminator()
output = discriminator(image1)
print(output.shape)

torch.Size([1, 1, 4, 8])
```

Model Parameters

```
In [20]: total_params = sum(p.numel() for p in discriminator.parameters())
total_params
```

```
Out[20]: 2766529
```

```
In [21]: discriminator = to_device(discriminator, device)
generator = to_device(generator, device)
```

Training

Loss Functions

```
In [22]: comparison_loss = nn.BCEWithLogitsLoss()
L1_loss_fn = nn.L1Loss()
```

PaddleOCR

```
In [23]: ocr_model = PaddleOCR(use_angle_cls=False, lang='en', det=False, rec=True, use_gpu=False)

def tensor_to_bgr_numpy(tensor_img, mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5)):
    if tensor_img.dim() == 4:
        tensor_img = tensor_img[0]

    if tensor_img.shape[0] != 3:
        raise ValueError(f"Expected 3 channels, got shape {tensor_img.shape}")

    unnorm = torch.zeros_like(tensor_img)
    for c in range(3):
        unnorm[c] = tensor_img[c] * float(std[c]) + float(mean[c])
    tensor_img = unnorm.clamp(0, 1)

    pil_img = TF.to_pil_image(tensor_img.cpu())
    np_img = np.array(pil_img)[:, :, ::-1]

    return np_img
```

```
[2025/09/01 07:00:04] ppocr DEBUG: Namespace(help='==SUPPRESS==', use_gpu=True, use_xpu=False, use_npu=False, use_mlu=False, use_gcu=False, ir_optim=True, use_tensorrt=False, min_subgraph_size=15, precision='fp32', gpu_mem=500, gpu_id=0, image_dir=None, page_num=0, det_algorithm='DB', det_model_dir='C:\\Users\\ADMIN/.paddleocr/whl\\det\\en\\PP-OCRv3_det_infer', det_limit_side_len=960, det_limit_type='max', det_box_type='quad', det_db_thresh=0.3, det_db_box_thresh=0.6, det_db_unclip_ratio=1.5, max_batch_size=10, use_dilation=False, det_db_score_mode='fast', det_east_score_thresh=0.8, det_east_cover_thresh=0.1, det_east_nms_thresh=0.2, det_sast_score_thresh=0.5, det_sast_nms_thresh=0.2, det_pse_thresh=0, det_pse_box_thresh=0.85, det_pse_min_area=16, det_pse_scale=1, scales=[8, 16, 32], alpha=1.0, beta=1.0, fourier_degree=5, rec_algorithm='SVTR_LCNet', rec_model_dir='C:\\Users\\ADMIN/.paddleocr/whl\\rec\\en\\PP-OCRv4_rec_infer', rec_image_inverse=True, rec_image_shape='3, 48, 320', rec_batch_num=6, max_text_length=25, rec_char_dict_path='C:\\\\Thesis\\\\LiPAD with Paddle\\\\lipadenv3.9\\\\lib\\\\site-packages\\\\paddleocr\\\\ppocr\\\\utils\\\\en_dict.txt', use_space_char=True, vis_font_path='./doc/fonts/simfang.ttf', drop_score=0.5, e2e_algorithm='PGNet', e2e_model_dir=None, e2e_limit_side_len=768, e2e_limit_type='max', e2e_pgnet_score_thresh=0.5, e2e_char_dict_path='./ppocr/utils/ic15_dict.txt', e2e_pgnet_valid_set='totaltext', e2e_pgnet_mode='fast', use_angle_cls=False, cls_model_dir='C:\\\\Users\\ADMIN/.paddleocr/whl\\cls\\ch_ppocr_mobile_v2.0_cls_infer', cls_image_shape='3, 48, 192', label_list=['0', '180'], cls_batch_num=6, cls_thresh=0.9, enable_mkldnn=False, cpu_threads=10, use_pderving=False, warmup=False, sr_model_dir=None, sr_image_shape='3, 32, 128', sr_batch_num=1, draw_img_save_dir='./inference_results', save_crop_res=False, crop_res_save_dir='./output', use_mp=False, total_process_num=1, process_id=0, benchmark=False, save_log_path='./log_output/', show_log=True, use_onnx=False, onnx_providers=False, onnx_sess_options=False, return_word_box=False, output='./output', table_max_len=488, table_algorithm='TableAttn', table_model_dir=None, merge_no_span_structure=True, table_char_dict_path=None, formula_algorithm='LaTeXOCR', formula_model_dir=None, formula_char_dict_path=None, formula_batch_num=1, layout_model_dir=None, layout_dict_path=None, layout_score_threshold=0.5, layout_nms_threshold=0.5, kie_algorithm='LayoutXLM', ser_model_dir=None, re_model_dir=None, use_visual Backbone=True, ser_dict_path='../train_data/XFUND/class_list_xfun.txt', ocr_order_method=None, mode='structure', image_orientation=False, layout=True, table=True, formula=False, ocr=True, recovery=False, recovery_to_markdown=False, use_pdf2docx_api=False, invert=False, binarize=False, alphacolor=(255, 255, 255), lang='en', det=False, rec=True, type='ocr', savefile=False, ocr_version='PP-OCRv4', structure_version='PP-StructureV2')
```

```
[2025/09/01 07:00:04] ppocr WARNING: The first GPU is used for inference by default, GPU ID: 0
```

```
[2025/09/01 07:00:05] ppocr WARNING: The first GPU is used for inference by default, GPU ID: 0
```

Discriminator Training Function

```
In [24]: def train_discriminator(discriminator, generator, inputs, targets, d_opt):
    discriminator.train()

    real_images = targets
    fake_images = generator(inputs).detach()

    pred_real = discriminator(real_images)
    pred_fake = discriminator(fake_images)

    real_labels = torch.ones_like(pred_real)
    fake_labels = torch.zeros_like(pred_fake)
```

```

real_score = torch.sigmoid(pred_real).mean().item()
fake_score = torch.sigmoid(pred_fake).mean().item()

loss_real = comparison_loss(pred_real, real_labels)
loss_fake = comparison_loss(pred_fake, fake_labels)

d_loss = (loss_real + loss_fake) / 2

d_opt.zero_grad()
d_loss.backward()
d_opt.step()

return d_loss.item(), real_score, fake_score

```

Generator Training Function

```

In [25]: def ssim(pred, target):
    return 1 - ssim_fn(pred, target, data_range=1.0, size_average=True)

lpips_loss = LPIPS(net='vgg').to(device)
lpips_loss.eval()
for param in lpips_loss.parameters():
    param.requires_grad = False

def train_generator(generator, discriminator, inputs, targets, g_opt, adv_lambda=1.
generator.train()

fake_images = generator(inputs)
pred_fake = discriminator(fake_images)

real_labels = torch.ones_like(pred_fake)

adv_loss = comparison_loss(pred_fake, real_labels)
l1_loss = L1_loss_fn(fake_images, targets)
perceptual_loss = lpips_loss(fake_images.clamp(-1, 1), targets.clamp(-1, 1)).me
ssim_loss = ssim(fake_images, targets)

# Text Loss
fake_np = tensor_to_bgr_numpy(fake_images)
real_np = tensor_to_bgr_numpy(targets)

with torch.no_grad():
    fake_text = ocr_model.ocr(fake_np, cls=False, det=False)[0]
    real_text = ocr_model.ocr(real_np, cls=False, det=False)[0]

fake_str = fake_text[0][0] if fake_text else ''
real_str = real_text[0][0] if real_text else ''

fake_str = fake_str.strip().replace(' ', '')
real_str = real_str.strip().replace(' ', '')

# Use normalized edit distance as loss
edit_dist = levenshtein_distance(fake_str, real_str)
norm_edit_dist = edit_dist / max(len(real_str), 1)
text_loss = torch.tensor(norm_edit_dist, device=device)

```

```

g_loss = (
    adv_loss * adv_lambda +
    l1_loss * l1_lambda +
    ssim_loss * ssim_lambda +
    perceptual_loss * perceptual_lambda +
    text_loss * text_lambda
)

g_opt.zero_grad()
g_loss.backward()
g_opt.step()

return g_loss.item(), fake_images, {
    "total_loss": g_loss.item(),
    "adv": adv_loss.item(),
    "l1": l1_loss.item(),
    "ssim": ssim_loss.item(),
    "perceptual": perceptual_loss.item(),
    "text": text_loss.item()
}

```

Setting up [LPIPS] perceptual loss: trunk [vgg], v[0.1], spatial [off]
Loading model from: C:\Thesis\LiPAD with Paddle\lipadenv3.9\lib\site-packages\lpips\weights\v0.1\vgg.pth

```

In [26]: psnr_metric = PeakSignalNoiseRatio(data_range=1.0).to(device)
ssim_metric = StructuralSimilarityIndexMeasure(data_range=1.0).to(device)
fid_metric = FrechetInceptionDistance(feature=2048).to(device)

def to_uint8(tensor):
    # Clamp to [0, 1], scale to [0, 255], then convert to uint8
    tensor = torch.clamp(tensor, 0.0, 1.0)
    tensor = (tensor * 255.0).to(torch.uint8)
    return tensor

def evaluate_test(generator):
    raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl,
                                                                fig, axes = plt.subplots(1, 2, figsize=(10, 8)) # 1 row, 2 columns

    # Raw Test Set
    axes[0].imshow(raw_grid_img)
    axes[0].axis("off")
    axes[0].set_title("Raw Test Set")

    # Generated Test Set
    axes[1].imshow(gen_grid_img)
    axes[1].axis("off")
    axes[1].set_title("Generated Test Set")

    plt.tight_layout()
    plt.show()

```

```

In [27]: def create_image_grid_from_loader(generator, dataloader, device, num_images=16, image_size=256):

```

```

generator.to(device)

raw_images = []
gen_images = []
gen_count = 0
raw_count = 0

for batch in dataloader:
    # Assume batch is either just images or (images, labels)
    if isinstance(batch, (tuple, list)):
        inputs = batch[0]
    else:
        inputs = batch

    inputs = inputs.to(device)
    outputs = generator(inputs)

    for img in inputs:
        if denormalize:
            img = denorm(img)

        img_resized = TF.resize(img, image_size)
        raw_images.append(img_resized)

        raw_count += 1
        if raw_count >= num_images:
            break

    for out_img in outputs:
        if denormalize:
            out_img = denorm(out_img)

        out_img_resized = TF.resize(out_img, image_size)
        gen_images.append(out_img_resized)

        gen_count += 1
        if gen_count >= num_images:
            break

    if raw_count >= num_images:
        break

raw_grid = make_grid(raw_images, nrow=4)
gen_grid = make_grid(gen_images, nrow=4)

raw_ndarr = raw_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()
gen_ndarr = gen_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()

raw_grid_image = Image.fromarray(raw_ndarr)
gen_grid_image = Image.fromarray(gen_ndarr)

if save_path:
    gen_grid_image.save(save_path)
    print(f"Saved image grid to {save_path}")

return raw_grid_image, gen_grid_image

```

```
In [28]: @torch.no_grad()
def evaluate_generator(generator, valid_dl, epoch):
    generator.eval()
    psnr_vals = []
    ssim_vals = []
    fid_metric.reset()
    shown = False

    for val_inputs, val_targets in valid_dl:
        val_inputs = val_inputs.to(device)
        val_targets = val_targets.to(device)

        val_outputs = generator(val_inputs)

        if not shown:
            print('Validation Data')
            print_images(val_inputs, 5)
            print_images(val_outputs, 5)
            print_images(val_targets, 5)
            shown = True

        # Convert to uint8 for FID
        gen_uint8 = to_uint8(val_outputs)
        target_uint8 = to_uint8(val_targets)

        fid_metric.update(gen_uint8, real=False)
        fid_metric.update(target_uint8, real=True)

        # Compute PSNR/SSIM
        psnr = psnr_metric(val_outputs, val_targets).item()
        ssim = ssim_metric(val_outputs, val_targets).item()
        psnr_vals.append(psnr)
        ssim_vals.append(ssim)

        mean_psnr = sum(psnr_vals) / len(psnr_vals)
        mean_ssime = sum(ssim_vals) / len(ssim_vals)
        fid = fid_metric.compute().item()

        print(f"Epoch [{epoch+1}/{epochs}] - Mean PSNR: {mean_psnr:.4f}, Mean SSIM: {me

    evaluate_test(generator)
    return mean_psnr, mean_ssime, fid

def fit(generator, discriminator, epochs, lr, adv_lambda=1.0, l1_lambda=100.0, ssim_perceptual_lambda=5.0, text_lambda=7.5, g_opt=None, d_opt=None, checkpoint=None):
    torch.cuda.empty_cache()

    losses_g = []
    losses_d = []
    real_scores = []
    fake_scores = []
    psnrs = []
    ssims = []
    fids = []
```

```

g_losses_trace = []

# Create optimizers
if d_opt is None:
    d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
if g_opt is None:
    g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))

# Load checkpoint state if resuming
if checkpoint is not None:
    print("Resuming training from checkpoint...")
    d_opt.load_state_dict(checkpoint['opt_d'])
    g_opt.load_state_dict(checkpoint['opt_g'])

# Train
for epoch in range(epochs):
    torch.cuda.empty_cache()
    generator.train()
    discriminator.train()
    for inputs, targets in tqdm(train_dl):
        inputs = inputs.to(device)
        targets = targets.to(device)

        # Train discriminator
        d_loss, real_score, fake_score = train_discriminator(discriminator, gen
                                                               ...

        # Train generator
        g_loss, fake_images, g_loss_trace = train_generator(generator, discrimi
                                                               ...
                                                               ...

        print('Training Data')
        print_images(inputs, 5)
        print_images(fake_images, 5)
        print_images(targets, 5)

        if (epoch + 1) % 5 == 0:
            torch.save(generator.state_dict(), f'[RAU-NET]generator-Datasetv2-{epoch}.pt')
            torch.save(discriminator.state_dict(), f'[RAU-NET]discriminator-Datasetv2-{epoch}.pt')
            torch.save({
                'generator': generator.state_dict(),
                'discriminator': discriminator.state_dict(),
                'opt_g': g_opt.state_dict(),
                'opt_d': d_opt.state_dict(),
                'losses_g': losses_g,
                'losses_d': losses_d,
                'real_score': real_scores,
                'fake_score': fake_scores,
                'psnrs': psnrs,
                'ssims': ssims,
                'fids': fids,
                'epoch': epoch + 1
            }, f'[RAU-NET]Checkpoint-Logs-Datasetv2-{epoch + 1} epoch.pt')

# Print progress
try:
    print('Epoch [{}/{}], loss_g: {:.4f}, loss_d: {:.4f}, real_score: {:.4f}

```

```

        epoch+1, epochs, g_loss, d_loss, real_score, fake_score))
except Exception as err:
    print('Error:', str(err))

mean_psnr, mean_ssim, fid = evaluate_generator(generator, valid_dl, epoch)

# Record losses and scores
losses_g.append(g_loss)
losses_d.append(d_loss)
real_scores.append(real_score)
fake_scores.append(fake_score)
psnrs.append(mean_psnr)
ssims.append(mean_ssim)
fids.append(fid)
g_losses_trace.append(g_loss_trace)

save_samples(generator, test_dl, epoch + 1, sample_dir='./visualization')

return losses_g, losses_d, real_scores, fake_scores, psnrs, ssims, fids, g_loss

```

Hyperparameter Configurations

```
In [29]: adv_lambda = 0.25 # Change to 0.25
l1_lambda = 100
ssim_lambda = 5 # Change to 5
perceptual_lambda = 5.0
text_lambda = 10 # Change to 10

lr = 0.0002
beta1 = 0.5
beta2 = 0.999

epochs = 100
```

```
In [30]: g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))
d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
```

```
In [31]: history = []
```

```
In [32]: %time
history += fit(generator,
                discriminator,
                epochs,
                lr,
                adv_lambda=adv_lambda,
                l1_lambda=l1_lambda,
                ssim_lambda=ssim_lambda,
                perceptual_lambda=perceptual_lambda,
                text_lambda=text_lambda,
                g_opt=g_opt,
                d_opt=d_opt)
```

100% |

1250/1250 [19:06<00:00, 1.09it/s]

Training Data



Epoch [1/100], loss_g: 13.7348, loss_d: 0.1594, real_score: 0.9313, fake_score: 0.21
12

Validation Data



Epoch [1/100] - Mean PSNR: 15.0363, Mean SSIM: 0.6222, FID: 88.5737075805664

Raw Test Set

Generated Test Set



100% |

1250/1250 [19:07<00:00, 1.09it/s]

Training Data



Epoch [2/100], loss_g: 15.3838, loss_d: 0.9125, real_score: 0.9743, fake_score: 0.8155

Validation Data



Epoch [2/100] - Mean PSNR: 15.5397, Mean SSIM: 0.6449, FID: 75.14977264404297

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:53<00:00, 1.10it/s]

Training Data





Epoch [3/100], loss_g: 18.3615, loss_d: 0.3157, real_score: 0.9976, fake_score: 0.4491

Validation Data



Epoch [3/100] - Mean PSNR: 15.3667, Mean SSIM: 0.6459, FID: 72.73918914794922

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]

Training Data





Epoch [4/100], loss_g: 18.2271, loss_d: 0.9220, real_score: 0.5310, fake_score: 0.6705

Validation Data



Epoch [4/100] - Mean PSNR: 15.9021, Mean SSIM: 0.6682, FID: 68.24333190917969

Raw Test Set

Generated Test Set



100% | 1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [5/100], loss_g: 11.9212, loss_d: 0.1475, real_score: 0.8396, fake_score: 0.10
92

Validation Data



Epoch [5/100] - Mean PSNR: 16.1336, Mean SSIM: 0.6742, FID: 65.99771881103516

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:35<00:00, 1.12it/s]

Training Data



Epoch [6/100], loss_g: 24.8820, loss_d: 0.2865, real_score: 0.9985, fake_score: 0.42
07

Validation Data



Epoch [6/100] - Mean PSNR: 16.1009, Mean SSIM: 0.6767, FID: 66.40089416503906

Raw Test Set

Generated Test Set



100% |
1250/1250 [19:30<00:00, 1.07it/s]

Training Data



Epoch [7/100], loss_g: 13.5259, loss_d: 0.1281, real_score: 0.9729, fake_score: 0.2005

Validation Data





Epoch [7/100] - Mean PSNR: 16.3089, Mean SSIM: 0.6765, FID: 64.02198791503906

Raw Test Set

Generated Test Set



100%

1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [8/100], loss_g: 12.2336, loss_d: 0.5234, real_score: 0.9916, fake_score: 0.62

38

Validation Data





Epoch [8/100] - Mean PSNR: 16.3767, Mean SSIM: 0.6826, FID: 61.568870544433594

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]
Training Data



Epoch [9/100], loss_g: 23.5805, loss_d: 0.3780, real_score: 0.5339, fake_score: 0.08
79

Validation Data



Epoch [9/100] - Mean PSNR: 16.4338, Mean SSIM: 0.6862, FID: 62.23197555541992



100% | 1250/1250 [18:31<00:00, 1.12it/s]
Training Data



Epoch [10/100], loss_g: 17.9833, loss_d: 0.5120, real_score: 0.9517, fake_score: 0.5866

Validation Data



Epoch [10/100] - Mean PSNR: 16.3268, Mean SSIM: 0.6792, FID: 60.522098541259766



100% | 1250/1250 [18:59<00:00, 1.10it/s]
Training Data



Epoch [11/100], loss_g: 13.2169, loss_d: 0.3397, real_score: 0.8349, fake_score: 0.3814

Validation Data



Epoch [11/100] - Mean PSNR: 16.3676, Mean SSIM: 0.6774, FID: 59.345088958740234

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:32<00:00, 1.12it/s]

Training Data





Epoch [12/100], loss_g: 13.1601, loss_d: 0.5280, real_score: 0.9245, fake_score: 0.6032

Validation Data



Epoch [12/100] - Mean PSNR: 16.6376, Mean SSIM: 0.6876, FID: 57.72936248779297

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:36<00:00, 1.12it/s]

Training Data





Epoch [13/100], loss_g: 10.1245, loss_d: 0.7096, real_score: 0.4820, fake_score: 0.4654

Validation Data



Epoch [13/100] - Mean PSNR: 16.4892, Mean SSIM: 0.6833, FID: 58.29381561279297

Raw Test Set

Generated Test Set



100% | 1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [14/100], loss_g: 18.7539, loss_d: 0.6209, real_score: 0.9795, fake_score: 0.6

467

Validation Data



Epoch [14/100] - Mean PSNR: 16.7355, Mean SSIM: 0.6924, FID: 57.52530288696289

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:28<00:00, 1.13it/s]

Training Data



Epoch [15/100], loss_g: 10.0308, loss_d: 0.3355, real_score: 0.7233, fake_score: 0.2

732

Validation Data



Epoch [15/100] - Mean PSNR: 16.7468, Mean SSIM: 0.6896, FID: 55.776702880859375

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [16/100], loss_g: 8.8388, loss_d: 0.5859, real_score: 0.4009, fake_score: 0.1757

Validation Data





Epoch [16/100] - Mean PSNR: 16.5479, Mean SSIM: 0.6874, FID: 56.621944427490234
Raw Test Set Generated Test Set



100% |
1250/1250 [18:33<00:00, 1.12it/s]
Training Data



Epoch [17/100], loss_g: 9.1669, loss_d: 0.3975, real_score: 0.5479, fake_score: 0.1433

Validation Data





Epoch [17/100] - Mean PSNR: 16.5876, Mean SSIM: 0.6920, FID: 56.302650451660156

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:29<00:00, 1.13it/s]

Training Data



Epoch [18/100], loss_g: 10.5401, loss_d: 0.8478, real_score: 0.9960, fake_score: 0.7920

Validation Data



Epoch [18/100] - Mean PSNR: 16.6729, Mean SSIM: 0.6938, FID: 58.11103439331055



100% |
1250/1250 [18:29<00:00, 1.13it/s]
Training Data



Epoch [19/100], loss_g: 11.3039, loss_d: 0.4145, real_score: 0.7972, fake_score: 0.4
311

Validation Data



Epoch [19/100] - Mean PSNR: 16.5853, Mean SSIM: 0.6894, FID: 55.42158889770508



100% |
1250/1250 [18:30<00:00, 1.13it/s]
Training Data



Epoch [20/100], loss_g: 10.0287, loss_d: 0.3361, real_score: 0.6721, fake_score: 0.2093

Validation Data



Epoch [20/100] - Mean PSNR: 16.6502, Mean SSIM: 0.6918, FID: 56.36728286743164

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]

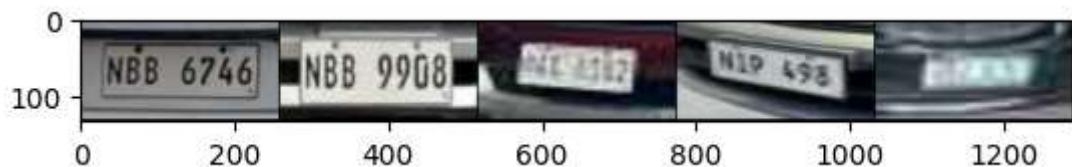
Training Data





Epoch [21/100], loss_g: 8.9730, loss_d: 1.3297, real_score: 0.0828, fake_score: 0.0494

Validation Data



Epoch [21/100] - Mean PSNR: 16.5212, Mean SSIM: 0.6864, FID: 54.23642349243164

Raw Test Set

Generated Test Set



100% | [1250/1250 [18:31<00:00, 1.12it/s]

Training Data





Epoch [22/100], loss_g: 9.5040, loss_d: 0.2122, real_score: 0.9590, fake_score: 0.30
90

Validation Data



Epoch [22/100] - Mean PSNR: 16.6216, Mean SSIM: 0.6889, FID: 54.598609924316406

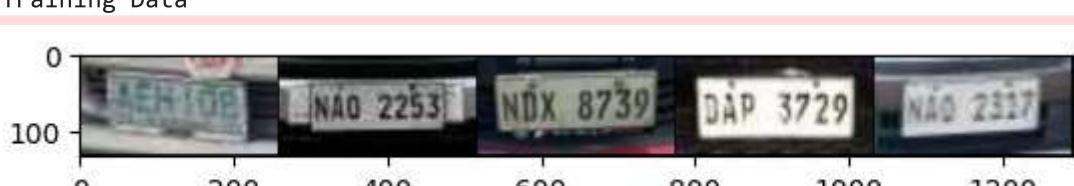
Raw Test Set

Generated Test Set



100% | 1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [23/100], loss_g: 12.8219, loss_d: 0.3694, real_score: 0.6633, fake_score: 0.2509

Validation Data



Epoch [23/100] - Mean PSNR: 16.7477, Mean SSIM: 0.6956, FID: 54.611846923828125

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [24/100], loss_g: 9.1771, loss_d: 0.4260, real_score: 0.6578, fake_score: 0.3227

Validation Data



Epoch [24/100] - Mean PSNR: 16.9083, Mean SSIM: 0.6952, FID: 53.58484649658203

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [25/100], loss_g: 8.3415, loss_d: 0.5109, real_score: 0.4099, fake_score: 0.05
73

Validation Data

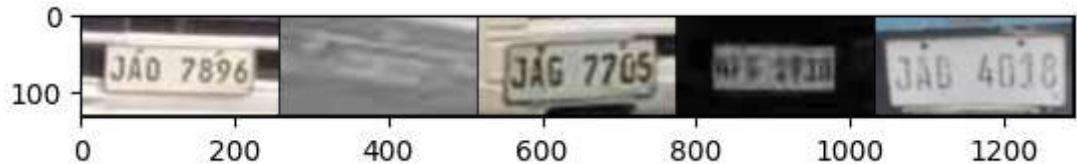




Epoch [25/100] - Mean PSNR: 16.7812, Mean SSIM: 0.6967, FID: 53.72076416015625
 Raw Test Set Generated Test Set



100% |
 1250/1250 [18:31<00:00, 1.12it/s]
 Training Data



Epoch [26/100], loss_g: 9.3117, loss_d: 1.8837, real_score: 0.0285, fake_score: 0.0254
 Validation Data





Epoch [26/100] - Mean PSNR: 16.8415, Mean SSIM: 0.6935, FID: 52.91928482055664

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [27/100], loss_g: 8.1864, loss_d: 0.2910, real_score: 0.7740, fake_score: 0.25
95

Validation Data



Epoch [27/100] - Mean PSNR: 16.6552, Mean SSIM: 0.6911, FID: 53.405426025390625



100% |
1250/1250 [18:30<00:00, 1.13it/s]
Training Data



Epoch [28/100], loss_g: 17.6641, loss_d: 0.4884, real_score: 0.5078, fake_score: 0.2
155

Validation Data



Epoch [28/100] - Mean PSNR: 16.7275, Mean SSIM: 0.6922, FID: 52.65266799926758



100% |
1250/1250 [18:31<00:00, 1.12it/s]
Training Data



Epoch [29/100], loss_g: 18.3687, loss_d: 0.3072, real_score: 0.7517, fake_score: 0.2611

Validation Data



Epoch [29/100] - Mean PSNR: 16.8364, Mean SSIM: 0.6956, FID: 53.03778076171875

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]

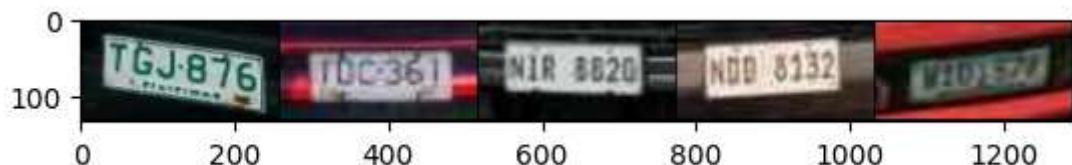
Training Data





Epoch [30/100], loss_g: 9.6385, loss_d: 0.3989, real_score: 0.7713, fake_score: 0.3899

Validation Data



Epoch [30/100] - Mean PSNR: 16.6369, Mean SSIM: 0.6930, FID: 52.135868072509766

Raw Test Set

Generated Test Set



100% | [1250/1250 [18:31<00:00, 1.12it/s]

Training Data





Epoch [31/100], loss_g: 9.6087, loss_d: 0.9785, real_score: 0.1785, fake_score: 0.07
80

Validation Data



Epoch [31/100] - Mean PSNR: 16.7697, Mean SSIM: 0.6900, FID: 51.67242431640625

Raw Test Set

Generated Test Set



100% | 1250/1250 [18:30<00:00, 1.13it/s]

Training Data



Epoch [32/100], loss_g: 13.5159, loss_d: 0.2031, real_score: 0.8943, fake_score: 0.2
428

Validation Data



Epoch [32/100] - Mean PSNR: 16.7448, Mean SSIM: 0.6914, FID: 52.12015151977539

Raw Test Set

Generated Test Set



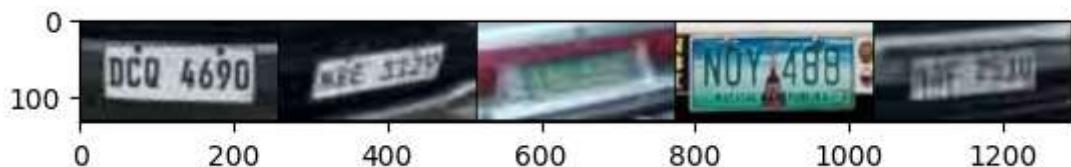
100% |
1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [33/100], loss_g: 15.4354, loss_d: 0.3949, real_score: 0.6614, fake_score: 0.2
810

Validation Data



Epoch [33/100] - Mean PSNR: 16.5437, Mean SSIM: 0.6859, FID: 51.8897819519043

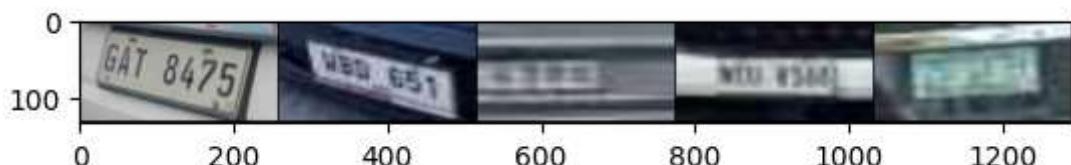
Raw Test Set

Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]

Training Data



Epoch [34/100], loss_g: 7.8392, loss_d: 0.4212, real_score: 0.7885, fake_score: 0.42
13

Validation Data

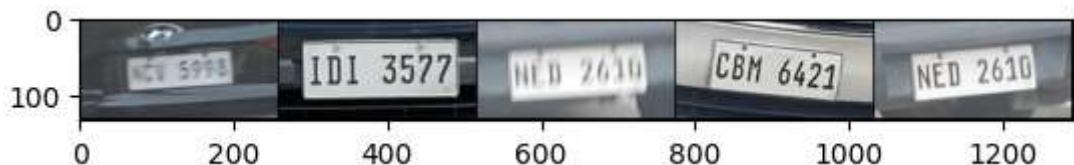




Epoch [34/100] - Mean PSNR: 16.7450, Mean SSIM: 0.6930, FID: 52.21112060546875
Raw Test Set Generated Test Set



100% |
1250/1250 [18:31<00:00, 1.12it/s]
Training Data



Epoch [35/100], loss_g: 6.5918, loss_d: 0.6630, real_score: 0.3682, fake_score: 0.20
56

Validation Data





Epoch [35/100] - Mean PSNR: 16.7751, Mean SSIM: 0.6902, FID: 51.03852462768555

Raw Test Set

Generated Test Set



100% |

1250/1250 [19:08<00:00, 1.09it/s]

Training Data



Epoch [36/100], loss_g: 7.3100, loss_d: 0.3983, real_score: 0.5890, fake_score: 0.19
24

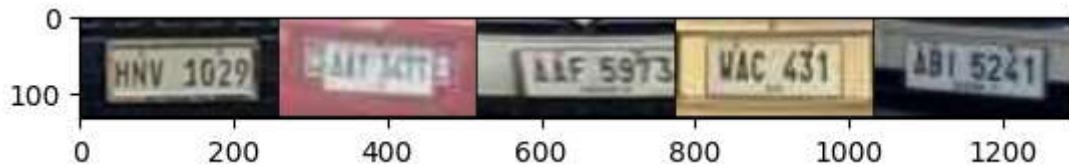
Validation Data



Epoch [36/100] - Mean PSNR: 16.8161, Mean SSIM: 0.6932, FID: 52.11528778076172



100% | 1250/1250 [18:32<00:00, 1.12it/s]
Training Data



Epoch [37/100], loss_g: 8.1810, loss_d: 0.3126, real_score: 0.9423, fake_score: 0.40
35

Validation Data



Epoch [37/100] - Mean PSNR: 16.6537, Mean SSIM: 0.6902, FID: 51.098419189453125

Raw Test Set

Generated Test Set



100% | 1250/1250 [18:28<00:00, 1.13it/s]
Training Data



Epoch [38/100], loss_g: 8.9853, loss_d: 0.4928, real_score: 0.5269, fake_score: 0.24
44

Validation Data



Epoch [38/100] - Mean PSNR: 16.8308, Mean SSIM: 0.6937, FID: 51.71621322631836

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:51<00:00, 1.10it/s]

Training Data





Epoch [39/100], loss_g: 8.7576, loss_d: 0.4374, real_score: 0.4941, fake_score: 0.0986

Validation Data



Epoch [39/100] - Mean PSNR: 16.6123, Mean SSIM: 0.6885, FID: 51.20871353149414

Raw Test Set

Generated Test Set



100% | [1250/1250 [19:02<00:00, 1.09it/s]

Training Data





Epoch [40/100], loss_g: 9.0074, loss_d: 0.5024, real_score: 0.5204, fake_score: 0.23
98

Validation Data



Epoch [40/100] - Mean PSNR: 16.7455, Mean SSIM: 0.6945, FID: 50.74522399902344

Raw Test Set

Generated Test Set



100% | 

1250/1250 [19:28<00:00, 1.07it/s]

Training Data



Epoch [41/100], loss_g: 14.7173, loss_d: 0.2921, real_score: 0.6945, fake_score: 0.1692

Validation Data



Epoch [41/100] - Mean PSNR: 16.7378, Mean SSIM: 0.6915, FID: 51.01935958862305

Raw Test Set

Generated Test Set



100%

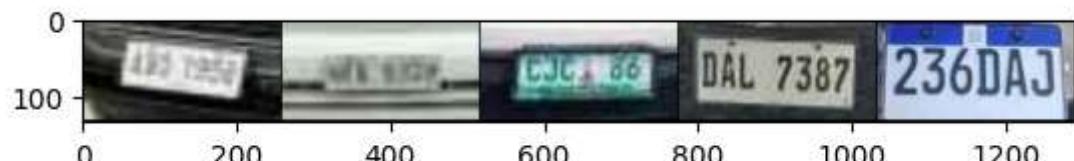
1250/1250 [19:20<00:00, 1.08it/s]

Training Data



Epoch [42/100], loss_g: 20.9016, loss_d: 0.5313, real_score: 0.5407, fake_score: 0.3079

Validation Data



Epoch [42/100] - Mean PSNR: 16.8750, Mean SSIM: 0.6949, FID: 52.39639663696289

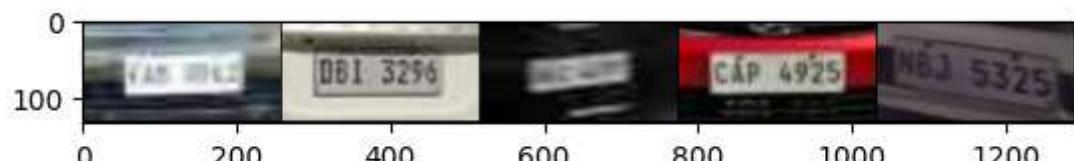
Raw Test Set

Generated Test Set



100% |
1250/1250 [19:26<00:00, 1.07it/s]

Training Data



Epoch [43/100], loss_g: 13.8600, loss_d: 0.5109, real_score: 0.5384, fake_score: 0.2821

Validation Data





Epoch [43/100] - Mean PSNR: 16.8274, Mean SSIM: 0.6907, FID: 50.910499572753906
Raw Test Set Generated Test Set



100% |
1250/1250 [18:52<00:00, 1.10it/s]

Training Data



Epoch [44/100], loss_g: 13.3424, loss_d: 0.4471, real_score: 0.6344, fake_score: 0.3

117

Validation Data





Epoch [44/100] - Mean PSNR: 16.6923, Mean SSIM: 0.6882, FID: 51.8835563659668

Raw Test Set

Generated Test Set



100% |

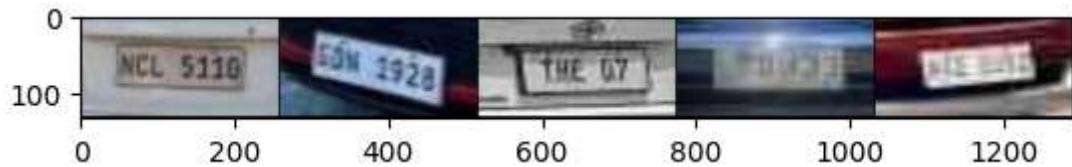
1250/1250 [18:32<00:00, 1.12it/s]

Training Data



Epoch [45/100], loss_g: 7.4127, loss_d: 1.0146, real_score: 0.1679, fake_score: 0.0858

Validation Data



Epoch [45/100] - Mean PSNR: 16.5438, Mean SSIM: 0.6854, FID: 51.01364517211914



100% |
1250/1250 [18:32<00:00, 1.12it/s]
Training Data



Epoch [46/100], loss_g: 7.7617, loss_d: 0.6117, real_score: 0.4447, fake_score: 0.24
78

Validation Data



Epoch [46/100] - Mean PSNR: 16.6507, Mean SSIM: 0.6854, FID: 51.20905303955078



100% |
1250/1250 [18:33<00:00, 1.12it/s]
Training Data



Epoch [47/100], loss_g: 6.7227, loss_d: 0.5923, real_score: 0.8769, fake_score: 0.5978

Validation Data



Epoch [47/100] - Mean PSNR: 16.6516, Mean SSIM: 0.6883, FID: 50.21856689453125

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:33<00:00, 1.12it/s]

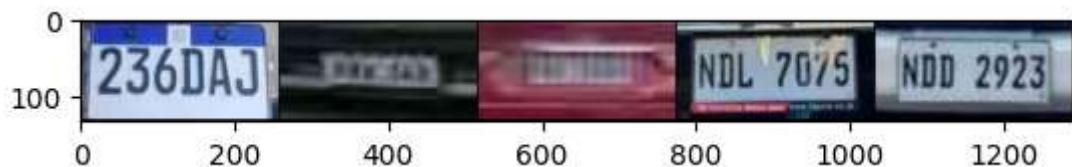
Training Data





Epoch [48/100], loss_g: 12.2718, loss_d: 0.2789, real_score: 0.7758, fake_score: 0.2349

Validation Data



Epoch [48/100] - Mean PSNR: 16.7517, Mean SSIM: 0.6905, FID: 51.063053131103516

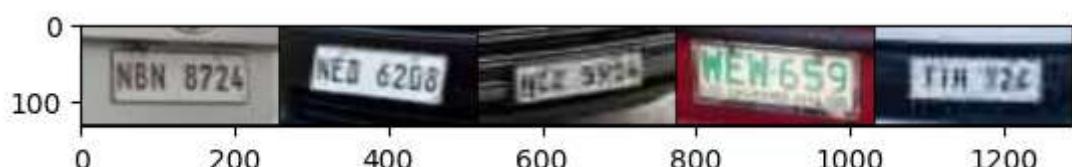
Raw Test Set

Generated Test Set



100% |
1250/1250 [18:32<00:00, 1.12it/s]

Training Data





Epoch [49/100], loss_g: 6.6962, loss_d: 0.4797, real_score: 0.4730, fake_score: 0.11
05

Validation Data



Epoch [49/100] - Mean PSNR: 16.7504, Mean SSIM: 0.6915, FID: 51.55146789550781
Raw Test Set Generated Test Set



100% |
1250/1250 [18:34<00:00, 1.12it/s]

Training Data



Epoch [50/100], loss_g: 10.4136, loss_d: 0.4632, real_score: 0.7812, fake_score: 0.4
484

Validation Data



Epoch [50/100] - Mean PSNR: 16.5840, Mean SSIM: 0.6856, FID: 50.91236877441406

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:32<00:00, 1.12it/s]

Training Data



Epoch [51/100], loss_g: 17.5845, loss_d: 0.2439, real_score: 0.8509, fake_score: 0.2
575

Validation Data



Epoch [51/100] - Mean PSNR: 16.7185, Mean SSIM: 0.6878, FID: 50.77513885498047

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:43<00:00, 1.11it/s]

Training Data



Epoch [52/100], loss_g: 8.0554, loss_d: 0.4911, real_score: 0.6765, fake_score: 0.38
62

Validation Data





Epoch [52/100] - Mean PSNR: 16.7433, Mean SSIM: 0.6878, FID: 50.79594039916992
Raw Test Set Generated Test Set



100% |
1250/1250 [19:03<00:00, 1.09it/s]
Training Data



Epoch [53/100], loss_g: 16.7479, loss_d: 0.4256, real_score: 0.8745, fake_score: 0.4683

Validation Data





Epoch [53/100] - Mean PSNR: 16.9420, Mean SSIM: 0.6901, FID: 50.3238639831543

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:37<00:00, 1.12it/s]

Training Data



Epoch [54/100], loss_g: 6.8816, loss_d: 0.4222, real_score: 0.7592, fake_score: 0.3785

Validation Data



Epoch [54/100] - Mean PSNR: 16.8547, Mean SSIM: 0.6906, FID: 51.44987869262695



100% | 1250/1250 [19:05<00:00, 1.09it/s]
Training Data



Epoch [55/100], loss_g: 8.0388, loss_d: 0.5639, real_score: 0.9150, fake_score: 0.59
92

Validation Data



Epoch [55/100] - Mean PSNR: 16.6732, Mean SSIM: 0.6885, FID: 50.865169525146484

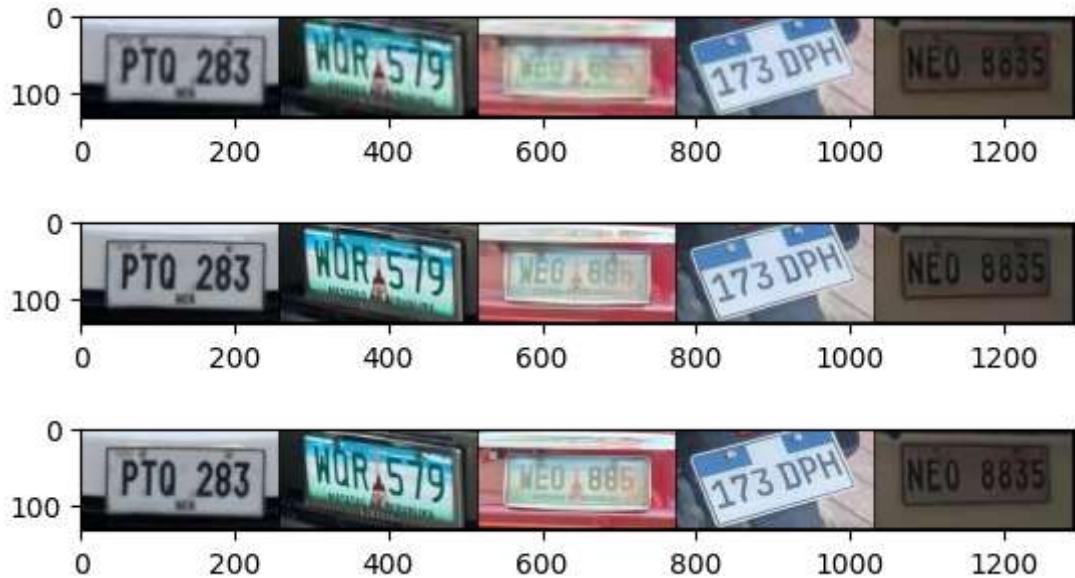


100% | 1250/1250 [19:13<00:00, 1.08it/s]
Training Data



Epoch [56/100], loss_g: 11.6357, loss_d: 0.2480, real_score: 0.9223, fake_score: 0.3124

Validation Data



Epoch [56/100] - Mean PSNR: 16.7014, Mean SSIM: 0.6880, FID: 51.10187530517578

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:36<00:00, 1.12it/s]

Training Data





Epoch [57/100], loss_g: 16.5970, loss_d: 0.5661, real_score: 0.9274, fake_score: 0.6107

Validation Data



Epoch [57/100] - Mean PSNR: 16.7396, Mean SSIM: 0.6861, FID: 51.19881820678711

Raw Test Set

Generated Test Set



100% | 1250/1250 [18:31<00:00, 1.12it/s]

Training Data





Epoch [58/100], loss_g: 7.0888, loss_d: 0.4762, real_score: 0.9241, fake_score: 0.53
65

Validation Data



Epoch [58/100] - Mean PSNR: 16.7776, Mean SSIM: 0.6880, FID: 51.152400970458984

Raw Test Set

Generated Test Set



100% | |

1250/1250 [18:36<00:00, 1.12it/s]

Training Data



Epoch [59/100], loss_g: 14.6100, loss_d: 0.5146, real_score: 0.7219, fake_score: 0.4507

Validation Data



Epoch [59/100] - Mean PSNR: 16.5912, Mean SSIM: 0.6844, FID: 50.85931396484375

Raw Test Set

Generated Test Set



100% |
1250/1250 [19:19<00:00, 1.08it/s]

Training Data



Epoch [60/100], loss_g: 6.2183, loss_d: 0.3801, real_score: 0.6578, fake_score: 0.2348

Validation Data



Epoch [60/100] - Mean PSNR: 16.6962, Mean SSIM: 0.6852, FID: 51.89915084838867

Raw Test Set

Generated Test Set



100% |
1250/1250 [19:11<00:00, 1.09it/s]

Training Data



Epoch [61/100], loss_g: 27.1288, loss_d: 0.6676, real_score: 0.8767, fake_score: 0.6443

Validation Data





Epoch [61/100] - Mean PSNR: 16.5761, Mean SSIM: 0.6839, FID: 51.232513427734375
Raw Test Set Generated Test Set



100% |
1250/1250 [18:35<00:00, 1.12it/s]
Training Data



Epoch [62/100], loss_g: 7.7187, loss_d: 0.6227, real_score: 0.3840, fake_score: 0.15
22

Validation Data





Epoch [62/100] - Mean PSNR: 16.7867, Mean SSIM: 0.6875, FID: 52.131446838378906

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:41<00:00, 1.11it/s]

Training Data



Epoch [63/100], loss_g: 9.8292, loss_d: 0.4396, real_score: 0.6724, fake_score: 0.32
10

Validation Data



Epoch [63/100] - Mean PSNR: 16.7710, Mean SSIM: 0.6878, FID: 51.19508743286133



100% |
1250/1250 [18:54<00:00, 1.10it/s]
Training Data



Epoch [64/100], loss_g: 7.8877, loss_d: 0.5136, real_score: 0.6005, fake_score: 0.34
63

Validation Data



Epoch [64/100] - Mean PSNR: 16.8282, Mean SSIM: 0.6896, FID: 51.36892318725586



100% |
1250/1250 [19:58<00:00, 1.04it/s]
Training Data



Epoch [65/100], loss_g: 9.8715, loss_d: 0.5169, real_score: 0.4627, fake_score: 0.1366

Validation Data



Epoch [65/100] - Mean PSNR: 16.7960, Mean SSIM: 0.6876, FID: 51.98481369018555

Raw Test Set

Generated Test Set



100% |
1250/1250 [21:03<00:00, 1.01s/it]

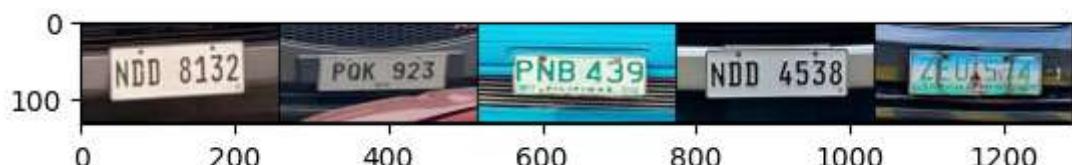
Training Data





Epoch [66/100], loss_g: 11.7402, loss_d: 0.3562, real_score: 0.7679, fake_score: 0.3145

Validation Data



Epoch [66/100] - Mean PSNR: 16.8601, Mean SSIM: 0.6882, FID: 51.82242965698242

Raw Test Set

Generated Test Set



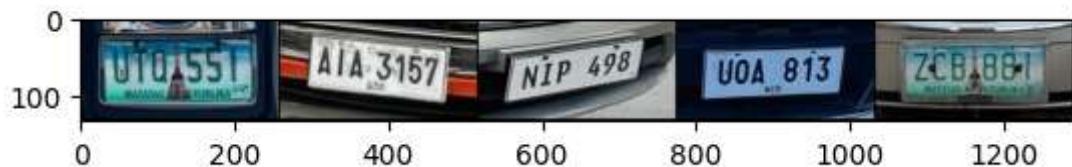
Training Data





Epoch [67/100], loss_g: 6.6485, loss_d: 0.4295, real_score: 0.8193, fake_score: 0.4205

Validation Data



Epoch [67/100] - Mean PSNR: 16.6254, Mean SSIM: 0.6850, FID: 50.46392822265625

Raw Test Set

Generated Test Set



1250/1250 [19:26<00:00, 1.07it/s]

Training Data



Epoch [68/100], loss_g: 14.4831, loss_d: 0.5181, real_score: 0.4964, fake_score: 0.2034

Validation Data



Epoch [68/100] - Mean PSNR: 16.5788, Mean SSIM: 0.6795, FID: 50.86103439331055

Raw Test Set

Generated Test Set



100% | 1250/1250 [18:44<00:00, 1.11it/s]

Training Data



Epoch [69/100], loss_g: 6.6487, loss_d: 0.3919, real_score: 0.8598, fake_score: 0.4204

Validation Data



Epoch [69/100] - Mean PSNR: 16.7326, Mean SSIM: 0.6862, FID: 50.71609115600586

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:45<00:00, 1.11it/s]

Training Data



Epoch [70/100], loss_g: 8.0742, loss_d: 0.3884, real_score: 0.8281, fake_score: 0.3952

Validation Data





Epoch [70/100] - Mean PSNR: 16.7447, Mean SSIM: 0.6862, FID: 51.09672927856445
Raw Test Set Generated Test Set



100% |
1250/1250 [18:42<00:00, 1.11it/s]
Training Data



Epoch [71/100], loss_g: 6.0035, loss_d: 0.4797, real_score: 0.6837, fake_score: 0.37
14

Validation Data





Epoch [71/100] - Mean PSNR: 16.6312, Mean SSIM: 0.6865, FID: 50.84730529785156

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:42<00:00, 1.11it/s]

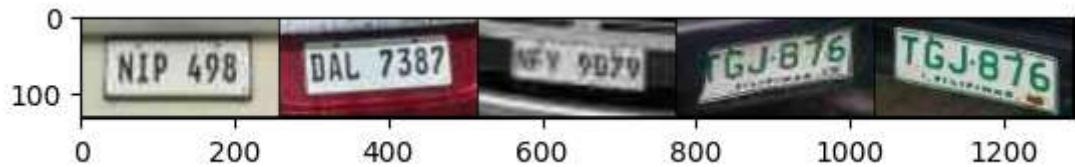
Training Data



Epoch [72/100], loss_g: 6.3332, loss_d: 0.5572, real_score: 0.4581, fake_score: 0.15

42

Validation Data



Epoch [72/100] - Mean PSNR: 16.6667, Mean SSIM: 0.6859, FID: 51.7897834777832



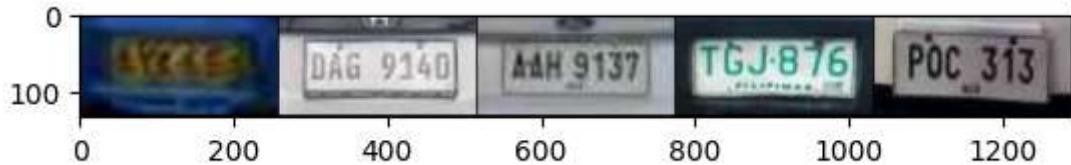
100% |
1250/1250 [18:43<00:00, 1.11it/s]

Training Data



Epoch [73/100], loss_g: 9.7908, loss_d: 0.3392, real_score: 0.8782, fake_score: 0.3789

Validation Data

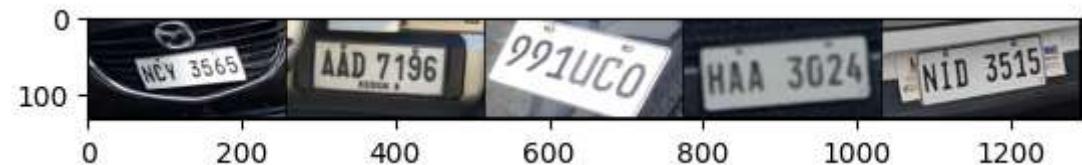


Epoch [73/100] - Mean PSNR: 16.8446, Mean SSIM: 0.6913, FID: 51.14340591430664



100% |
1250/1250 [19:10<00:00, 1.09it/s]

Training Data



Validation Data



Epoch [74/100] - Mean PSNR: 16.7647, Mean SSIM: 0.6874, FID: 51.344581604003906

Raw Test Set

Generated Test Set



100% |
1250/1250 [19:09<00:00, 1.09it/s]

Training Data





Epoch [75/100], loss_g: 6.3994, loss_d: 0.5533, real_score: 0.4785, fake_score: 0.20
17

Validation Data



Epoch [75/100] - Mean PSNR: 16.8720, Mean SSIM: 0.6894, FID: 50.776729583740234

Raw Test Set

Generated Test Set



0%
| 0/1250 [00:05<?, ?it/s]

```
-----  
KeyboardInterrupt                                     Traceback (most recent call last)  
File <timed exec>:1  
  
Cell In[28], line 75, in fit(generator, discriminator, epochs, lr, adv_lambda, l1_lambda, ssim_lambda, perceptual_lambda, text_lambda, g_opt, d_opt, checkpoint)  
    73     generator.train()  
    74     discriminator.train()  
--> 75 for inputs, targets in tqdm(train_dl):  
    76         inputs = inputs.to(device)  
    77         targets = targets.to(device)  
  
File C:\Thesis\LiPAD with Paddle\lipadvenv3.9\lib\site-packages\tqdm\std.py:1181, in  
tqdm.__iter__(self)  
1178     time = self._time  
1180     try:  
-> 1181         for obj in iterable:  
    1182             yield obj  
    1183             # Update and possibly print the progressbar.  
    1184             # Note: does not call self.update(1) for speed optimisation.  
  
Cell In[13], line 18, in DeviceDataLoader.__iter__(self)  
    17 def __iter__(self):  
--> 18     for b in self.dl:  
    19         yield to_device(b, self.device)  
  
File C:\Thesis\LiPAD with Paddle\lipadvenv3.9\lib\site-packages\torch\utils\data\dataloader.py:441, in DataLoader.__iter__(self)  
    439     return self._iterator  
    440 else:  
--> 441     return self._get_iterator()  
  
File C:\Thesis\LiPAD with Paddle\lipadvenv3.9\lib\site-packages\torch\utils\data\dataloader.py:388, in DataLoader._get_iterator(self)  
    386 else:  
    387     self.check_worker_number_rationality()  
--> 388     return _MultiProcessingDataLoaderIter(self)  
  
File C:\Thesis\LiPAD with Paddle\lipadvenv3.9\lib\site-packages\torch\utils\data\dataloader.py:1042, in _MultiProcessingDataLoaderIter.__init__(self, loader)  
    1035 w.daemon = True  
    1036 # NB: Process.start() actually take some time as it needs to  
    1037 #      start a process and pass the arguments over via a pipe.  
    1038 #      Therefore, we only add a worker to self._workers list after  
    1039 #      it started, so that we do not call .join() if program dies  
    1040 #      before it starts, and __del__ tries to join but will get:  
    1041 #          AssertionError: can only join a started process.  
-> 1042 w.start()  
    1043 self._index_queues.append(index_queue)  
    1044 self._workers.append(w)  
  
File ~\AppData\Local\Programs\Python\Python39\lib\multiprocessing\process.py:121, in  
BaseProcess.start(self)  
118 assert not _current_process._config.get('daemon'), \  
119             'daemonic processes are not allowed to have children'  
120 _cleanup()
```

```

--> 121 self._popen = self._Popen(self)
  122 self._sentinel = self._popen.sentinel
  123 # Avoid a refcycle if the target function holds an indirect
  124 # reference to the process object (see bpo-30775)

File ~\AppData\Local\Programs\Python\Python39\lib\multiprocessing\context.py:224, in
Process._Popen(process_obj)
  222 @staticmethod
  223 def _Popen(process_obj):
--> 224     return _default_context.get_context().Process._Popen(process_obj)

File ~\AppData\Local\Programs\Python\Python39\lib\multiprocessing\context.py:327, in
SpawnProcess._Popen(process_obj)
  324 @staticmethod
  325 def _Popen(process_obj):
  326     from .popen_spawn_win32 import Popen
--> 327     return Popen(process_obj)

File ~\AppData\Local\Programs\Python\Python39\lib\multiprocessing\popen_spawn_win32.
py:93, in Popen.__init__(self, process_obj)
  91 try:
  92     reduction.dump(prep_data, to_child)
--> 93     reduction.dump(process_obj, to_child)
  94 finally:
  95     set_spawning_popen(None)

File ~\AppData\Local\Programs\Python\Python39\lib\multiprocessing\reduction.py:60, i
n dump(obj, file, protocol)
  58 def dump(obj, file, protocol=None):
  59     '''Replacement for pickle.dump() using ForkingPickler.'''
--> 60     ForkingPickler(file, protocol).dump(obj)

KeyboardInterrupt:

```

Post Training

In [34]: `torch.cuda.empty_cache()`

Save Model

In [34]: `torch.save(generator.state_dict(), "[RAU-NET]generator-Datasetv2-5epoch.pt")
torch.save(discriminator.state_dict(), "[RAU-NET]discriminator-Datasetv2-5epoch.pt")`

In [35]: `torch.save({
 'generator_state_dict': generator.state_dict(),
 'discriminator_state_dict': discriminator.state_dict(),
 'opt_g_state_dict': g_opt.state_dict(),
 'opt_d_state_dict': d_opt.state_dict(),
 'history': history
}, f'[RAU-NET]Checkpoint-Datasetv2-30epoch.pt')`

Model Evaluation

```
In [ ]: psnr_vals = []
ssim_vals = []

fid_metric = FrechetInceptionDistance(feature=2048).to(device)
fid_metric.reset()

generator.eval()
for val_input, val_target in valid_dl:
    val_input = val_input.to(device)
    val_target = val_target.to(device)

    with torch.no_grad():
        gen_output = generator(val_input)

    # Convert to uint8 for FID
    gen_uint8 = to_uint8(gen_output)
    target_uint8 = to_uint8(val_target)

    fid_metric.update(gen_uint8, real=False)
    fid_metric.update(target_uint8, real=True)

    show_generated_images(gen_output)

    # Compute PSNR/SSIM
    psnr = psnr_metric(gen_output, val_target).item()
    ssim = ssim_metric(gen_output, val_target).item()
    psnr_vals.append(psnr)
    ssim_vals.append(ssim)

print("Mean PSNR:", sum(psnr_vals) / len(psnr_vals))
print("Mean SSIM:", sum(ssim_vals) / len(ssim_vals))
print("FID:", fid_metric.compute().item())
```











```
In [31]: def deblur_image(image_path, generator, device):
    image = Image.open(image_path).convert('RGB')
    input_tensor = valid_transform(image).unsqueeze(0).to(device)

    generator.eval()
    with torch.no_grad():
        gen_image = generator(input_tensor)

    output = denorm(gen_image.squeeze(0).cpu())
```

```
plt.imshow(output.permute(1, 2, 0)) # CHW → HWC
plt.axis('off')
plt.title("Generated Image")
plt.show()
```

In [32]: deblur_image(os.path.join(DATA_DIR, 'low_qual', 'test', 'low_qual', '48.jpg'), gene

Generated Image



```
raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl, devi
fig, axes = plt.subplots(1, 2, figsize=(15, 15)) # 1 row, 2 columns

# Raw Test Set
axes[0].imshow(raw_grid_img)
axes[0].axis("off")
axes[0].set_title("Raw Test Set")

# Generated Test Set
axes[1].imshow(gen_grid_img)
axes[1].axis("off")
axes[1].set_title("Generated Test Set")

plt.tight_layout()
plt.show()
```



```
In [36]: torch.cuda.empty_cache()
```

Load Model

```
In [34]: MODELS_DIR = os.path.join('..', '..', '..', 'models', 'h_blur')
```

```
In [37]: generator = AttentionResUNetGenerator()
loaded_g_wts = torch.load('[RAU-NET]generator-Datasetv2-55 epoch.pt')
generator.load_state_dict(loaded_g_wts)
generator = to_device(generator, device)

discriminator = PatchDiscriminator()
loaded_d_wts = torch.load('[RAU-NET]discriminator-Datasetv2-55 epoch.pt')
discriminator.load_state_dict(loaded_d_wts)
discriminator = to_device(discriminator, device)
```

```
In [29]: checkpoint = torch.load('[RAU-NET]Checkpoint-Logs-Datasetv2-75 epoch.pt')
#g_opt.load_state_dict(checkpoint['opt_g_state_dict'])
#d_opt.load_state_dict(checkpoint['opt_d_state_dict'])
#history = checkpoint['losses_g']
```

```
In [39]: def are_models_identical(model1, model2):
    # Check if state dictionaries have the same keys
    if model1.state_dict().keys() != model2.state_dict().keys():
        return False
```

```
# Compare each parameter's values
for key in model1.state_dict():
    if not torch.equal(model1.state_dict()[key], model2.state_dict()[key]):
        return False
return True
```

In [45]:

```
if are_models_identical(generator1, generator2):
    print("Models have identical weights and biases")
else:
    print("Models have different weights or biases")
```

Models have different weights or biases

In [9]:

```
len(history)
```

Out[9]: 79

Save Reconstructed Test Set

In [36]:

```
def save_generated_tests(generator, test_dl, sample_dir='./results'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)

            for b in range(batch_size):
                img = fake_images[b]

                file_name = f"{global_idx}.jpg"
                save_path = os.path.join(sample_dir, file_name)

                save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
                global_idx += 1
```

In [37]:

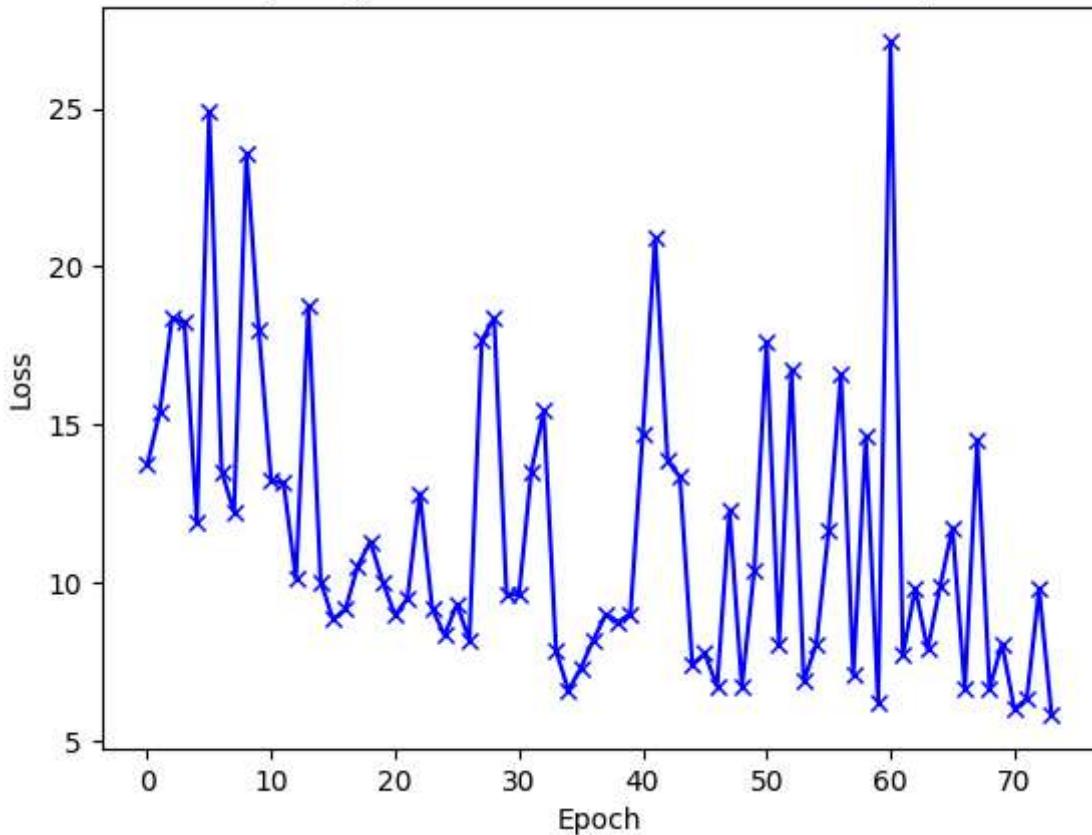
```
save_generated_tests(generator, test_dl)
```

In [38]:

```
def plot_g_losses(history):
    losses_g = [x for x in history['losses_g']]
    plt.plot(losses_g, '-bx')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Low Quality Blur - Generator Loss vs No. of epochs')
```

```
plot_g_losses(checkpoint)
```

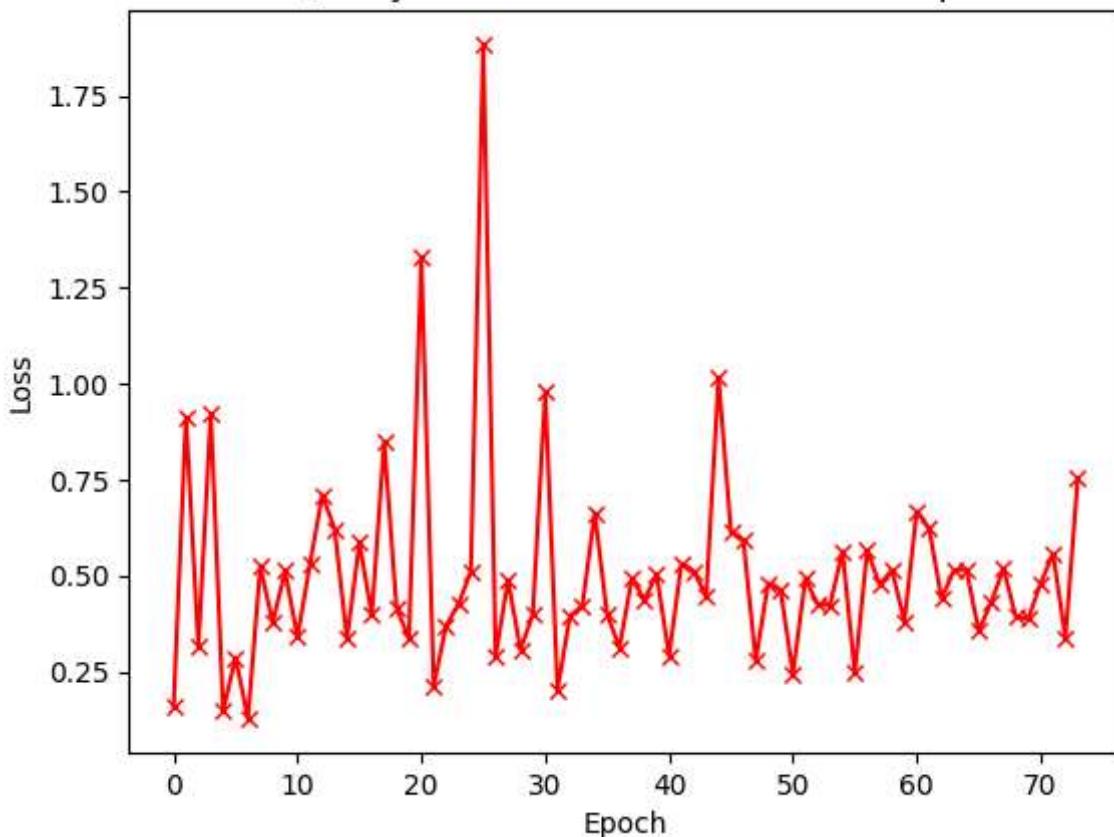
Low Quality Blur - Generator Loss vs No. of epochs



```
In [39]: def plot_d_losses(history):
    losses_d = [x for x in history['losses_d']]
    plt.plot(losses_d, '-rx')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Low Quality - Discriminator Loss vs No. of epochs')

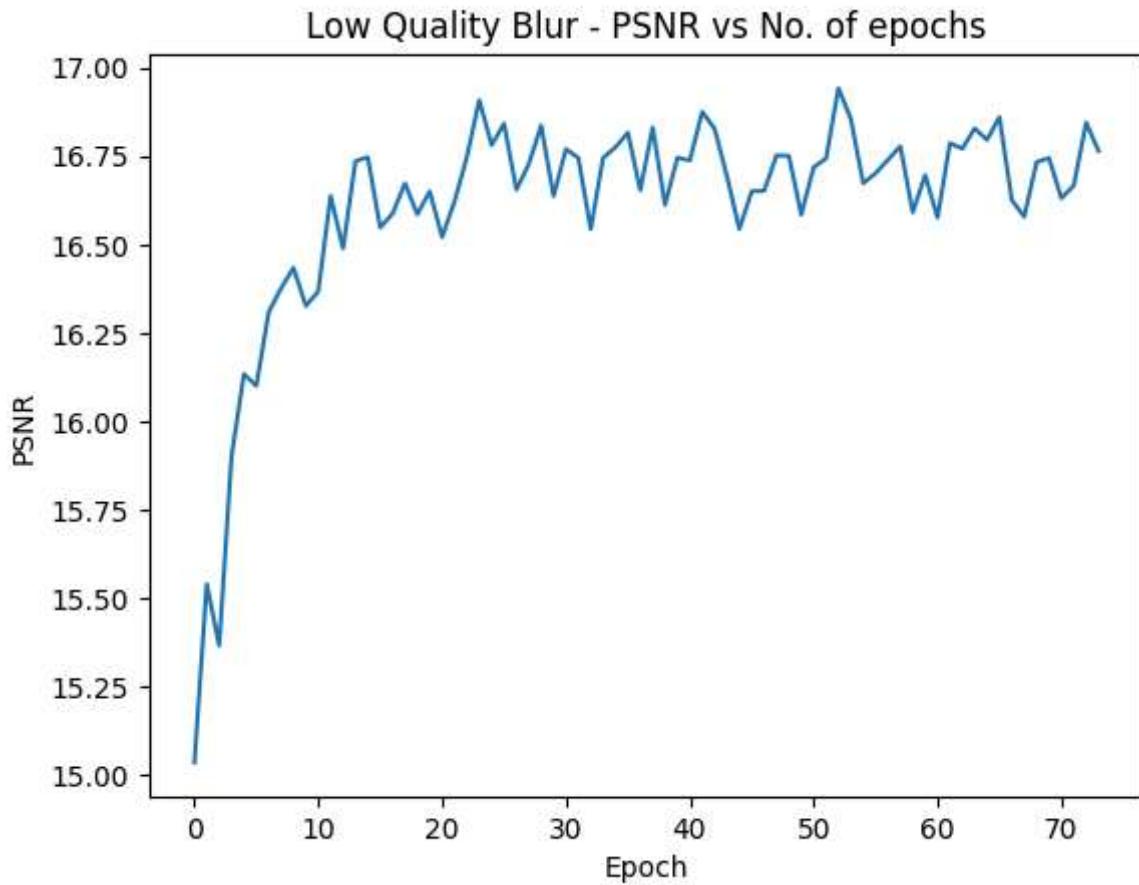
plot_d_losses(checkpoint)
```

Low Quality - Discriminator Loss vs No. of epochs



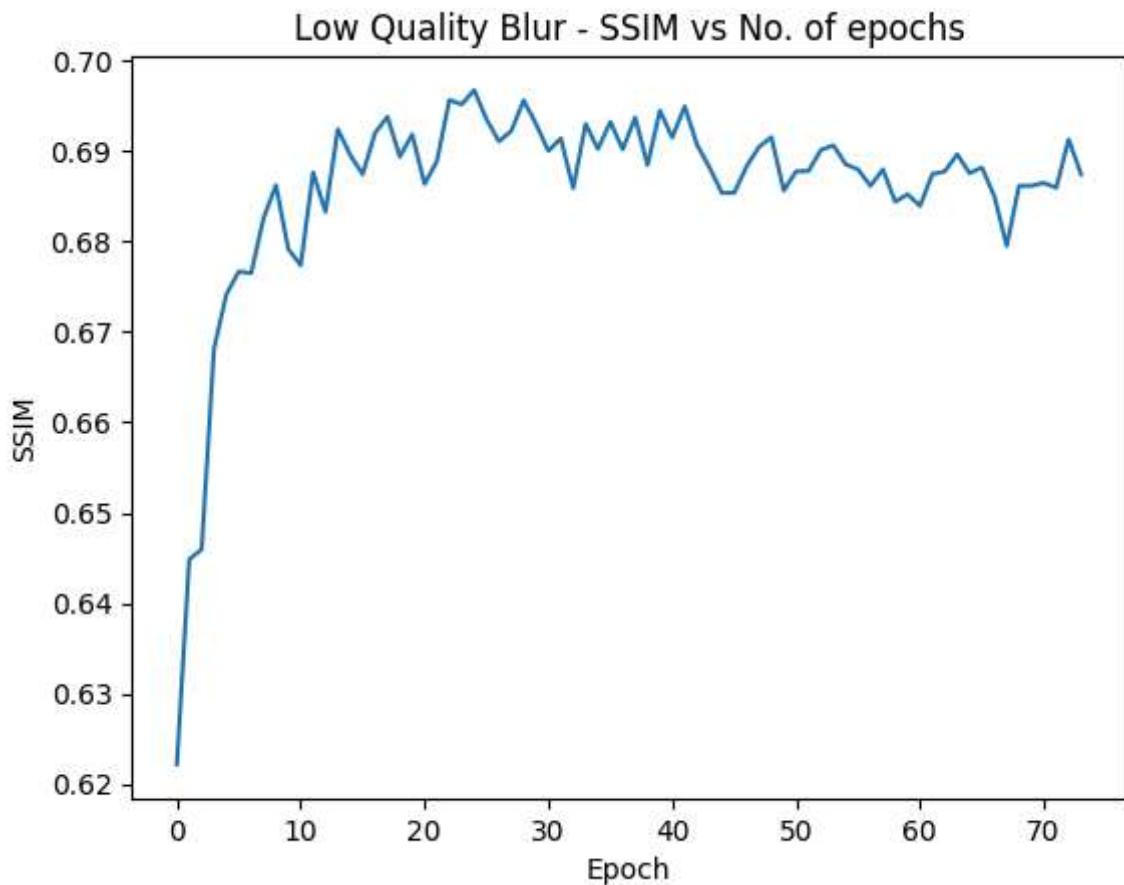
```
In [34]: def plot_psnrs(history):
    psnrs = [x for x in history['psnrs']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('PSNR')
    plt.title('Low Quality Blur - PSNR vs No. of epochs')

plot_psnrs(checkpoint)
```



```
In [36]: def plot_ssims(history):
    psnrs = [x for x in history['ssims']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('SSIM')
    plt.title('Low Quality Blur - SSIM vs No. of epochs')

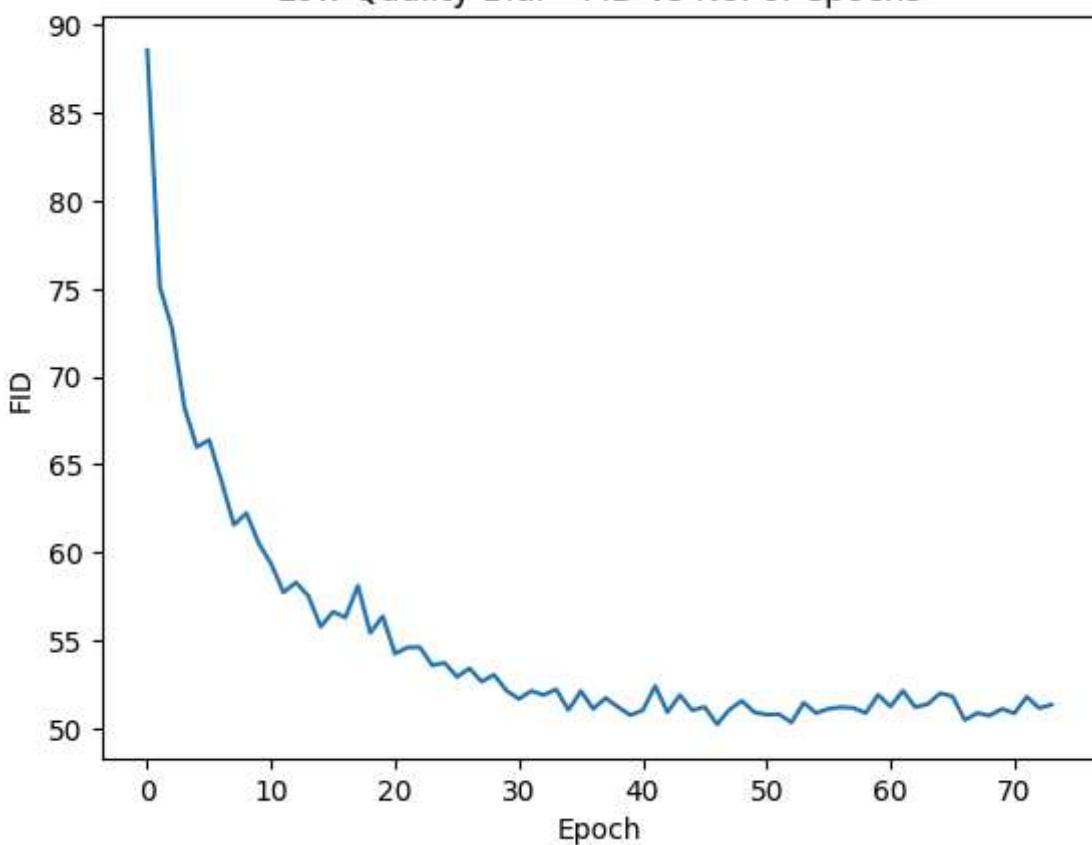
plot_ssims(checkpoint)
```



```
In [37]: def plot_fids(history):
    psnrs = [x for x in history['fids']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('FID')
    plt.title('Low Quality Blur - FID vs No. of epochs')

plot_fids(checkpoint)
```

Low Quality Blur - FID vs No. of epochs



In []: