

```
In [1]: import os
import sys
from tqdm import tqdm
import numpy as np
import torch
import torch.nn as nn
from torch.optim.lr_scheduler import ReduceLROnPlateau
from torch import optim
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
from torchvision.utils import make_grid
import torchvision.utils as vutils
import torchvision.transforms as T
import torchvision.transforms.functional as TF
import torch.nn.functional as F
from torchvision.utils import save_image
from PIL import Image
import matplotlib.pyplot as plt
%matplotlib inline
from paddleocr import PaddleOCR, draw_ocr
import paddle
from Levenshtein import distance as levenshtein_distance
from pytorch_msssim import ssim as ssim_fn
from lpips import LPIPS
from torchmetrics.image.psnr import PeakSignalNoiseRatio
from torchmetrics.image.ssim import StructuralSimilarityIndexMeasure
from torchmetrics.image.fid import FrechetInceptionDistance
```

```
In [2]: sys.path.append(os.path.abspath('../..../src/dataset'))
from paired_image_dataset import PairedImageDataset

DATA_DIR = os.path.join('..', '..', '..', 'data')
print(os.listdir(DATA_DIR))

['h_blur', 'low_light', 'low_qual', 'test', 'v_blur']
```

```
In [3]: image_width = 256
image_height = 128
batch_size = 16
stats = (0.5, 0.5, 0.5), (0.5, 0.5, 0.5)
```

```
In [4]: train_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)])

valid_transform = T.Compose([
    T.Resize((image_height, image_width)),
    T.ToTensor(),
    T.Normalize(*stats)
])
```

```
In [60]: train_ds = PairedImageDataset(blur_dir=os.path.join(DATA_DIR, 'v_blur', 'train', 'v',
normal_dir=os.path.join(DATA_DIR, 'v_blur', 'train',
```

```

        transform=train_transform)

valid_ds = PairedImageDataset(blur_dir=os.path.join(DATA_DIR, 'v.blur', 'valid', 'v',
                                                    normal_dir=os.path.join(DATA_DIR, 'v.blur', 'valid',
                                                    transform=valid_transform))

test_ds = ImageFolder(os.path.join(DATA_DIR, 'v.blur', 'test'), transform=valid_tra

```

In [61]:

```

train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=4)
valid_dl = DataLoader(valid_ds, batch_size=batch_size*2, shuffle=True, num_workers=4)
test_dl = DataLoader(test_ds, batch_size=batch_size//2, shuffle=False, num_workers=2)

```

Helper functions

In [62]:

```

def denorm(img_tensors):
    return img_tensors * stats[1][0] + stats[0][0]

def show_images(input_images, target_images, nmax=16):
    input_images = denorm(input_images.detach()[:nmax])
    target_images = denorm(target_images.detach()[:nmax])

    # Combine into a single tensor for visualization
    combined = torch.cat((input_images, target_images), 0)
    grid = make_grid(combined, nrow=nmax)

    plt.figure(figsize=(nmax, 4))
    plt.imshow(grid.permute(1, 2, 0))
    plt.axis("off")
    plt.title("Top: Horizontal Blur / Blurred | Bottom: Normal / Target")
    plt.show()

def show_batch(dl, nmax=16):
    for input_batch, target_batch in dl:
        show_images(input_batch, target_batch, nmax)
        break

```

In [63]:

```

def show_paired_samples(dataset, num_samples=4):
    fig, axes = plt.subplots(num_samples, 2, figsize=(5, 2 * num_samples))

    for i in range(num_samples):
        input_img, target_img = dataset[i]

        input_img = denorm(input_img)
        target_img = denorm(target_img)

        if isinstance(input_img, torch.Tensor):
            input_img = input_img.permute(1, 2, 0).numpy()
            target_img = target_img.permute(1, 2, 0).numpy()

        axes[i, 0].imshow(input_img)
        axes[i, 0].set_title("Horizontal Blur (Input)")
        axes[i, 0].axis("off")

```

```
axes[i, 1].imshow(target_img)
axes[i, 1].set_title("Normal (Target)")
axes[i, 1].axis("off")

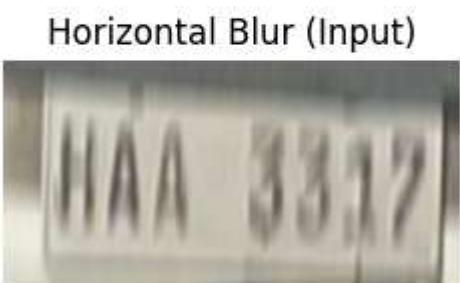
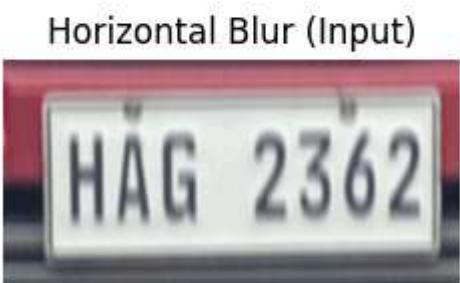
plt.tight_layout()
plt.show()

def show_generated_images(images, nrow=4):
    images = denorm(images.detach().cpu())

    # Make grid
    grid_img = vutils.make_grid(images, nrow=nrow, normalize=False)

    # Show image
    plt.figure(figsize=(nrow * 2, 2 * (len(images) // nrow)))
    plt.axis('off')
    plt.imshow(grid_img.permute(1, 2, 0))
    plt.show()
```

```
In [64]: show_paired_samples(train_ds)
```



```
In [65]: show_batch(valid_dl, nmax=14)
```

Top: Horizontal Blur / Blurred | Bottom: Normal / Target



```
In [66]: def print_images(image_tensor, num_images):
```

```
    images = denorm(image_tensor)
    images = images.detach().cpu()
    image_grid = make_grid(images[:num_images], nrow=5)
```

```
plt.imshow(image_grid.permute(1, 2, 0).squeeze())
plt.show()
```

```
In [67]: def save_samples(generator, test_dl, epoch, sample_dir='./generated_test_outputs'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)

            for b in range(batch_size):
                img = fake_images[b]

                # create folder per image
                folder_path = os.path.join(sample_dir, str(global_idx))
                os.makedirs(folder_path, exist_ok=True)

                file_name = f'{global_idx}_{epoch}epoch.jpg'
                save_path = os.path.join(folder_path, file_name)

                save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
                global_idx += 1
```

```
In [68]: def get_default_device():
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    def __init__(self, dl, device):
        self(dl = dl
        self.device = device

    def __iter__(self):
        for b in self(dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self(dl)
```

```
In [69]: device = get_default_device()
device
```

```
Out[69]: device(type='cuda')
```

```
In [70]: train_dl = DeviceDataLoader(train_dl, device)
valid_dl = DeviceDataLoader(valid_dl, device)
test_dl = DeviceDataLoader(test_dl, device)
```

Building the Model

Generator

```
In [16]: class ResConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3, dropout=0.0, batch_norm=False):
        super().__init__()
        padding = kernel_size // 2

        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size, padding=padding)
        self.bn1 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size, padding=padding)
        self.bn2 = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

        self.drop = nn.Dropout(p=dropout) if dropout > 0 else nn.Identity()

        self.shortcut = nn.Conv2d(in_channels, out_channels, kernel_size=1, padding=0)
        self.bn_sc = nn.BatchNorm2d(out_channels) if batch_norm else nn.Identity()

    def forward(self, x):
        residual = x

        out = self.conv1(x)
        out = self.bn1(out)
        out = F.relu(out)

        out = self.conv2(out)
        out = self.bn2(out)
        out = self.drop(out)

        shortcut = self.shortcut(residual)
        shortcut = self.bn_sc(shortcut)

        out += shortcut
        out = F.relu(out)
        return out

class GatingSignal(nn.Module):
    def __init__(self, in_channels, out_channels, batch_norm=False):
        super().__init__()
```

```

        layers = [nn.Conv2d(in_channels, out_channels, kernel_size=1, padding=0)]
        if batch_norm:
            layers.append(nn.BatchNorm2d(out_channels))
        layers.append(nn.ReLU(inplace=True))
        self.gate = nn.Sequential(*layers)

    def forward(self, x):
        return self.gate(x)

class AttentionBlock(nn.Module):
    def __init__(self, in_channels, gating_channels, inter_channels):
        super().__init__()

        self.theta_x = nn.Conv2d(in_channels, inter_channels, kernel_size=2, stride=2)
        self.phi_g = nn.Conv2d(gating_channels, inter_channels, kernel_size=1, padding=0)

        self.upsample_g = nn.ConvTranspose2d(inter_channels, inter_channels, kernel_size=2, stride=2)
        self.combine = nn.Sequential(
            nn.ReLU(inplace=True),
            nn.Conv2d(inter_channels, 1, kernel_size=1),
            nn.Sigmoid())

        self.final_conv = nn.Sequential(
            nn.Conv2d(in_channels, in_channels, kernel_size=1),
            nn.BatchNorm2d(in_channels))

    def forward(self, x, g):
        # x: skip connection feature map
        # g: gating signal (decoder feature map)
        theta_x = self.theta_x(x)
        phi_g = self.phi_g(g)

        # Upsample gating to match theta_x
        upsample_g = self.upsample_g(phi_g)

        if upsample_g.shape != theta_x.shape:
            upsample_g = F.interpolate(upsample_g, size=theta_x.shape[2:], mode='bilinear')

        psi = self.combine(theta_x + upsample_g)
        psi = F.interpolate(psi, size=x.shape[2:], mode='bilinear', align_corners=True)
        psi = psi.expand(-1, x.shape[1], -1, -1)

        y = x * psi
        return self.final_conv(y)

class AttentionResUNetGenerator(nn.Module):
    def __init__(self, in_channels=3, out_channels=3, base_filters=64, dropout=0.0,
                 super().__init__()

        # Encoder
        self.down1 = ResConvBlock(in_channels, base_filters, dropout=dropout, batch_norm=True)
        self.pool1 = nn.MaxPool2d(2)
        self.down2 = ResConvBlock(base_filters, base_filters * 2, dropout=dropout, batch_norm=True)
        self.pool2 = nn.MaxPool2d(2)
        self.down3 = ResConvBlock(base_filters * 2, base_filters * 4, dropout=dropout, batch_norm=True)
        self.pool3 = nn.MaxPool2d(2)
        self.up1 = ResConvBlock(base_filters * 4, base_filters * 2, dropout=dropout, batch_norm=True)
        self.up2 = ResConvBlock(base_filters * 2, base_filters, dropout=dropout, batch_norm=True)
        self.up3 = ResConvBlock(base_filters, base_filters, dropout=dropout, batch_norm=True)
        self.out = nn.Conv2d(base_filters, out_channels, kernel_size=1)

```

```

        self.pool3 = nn.MaxPool2d(2)
        self.down4 = ResConvBlock(base_filters * 4, base_filters * 8, dropout=dropout)
        self.pool4 = nn.MaxPool2d(2)

    # Bottleneck
    self.bottleneck = ResConvBlock(base_filters * 8, base_filters * 16, dropout=dropout)

    # Gating and Attention
    self.gate4 = GatingSignal(base_filters * 16, base_filters * 8, batch_norm)
    self.attn4 = AttentionBlock(base_filters * 8, base_filters * 8, base_filters * 16)

    self.gate3 = GatingSignal(base_filters * 8, base_filters * 4, batch_norm) #
    self.attn3 = AttentionBlock(base_filters * 4, base_filters * 4, base_filters * 8)

    self.gate2 = GatingSignal(base_filters * 4, base_filters * 2, batch_norm) #
    self.attn2 = AttentionBlock(base_filters * 2, base_filters * 2, base_filters * 4)

    self.gate1 = GatingSignal(base_filters * 2, base_filters, batch_norm) # 128
    self.attn1 = AttentionBlock(base_filters, base_filters, base_filters) # 64

    # Decoder
    self.up4 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec4 = ResConvBlock(base_filters * 16 + base_filters * 8, base_filters * 8)

    self.up3 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec3 = ResConvBlock(base_filters * 8 + base_filters * 4, base_filters * 4)

    self.up2 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec2 = ResConvBlock(base_filters * 4 + base_filters * 2, base_filters * 2)

    self.up1 = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
    self.dec1 = ResConvBlock(base_filters * 2 + base_filters, base_filters, dropout=dropout)

    # Output Layer
    self.final_conv = nn.Sequential(
        nn.Conv2d(base_filters, out_channels, kernel_size=1),
        nn.Tanh())

def forward(self, x):
    d1 = self.down1(x)
    p1 = self.pool1(d1)
    d2 = self.down2(p1)
    p2 = self.pool2(d2)
    d3 = self.down3(p2)
    p3 = self.pool3(d3)
    d4 = self.down4(p3)
    p4 = self.pool4(d4)

    bn = self.bottleneck(p4)

    g4 = self.gate4(bn)
    a4 = self.attn4(d4, g4)
    u4 = self.up4(bn)
    u4 = torch.cat([u4, a4], dim=1)
    d5 = self.dec4(u4)

```

```

        g3 = self.gate3(d5)
        a3 = self.attn3(d3, g3)
        u3 = self.up3(d5)
        u3 = torch.cat([u3, a3], dim=1)
        d6 = self.dec3(u3)

        g2 = self.gate2(d6)
        a2 = self.attn2(d2, g2)
        u2 = self.up2(d6)
        u2 = torch.cat([u2, a2], dim=1)
        d7 = self.dec2(u2)

        g1 = self.gate1(d7)
        a1 = self.attn1(d1, g1)
        u1 = self.up1(d7)
        u1 = torch.cat([u1, a1], dim=1)
        d8 = self.dec1(u1)

    return self.final_conv(d8)

```

In [17]:

```

generator = AttentionResUNetGenerator()
x = torch.randn(4, 3, 128, 256) # Batch of horizontal motion blurred images
y = generator(x) # Deblurred output
print(y.shape) # Should be (4, 3, 128, 256)

```

torch.Size([4, 3, 128, 256])

Discriminator

In [18]:

```

class PatchDiscriminator(nn.Module):
    def __init__(self, in_channels=3, base_filters=64):
        super().__init__()
        layers = [
            nn.Conv2d(in_channels, base_filters, kernel_size=4, stride=2, padding=1,
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters, base_filters * 2, kernel_size=4, stride=2, padding=1,
                     nn.BatchNorm2d(base_filters * 2),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 2, base_filters * 4, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 4),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 4, base_filters * 8, kernel_size=4, stride=2,
                     nn.BatchNorm2d(base_filters * 8),
                     nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(base_filters * 8, 1, kernel_size=4, stride=2, padding=1),
        ]
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)

```

```
In [19]: image1 = torch.rand((1, 3, 128, 256))

discriminator = PatchDiscriminator()
output = discriminator(image1)
print(output.shape)

torch.Size([1, 1, 4, 8])
```

Model Parameters

```
In [20]: total_params = sum(p.numel() for p in discriminator.parameters())
total_params
```

```
Out[20]: 2766529
```

```
In [21]: discriminator = to_device(discriminator, device)
generator = to_device(generator, device)
```

Training

Loss Functions

```
In [22]: comparison_loss = nn.BCEWithLogitsLoss()
L1_loss_fn = nn.L1Loss()
```

PaddleOCR

```
In [23]: ocr_model = PaddleOCR(use_angle_cls=True, lang='en', use_gpu=True, show_log=False)

def tensor_to_bgr_numpy(tensor_img, mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5)):
    if tensor_img.dim() == 4:
        tensor_img = tensor_img[0]

    if tensor_img.shape[0] != 3:
        raise ValueError(f"Expected 3 channels, got shape {tensor_img.shape}")

    unnorm = torch.zeros_like(tensor_img)
    for c in range(3):
        unnorm[c] = tensor_img[c] * float(std[c]) + float(mean[c])
    tensor_img = unnorm.clamp(0, 1)

    pil_img = TF.to_pil_image(tensor_img.cpu())
    np_img = np.array(pil_img)[:, :, ::-1]

    return np_img
```

```
[2025/09/10 17:28:52] ppocr WARNING: The first GPU is used for inference by default,  
GPU ID: 0  
[2025/09/10 17:28:53] ppocr WARNING: The first GPU is used for inference by default,  
GPU ID: 0  
[2025/09/10 17:28:56] ppocr WARNING: The first GPU is used for inference by default,  
GPU ID: 0
```

Discriminator Training Function

```
In [24]: def train_discriminator(discriminator, generator, inputs, targets, d_opt):  
    discriminator.train()  
  
    real_images = targets  
    fake_images = generator(inputs).detach()  
  
    pred_real = discriminator(real_images)  
    pred_fake = discriminator(fake_images)  
  
    real_labels = torch.ones_like(pred_real)  
    fake_labels = torch.zeros_like(pred_fake)  
  
    real_score = torch.sigmoid(pred_real).mean().item()  
    fake_score = torch.sigmoid(pred_fake).mean().item()  
  
    loss_real = comparison_loss(pred_real, real_labels)  
    loss_fake = comparison_loss(pred_fake, fake_labels)  
  
    d_loss = (loss_real + loss_fake) / 2  
  
    d_opt.zero_grad()  
    d_loss.backward()  
    d_opt.step()  
  
    return d_loss.item(), real_score, fake_score
```

Generator Training Function

```
In [25]: def ssim(pred, target):  
    return 1 - ssim_fn(pred, target, data_range=1.0, size_average=True)  
  
lpips_loss = LPIPS(net='vgg').to(device)  
lpips_loss.eval()  
for param in lpips_loss.parameters():  
    param.requires_grad = False  
  
def train_generator(generator, discriminator, inputs, targets, g_opt, adv_lambda=1.  
    generator.train()  
  
    fake_images = generator(inputs)  
    pred_fake = discriminator(fake_images)  
  
    real_labels = torch.ones_like(pred_fake)  
  
    adv_loss = comparison_loss(pred_fake, real_labels)
```

```

l1_loss = L1_loss_fn(fake_images, targets)
perceptual_loss = lpips_loss(fake_images.clamp(-1, 1), targets.clamp(-1, 1)).mse
ssim_loss = ssim(fake_images, targets)

# Text Loss
fake_np = tensor_to_bgr_numpy(fake_images)
real_np = tensor_to_bgr_numpy(targets)

with torch.no_grad():
    fake_text = ocr_model.ocr(fake_np, cls=False)[0]
    real_text = ocr_model.ocr(real_np, cls=False)[0]

fake_str = fake_text[0][1][0] if fake_text else ''
real_str = real_text[0][1][0] if real_text else ''

fake_str = fake_str.strip().replace(' ', '')
real_str = real_str.strip().replace(' ', '')

# Use normalized edit distance as loss
edit_dist = levenshtein_distance(fake_str, real_str)
norm_edit_dist = edit_dist / max(len(real_str), 1)
text_loss = torch.tensor(norm_edit_dist, device=device)

g_loss = (
    adv_loss * adv_lambda +
    l1_loss * l1_lambda +
    ssim_loss * ssim_lambda +
    perceptual_loss * perceptual_lambda +
    text_loss * text_lambda
)

g_opt.zero_grad()
g_loss.backward()
g_opt.step()

return g_loss.item(), fake_images, {
    "total_loss": g_loss.item(),
    "adv": adv_loss.item(),
    "l1": l1_loss.item(),
    "ssim": ssim_loss.item(),
    "perceptual": perceptual_loss.item(),
    "text": text_loss.item()
}

```

Setting up [LPIPS] perceptual loss: trunk [vgg], v[0.1], spatial [off]
Loading model from: C:\Thesis\LiPAD with Paddle\lipadvenv3.9\lib\site-packages\lpips\weights\v0.1\vgg.pth

```

In [26]: psnr_metric = PeakSignalNoiseRatio(data_range=1.0).to(device)
ssim_metric = StructuralSimilarityIndexMeasure(data_range=1.0).to(device)
fid_metric = FrechetInceptionDistance(feature=2048).to(device)

def to_uint8(tensor):
    # Clamp to [0, 1], scale to [0, 255], then convert to uint8
    tensor = torch.clamp(tensor, 0.0, 1.0)
    tensor = (tensor * 255.0).to(torch.uint8)

```

```

    return tensor

def evaluate_test(generator):
    raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl,
                                                               fig, axes = plt.subplots(1, 2, figsize=(10, 8)) # 1 row, 2 columns

    # Raw Test Set
    axes[0].imshow(raw_grid_img)
    axes[0].axis("off")
    axes[0].set_title("Raw Test Set")

    # Generated Test Set
    axes[1].imshow(gen_grid_img)
    axes[1].axis("off")
    axes[1].set_title("Generated Test Set")

    plt.tight_layout()
    plt.show()

```

In [27]:

```

def create_image_grid_from_loader(generator, dataloader, device, num_images=16, image_size=(64, 64)):
    generator.eval()
    generator.to(device)

    raw_images = []
    gen_images = []
    gen_count = 0
    raw_count = 0

    for batch in dataloader:
        # Assume batch is either just images or (images, labels)
        if isinstance(batch, (tuple, list)):
            inputs = batch[0]
        else:
            inputs = batch

        inputs = inputs.to(device)
        outputs = generator(inputs)

        for img in inputs:
            if denormalize:
                img = denorm(img)

            img_resized = TF.resize(img, image_size)
            raw_images.append(img_resized)

            raw_count += 1
            if raw_count >= num_images:
                break

        for out_img in outputs:
            if denormalize:
                out_img = denorm(out_img)

            out_img_resized = TF.resize(out_img, image_size)
            gen_images.append(out_img_resized)

```

```

        gen_count += 1
        if gen_count >= num_images:
            break

    if raw_count >= num_images:
        break

raw_grid = make_grid(raw_images, nrow=4)
gen_grid = make_grid(gen_images, nrow=4)

raw_ndarr = raw_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()
gen_ndarr = gen_grid.mul(255).byte().cpu().permute(1, 2, 0).numpy()

raw_grid_image = Image.fromarray(raw_ndarr)
gen_grid_image = Image.fromarray(gen_ndarr)

if save_path:
    gen_grid_image.save(save_path)
    print(f"Saved image grid to {save_path}")

return raw_grid_image, gen_grid_image

```

In [28]:

```

@torch.no_grad()
def evaluate_generator(generator, valid_dl, epoch):
    generator.eval()
    psnr_vals = []
    ssim_vals = []
    fid_metric.reset()
    shown = False

    for val_inputs, val_targets in valid_dl:
        val_inputs = val_inputs.to(device)
        val_targets = val_targets.to(device)

        val_outputs = generator(val_inputs)

        if not shown:
            print('Validation Data')
            print_images(val_inputs, 5)
            print_images(val_outputs, 5)
            print_images(val_targets, 5)
            shown = True

        # Convert to uint8 for FID
        gen_uint8 = to_uint8(val_outputs)
        target_uint8 = to_uint8(val_targets)

        fid_metric.update(gen_uint8, real=False)
        fid_metric.update(target_uint8, real=True)

        # Compute PSNR/SSIM
        psnr = psnr_metric(val_outputs, val_targets).item()
        ssim = ssim_metric(val_outputs, val_targets).item()
        psnr_vals.append(psnr)
        ssim_vals.append(ssim)

```

```

        mean_psnr = sum(psnr_vals) / len(psnr_vals)
        mean_ssim = sum(ssim_vals) / len(ssim_vals)
        fid = fid_metric.compute().item()

        print(f"Epoch [{epoch+1}/{epochs}] - Mean PSNR: {mean_psnr:.4f}, Mean SSIM: {mean_ssim:.4f}, FID: {fid:.4f}")

    evaluate_test(generator)
    return mean_psnr, mean_ssim, fid

def get_lr(optimizer):
    return optimizer.param_groups[0]['lr']

def fit(generator, discriminator, epochs, lr, adv_lambda=1.0, l1_lambda=100.0, ssim_lambda=1.0,
        perceptual_lambda=5.0, text_lambda=7.5, g_opt=None, d_opt=None, checkpoint=None):
    torch.cuda.empty_cache()

    losses_g = []
    losses_d = []
    real_scores = []
    fake_scores = []
    psnrs = []
    ssims = []
    fids = []
    g_losses_trace = []
    g_lrs = []
    d_lrs = []

    # Create optimizers
    if d_opt is None:
        d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
    if g_opt is None:
        g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))

    # Load checkpoint state if resuming
    if checkpoint is not None:
        print("Resuming training from checkpoint...")
        d_opt.load_state_dict(checkpoint['opt_d'])
        g_opt.load_state_dict(checkpoint['opt_g'])

    g_scheduler = ReduceLROnPlateau(g_opt, mode='min', factor=0.5, patience=10, verbose=False)
    d_scheduler = ReduceLROnPlateau(d_opt, mode='min', factor=0.5, patience=12, verbose=False)

    # Train
    for epoch in range(epochs):
        torch.cuda.empty_cache()
        generator.train()
        discriminator.train()
        for inputs, targets in tqdm(train_dl):
            inputs = inputs.to(device)
            targets = targets.to(device)

            # Train discriminator
            d_loss, real_score, fake_score = train_discriminator(discriminator, generator, inputs, targets, device, g_opt, d_opt, adv_lambda, l1_lambda, ssim_lambda, perceptual_lambda, text_lambda)
            losses_d.append(d_loss.item())
            real_scores.append(real_score.item())
            fake_scores.append(fake_score.item())
            psnrs.append(psnr_fn(inputs, targets).item())
            ssims.append(ssim_fn(inputs, targets).item())
            fids.append(fid_fn(generator, discriminator, inputs, device).item())
            g_losses_trace.append(g_opt.state_dict()['state'][0]['step'])

        # Train generator
        g_opt.zero_grad()
        generator.train()
        for inputs, targets in tqdm(train_dl):
            inputs = inputs.to(device)
            targets = targets.to(device)

            # Train generator
            g_loss = train_generator(generator, discriminator, inputs, targets, device, g_opt, d_opt, adv_lambda, l1_lambda, ssim_lambda, perceptual_lambda, text_lambda)
            losses_g.append(g_loss.item())
            g_lrs.append(g_opt.state_dict()['state'][0]['step'])

        # Save checkpoint
        if (epoch + 1) % 5 == 0:
            checkpoint = {
                'epoch': epoch + 1,
                'opt_d': d_opt.state_dict(),
                'opt_g': g_opt.state_dict(),
                'losses_d': losses_d,
                'real_scores': real_scores,
                'fake_scores': fake_scores,
                'psnrs': psnrs,
                'ssims': ssims,
                'fids': fids,
                'g_losses_trace': g_losses_trace,
                'g_lrs': g_lrs,
                'd_lrs': d_lrs
            }
            torch.save(checkpoint, f'checkpoint_{epoch+1}.pt')

    # Evaluate test set
    mean_psnr, mean_ssim, fid = evaluate_test(generator)
    print(f"Final Mean PSNR: {mean_psnr:.4f}, Mean SSIM: {mean_ssim:.4f}, FID: {fid:.4f}")

```

```

# Train generator
g_loss, fake_images, g_loss_trace = train_generator(generator, discriminator,
                                                    g1_lambda, ssim_lambda, perceptron)

print('Training Data')
print_images(inputs, 5)
print_images(fake_images, 5)
print_images(targets, 5)

if (epoch + 1) % 5 == 0:
    torch.save(generator.state_dict(), f'[RAU-NET OCRDET LR]generator-Datas')
    torch.save(discriminator.state_dict(), f'[RAU-NET OCRDET LR]discriminator-Datas')
    torch.save({
        'generator': generator.state_dict(),
        'discriminator': discriminator.state_dict(),
        'opt_g': g_opt.state_dict(),
        'opt_d': d_opt.state_dict(),
        'losses_g': losses_g,
        'losses_d': losses_d,
        'real_score': real_scores,
        'fake_score': fake_scores,
        'psnrs': psnrs,
        'ssims': ssims,
        'fids': fids,
        'g_lrs': g_lrs,
        'd_lrs': d_lrs,
        'epoch': epoch + 1
    }, f'[RAU-NET OCRDET LR]Checkpoint-Logs-Datasetv2-{epoch + 1} epoch.pt')

# Print progress
try:
    curr_glr = get_lr(g_opt)
    curr_dlr = get_lr(d_opt)

    print(f'Epoch [{epoch}/{epochs}], loss_g: {g_loss:.4f}, loss_d: {d_loss:.4f}, real_score: {real_scores:.4f}')
    print(f"Current LRs - Generator: {curr_glr:.6f}, Discriminator: {curr_dlr:.6f}")
except Exception as err:
    print('Error:', str(err))

mean_psnr, mean_ssims, fid = evaluate_generator(generator, valid_dl, epoch)

g_scheduler.step(g_loss)
d_scheduler.step(d_loss)

# Record Losses and scores
losses_g.append(g_loss)
losses_d.append(d_loss)
real_scores.append(real_score)
fake_scores.append(fake_score)
psnrs.append(mean_psnr)
ssims.append(mean_ssims)
fids.append(fid)
g_losses_trace.append(g_loss_trace)
g_lrs.append(curr_glr)
d_lrs.append(curr_dlr)

```

```
    save_samples(generator, test_dl, epoch + 1, sample_dir='./visualization-ocr')

    return losses_g, losses_d, real_scores, fake_scores, psnrs, ssims, fids, g_loss
```

Hyperparameter Configurations

```
In [29]: adv_lambda = 0.25 # Change to 0.25
    l1_lambda = 100
    ssim_lambda = 5 # Change to 5
    perceptual_lambda = 5.0
    text_lambda = 10 # Change to 10

    lr = 0.0002
    beta1 = 0.5
    beta2 = 0.999

    epochs = 200
```

```
In [30]: g_opt = optim.Adam(generator.parameters(), lr=lr, betas=(beta1, beta2))
    d_opt = optim.Adam(discriminator.parameters(), lr=lr, betas=(beta1, beta2))
```

```
In [31]: history = []
```

```
In [ ]: %time
    history += fit(generator,
                    discriminator,
                    epochs,
                    lr,
                    adv_lambda=adv_lambda,
                    l1_lambda=l1_lambda,
                    ssim_lambda=ssim_lambda,
                    perceptual_lambda=perceptual_lambda,
                    text_lambda=text_lambda,
                    g_opt=g_opt,
                    d_opt=d_opt)
```

100% |
1250/1250 [18:54<00:00, 1.10it/s]
Training Data





Epoch [1/200], loss_g: 21.7972, loss_d: 0.3590, real_score: 0.9407, fake_score: 0.4576

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [1/200] - Mean PSNR: 12.7310, Mean SSIM: 0.5079, FID: 98.25654602050781

Raw Test Set

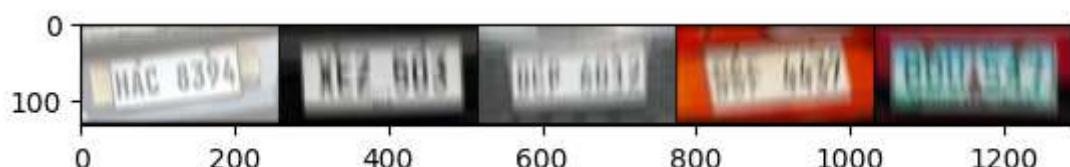
Generated Test Set



100% |

1250/1250 [19:02<00:00, 1.09it/s]

Training Data



Epoch [2/200], loss_g: 15.5627, loss_d: 0.4916, real_score: 0.9972, fake_score: 0.59
98

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [2/200] - Mean PSNR: 13.7796, Mean SSIM: 0.5707, FID: 86.4118423461914

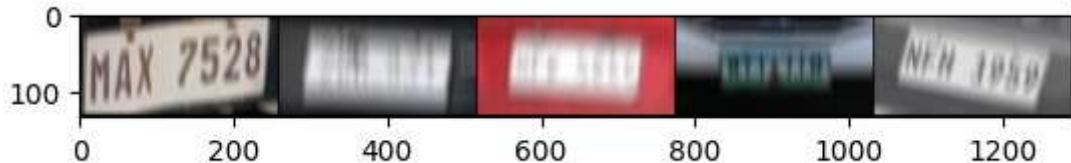
Raw Test Set

Generated Test Set



100% |
1250/1250 [18:59<00:00, 1.10it/s]

Training Data



Epoch [3/200], loss_g: 15.1074, loss_d: 1.3616, real_score: 0.0743, fake_score: 0.02
42

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [3/200] - Mean PSNR: 14.5194, Mean SSIM: 0.6036, FID: 80.87458038330078

Raw Test Set

Generated Test Set



100% |
1250/1250 [19:00<00:00, 1.10it/s]

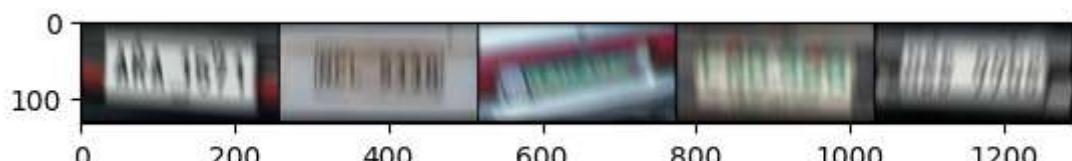
Training Data



Epoch [4/200], loss_g: 15.0530, loss_d: 0.2520, real_score: 0.9342, fake_score: 0.34
37

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [4/200] - Mean PSNR: 14.8999, Mean SSIM: 0.6257, FID: 77.60379028320312

Raw Test Set

Generated Test Set



100%

1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [5/200], loss_g: 15.6081, loss_d: 0.4763, real_score: 0.4459, fake_score: 0.10

43

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [5/200] - Mean PSNR: 15.3738, Mean SSIM: 0.6433, FID: 75.07611846923828

Raw Test Set

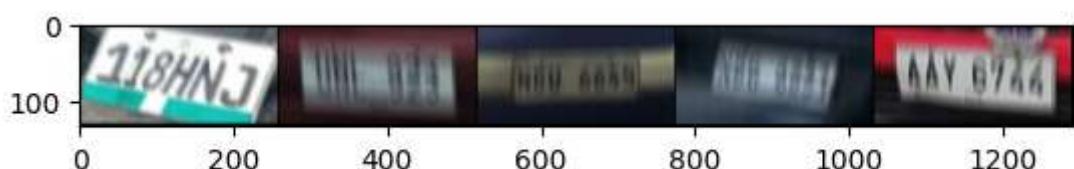
Generated Test Set



100% |

1250/1250 [18:55<00:00, 1.10it/s]

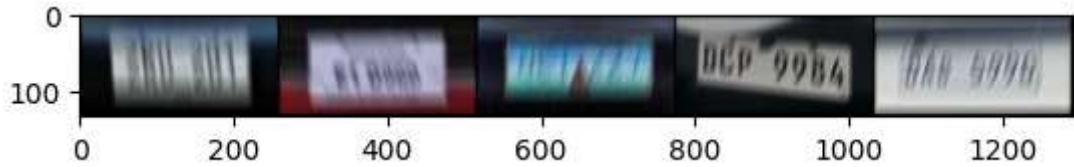
Training Data



Epoch [6/200], loss_g: 12.2404, loss_d: 0.7754, real_score: 0.2958, fake_score: 0.2357

Current LRs – Generator: 0.000200, Discriminator: 0.000200

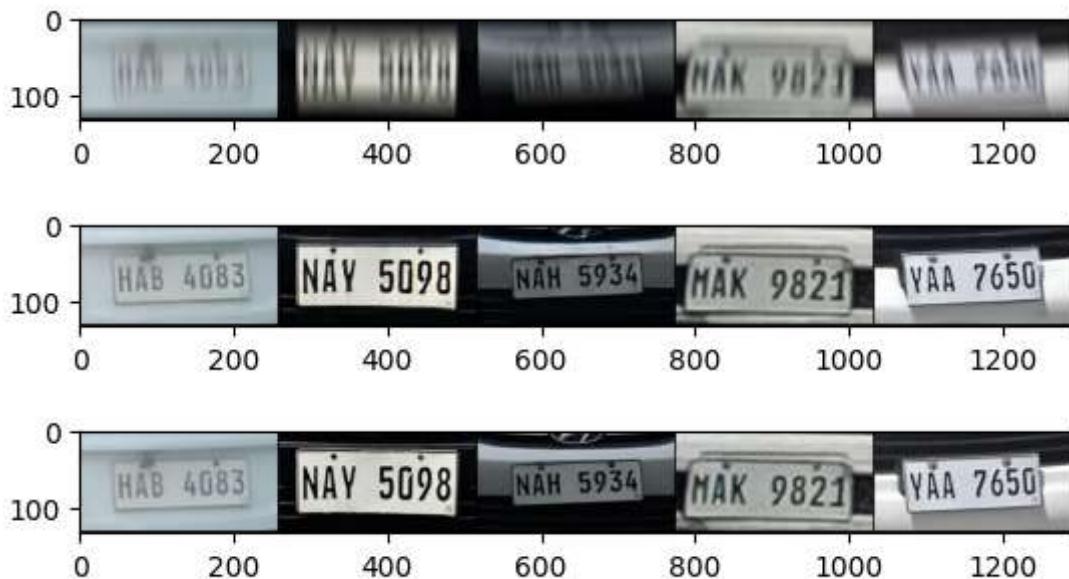
Validation Data



Epoch [6/200] - Mean PSNR: 15.7332, Mean SSIM: 0.6593, FID: 74.7110595703125



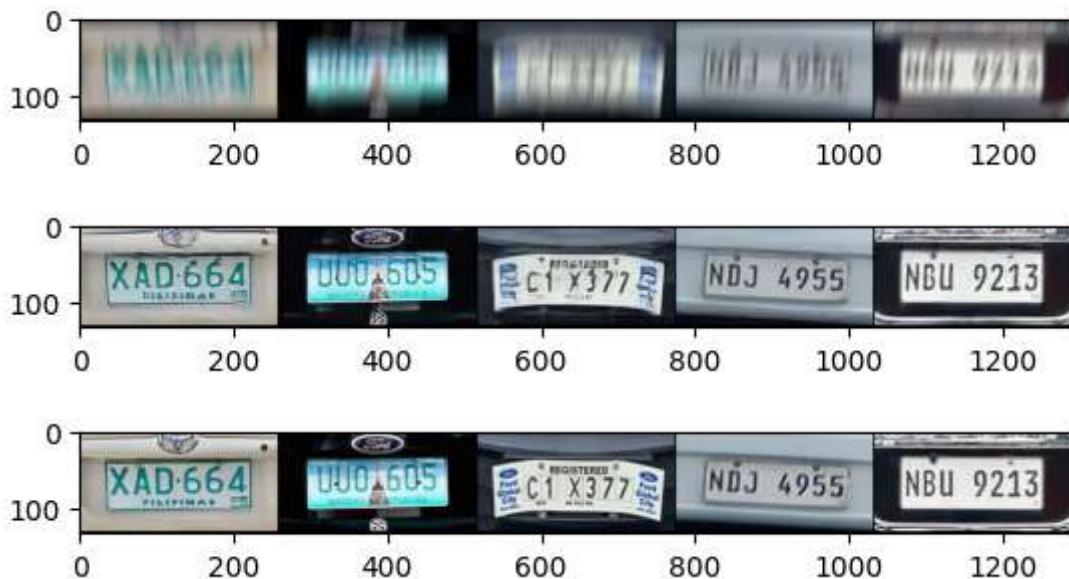
100% |
1250/1250 [18:55<00:00, 1.10it/s]
Training Data



Epoch [7/200], loss_g: 9.9165, loss_d: 0.2968, real_score: 0.7001, fake_score: 0.193
8

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [7/200] - Mean PSNR: 15.7611, Mean SSIM: 0.6597, FID: 73.96260833740234



100% |
1250/1250 [18:53<00:00, 1.10it/s]
Training Data



Epoch [8/200], loss_g: 11.6706, loss_d: 0.6853, real_score: 0.7248, fake_score: 0.6163

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [8/200] - Mean PSNR: 16.0421, Mean SSIM: 0.6659, FID: 70.47030639648438

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:53<00:00, 1.10it/s]

Training Data





Epoch [9/200], loss_g: 11.1505, loss_d: 0.3645, real_score: 0.9692, fake_score: 0.4768

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [9/200] - Mean PSNR: 16.1134, Mean SSIM: 0.6726, FID: 69.4925308227539

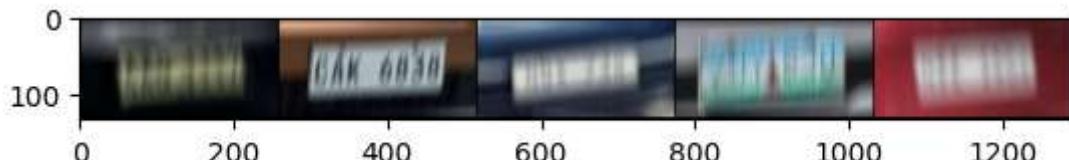
Raw Test Set

Generated Test Set



100% |
1250/1250 [18:55<00:00, 1.10it/s]

Training Data

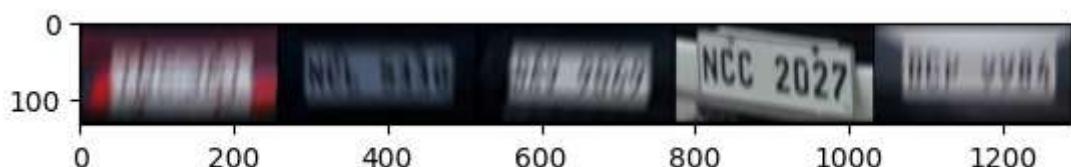




Epoch [10/200], loss_g: 13.0027, loss_d: 0.5434, real_score: 0.3987, fake_score: 0.0
805

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [10/200] - Mean PSNR: 16.4206, Mean SSIM: 0.6815, FID: 66.58281707763672

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:55<00:00, 1.10it/s]

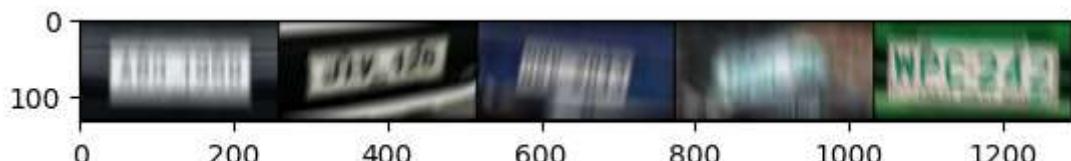
Training Data



Epoch [11/200], loss_g: 9.7814, loss_d: 0.3423, real_score: 0.9229, fake_score: 0.42
71

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [11/200] - Mean PSNR: 16.7512, Mean SSIM: 0.6970, FID: 65.14450073242188

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [12/200], loss_g: 8.4666, loss_d: 1.1367, real_score: 0.1356, fake_score: 0.10
16

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [12/200] - Mean PSNR: 16.6780, Mean SSIM: 0.6930, FID: 64.80681610107422

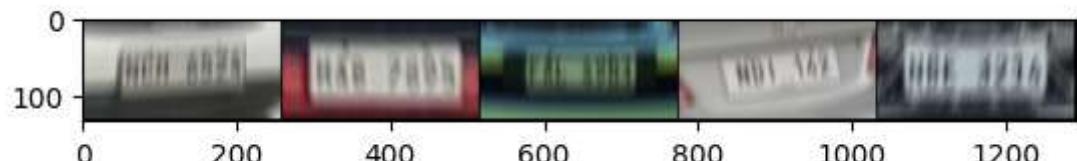
Raw Test Set

Generated Test Set



100% |
1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [13/200], loss_g: 10.5010, loss_d: 0.5308, real_score: 0.4480, fake_score: 0.1809

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



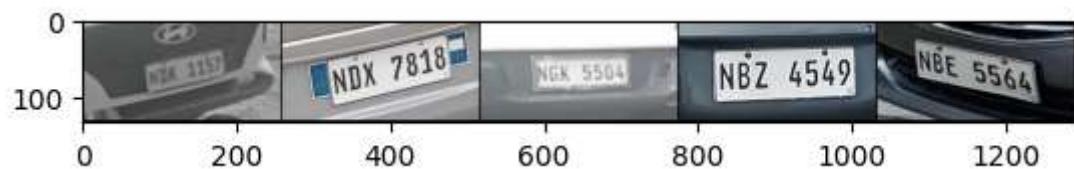
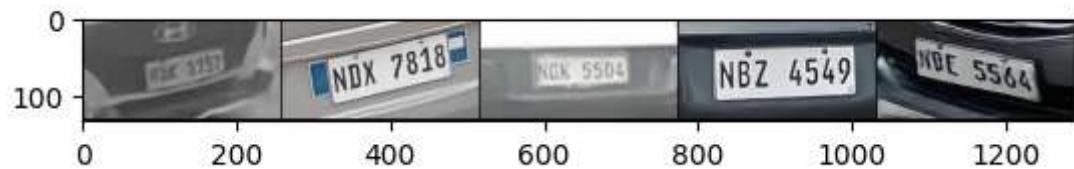
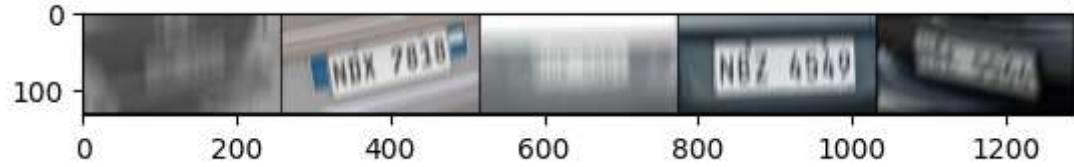


Epoch [13/200] - Mean PSNR: 17.0537, Mean SSIM: 0.7073, FID: 63.35651397705078
 Raw Test Set Generated Test Set



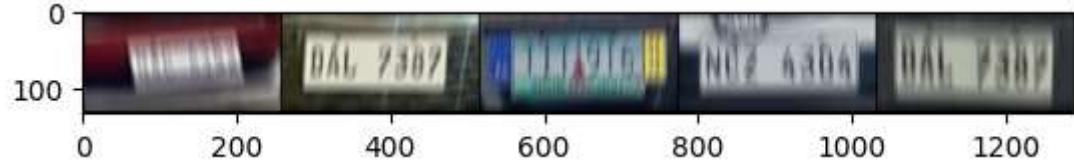
100% |
 1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [14/200], loss_g: 12.9691, loss_d: 0.2707, real_score: 0.6972, fake_score: 0.1432

Current LRs - Generator: 0.000200, Discriminator: 0.000200
 Validation Data





Epoch [14/200] - Mean PSNR: 16.6548, Mean SSIM: 0.6982, FID: 62.315162658691406

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:56<00:00, 1.10it/s]

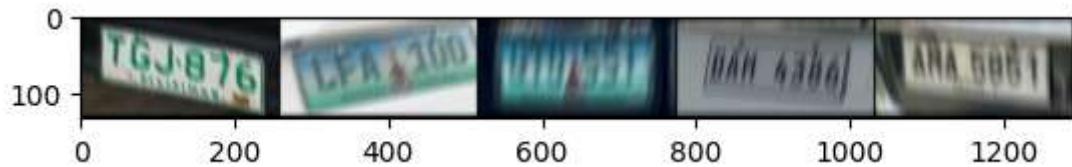
Training Data



Epoch [15/200], loss_g: 10.5756, loss_d: 0.3018, real_score: 0.9042, fake_score: 0.3701

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [15/200] - Mean PSNR: 16.8709, Mean SSIM: 0.7059, FID: 62.25211715698242



100% |
1250/1250 [18:57<00:00, 1.10it/s]
Training Data



Epoch [16/200], loss_g: 7.8930, loss_d: 0.5079, real_score: 0.4588, fake_score: 0.1655

Current LRs – Generator: 0.000200, Discriminator: 0.000200

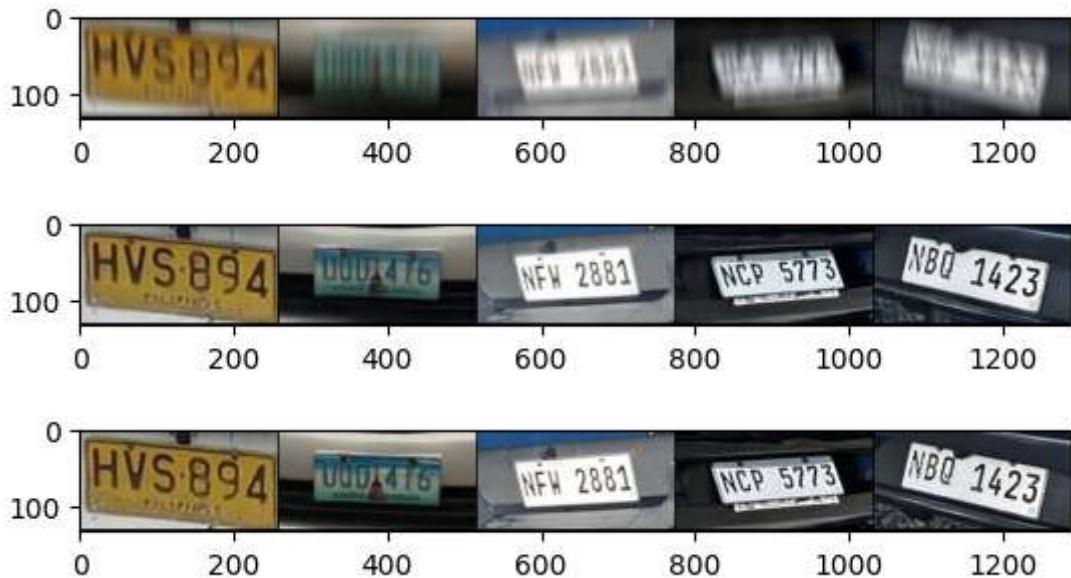
Validation Data



Epoch [16/200] - Mean PSNR: 17.1057, Mean SSIM: 0.7109, FID: 59.854759216308594



100% |
1250/1250 [18:57<00:00, 1.10it/s]
Training Data



Epoch [17/200], loss_g: 10.7309, loss_d: 0.2225, real_score: 0.9284, fake_score: 0.2959

Current LRs – Generator: 0.000200, Discriminator: 0.000200

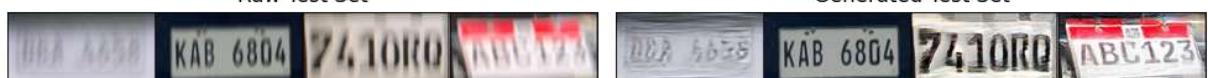
Validation Data



Epoch [17/200] - Mean PSNR: 17.0862, Mean SSIM: 0.7147, FID: 58.92013931274414

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:57<00:00, 1.10it/s]

Training Data





Epoch [18/200], loss_g: 8.8675, loss_d: 0.9229, real_score: 0.1941, fake_score: 0.0491

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [18/200] - Mean PSNR: 17.3312, Mean SSIM: 0.7169, FID: 58.23484420776367

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:57<00:00, 1.10it/s]

Training Data

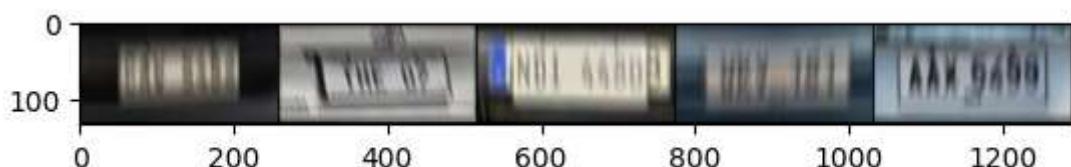




Epoch [19/200], loss_g: 7.1087, loss_d: 0.3116, real_score: 0.8900, fake_score: 0.37
19

Current LRs – Generator: 0.000200, Discriminator: 0.000200

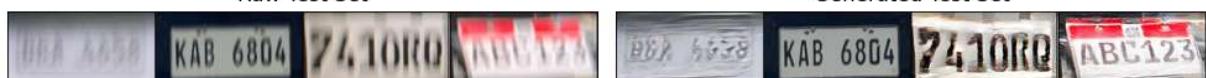
Validation Data



Epoch [19/200] - Mean PSNR: 17.0421, Mean SSIM: 0.7020, FID: 58.55329513549805

Raw Test Set

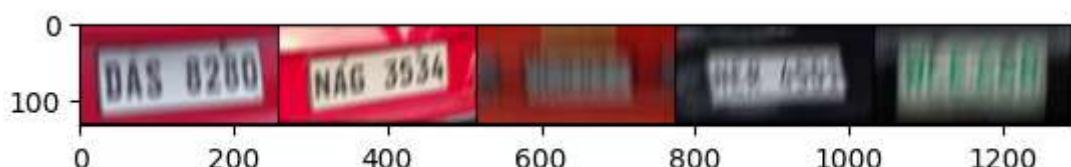
Generated Test Set



100% |

1250/1250 [18:55<00:00, 1.10it/s]

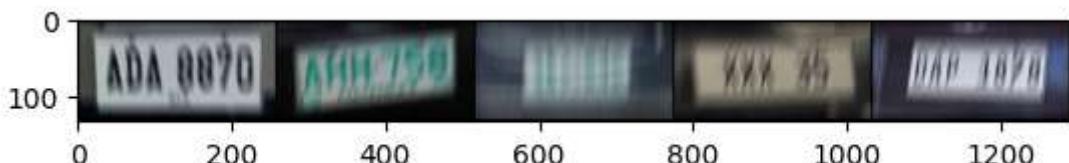
Training Data



Epoch [20/200], loss_g: 8.6885, loss_d: 0.6190, real_score: 0.3975, fake_score: 0.20
71

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [20/200] - Mean PSNR: 17.3240, Mean SSIM: 0.7178, FID: 56.538368225097656

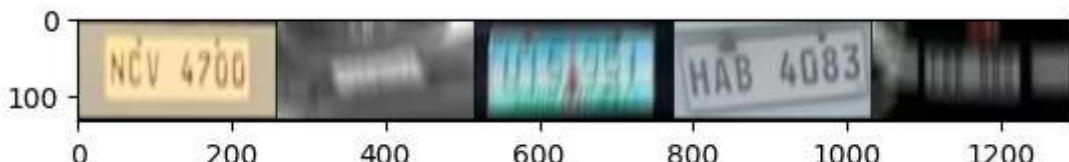
Raw Test Set

Generated Test Set



100% |
1250/1250 [18:55<00:00, 1.10it/s]

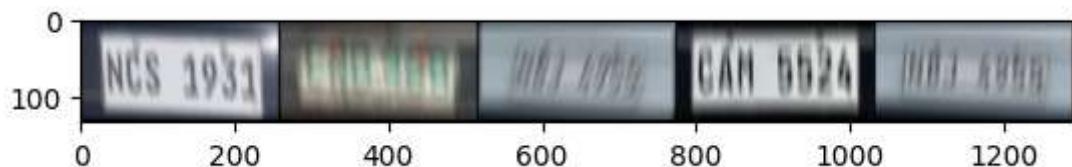
Training Data



Epoch [21/200], loss_g: 9.2200, loss_d: 0.1602, real_score: 0.9003, fake_score: 0.18
46

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [21/200] - Mean PSNR: 17.5276, Mean SSIM: 0.7239, FID: 55.275970458984375

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:54<00:00, 1.10it/s]

Training Data



Epoch [22/200], loss_g: 8.4237, loss_d: 0.3641, real_score: 0.9452, fake_score: 0.46

55

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [22/200] - Mean PSNR: 17.5002, Mean SSIM: 0.7235, FID: 55.5517463684082

Raw Test Set

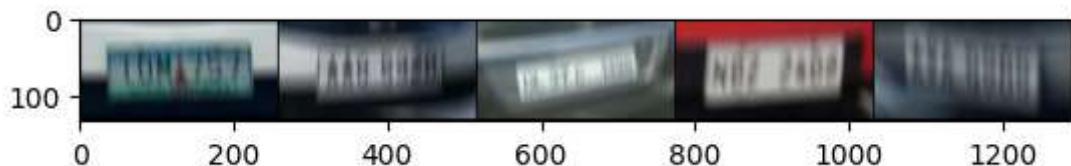
Generated Test Set



100%

1250/1250 [18:56<00:00, 1.10it/s]

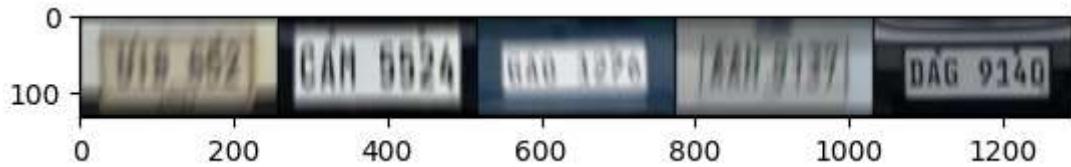
Training Data



Epoch [23/200], loss_g: 10.4220, loss_d: 0.2291, real_score: 0.8484, fake_score: 0.2371

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [23/200] - Mean PSNR: 17.5685, Mean SSIM: 0.7260, FID: 54.107852935791016

Raw Test Set

Generated Test Set



100%

1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [24/200], loss_g: 7.9200, loss_d: 0.4020, real_score: 0.5735, fake_score: 0.1784

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [24/200] - Mean PSNR: 17.4107, Mean SSIM: 0.7240, FID: 54.32588577270508



100% | 1250/1250 [18:54<00:00, 1.10it/s]
Training Data



Epoch [25/200], loss_g: 8.1682, loss_d: 0.6896, real_score: 0.3011, fake_score: 0.0653

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [25/200] - Mean PSNR: 17.4726, Mean SSIM: 0.7229, FID: 55.02034378051758



100% | 1250/1250 [18:55<00:00, 1.10it/s]
Training Data



Epoch [26/200], loss_g: 78.0004, loss_d: 0.4434, real_score: 0.5355, fake_score: 0.1602

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [26/200] - Mean PSNR: 17.1482, Mean SSIM: 0.7182, FID: 54.91360092163086

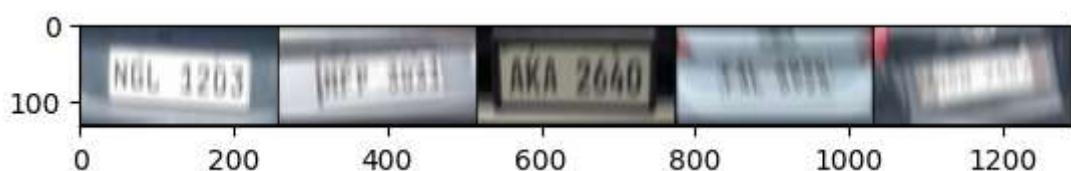
Raw Test Set

Generated Test Set



100% |
1250/1250 [18:55<00:00, 1.10it/s]

Training Data





Epoch [27/200], loss_g: 7.8816, loss_d: 0.3581, real_score: 0.7307, fake_score: 0.3005

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [27/200] - Mean PSNR: 17.5700, Mean SSIM: 0.7278, FID: 53.796043395996094

Raw Test Set

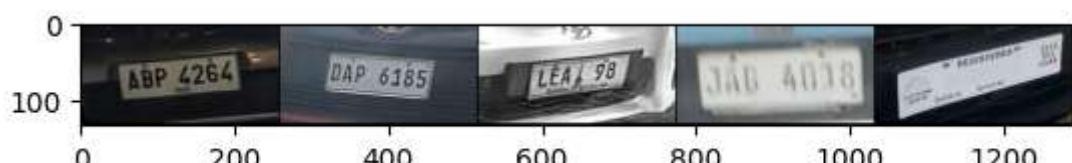
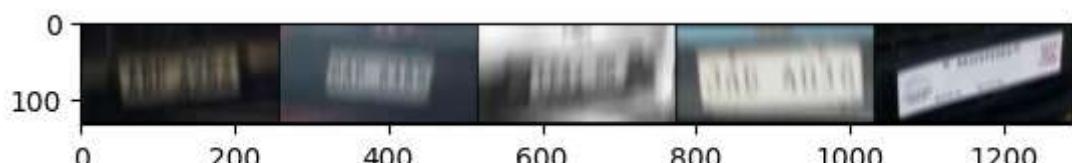
Generated Test Set

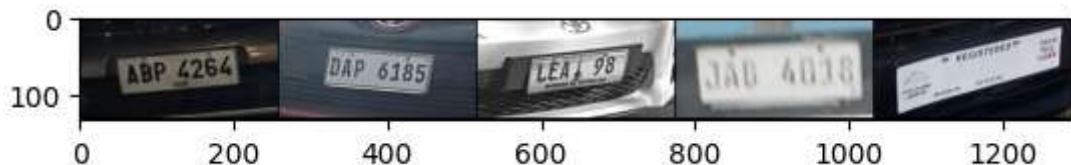


100%

1250/1250 [18:55<00:00, 1.10it/s]

Training Data





Epoch [28/200], loss_g: 7.8280, loss_d: 0.3021, real_score: 0.9086, fake_score: 0.37
81

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [28/200] - Mean PSNR: 17.5111, Mean SSIM: 0.7213, FID: 52.690799713134766

Raw Test Set

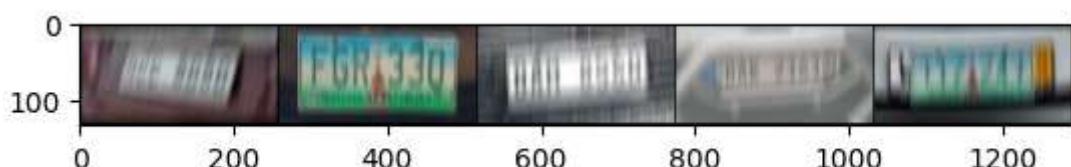
Generated Test Set



100% |

1250/1250 [18:55<00:00, 1.10it/s]

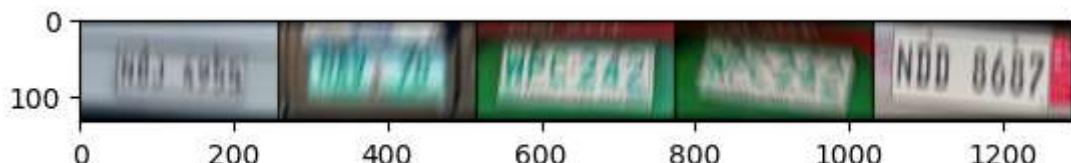
Training Data



Epoch [29/200], loss_g: 7.7533, loss_d: 0.2176, real_score: 0.9296, fake_score: 0.28
45

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [29/200] - Mean PSNR: 17.7545, Mean SSIM: 0.7321, FID: 52.654701232910156

Raw Test Set

Generated Test Set



100% |
1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [30/200], loss_g: 7.0580, loss_d: 0.5414, real_score: 0.8560, fake_score: 0.56
78

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [30/200] - Mean PSNR: 17.6164, Mean SSIM: 0.7283, FID: 52.9904899597168

Raw Test Set

Generated Test Set



100% |

1250/1250 [18:55<00:00, 1.10it/s]

Training Data



Epoch [31/200], loss_g: 8.1779, loss_d: 0.3211, real_score: 0.8771, fake_score: 0.3775

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [31/200] - Mean PSNR: 17.6180, Mean SSIM: 0.7337, FID: 52.106353759765625

Raw Test Set

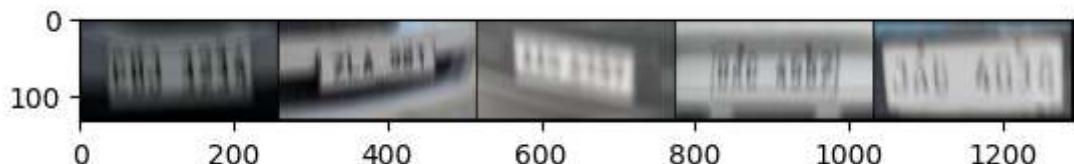
Generated Test Set



100%

1250/1250 [19:26<00:00, 1.07it/s]

Training Data



Epoch [32/200], loss_g: 7.5448, loss_d: 0.4539, real_score: 0.7369, fake_score: 0.4199

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data





Epoch [32/200] - Mean PSNR: 17.5873, Mean SSIM: 0.7309, FID: 51.80231475830078

Raw Test Set

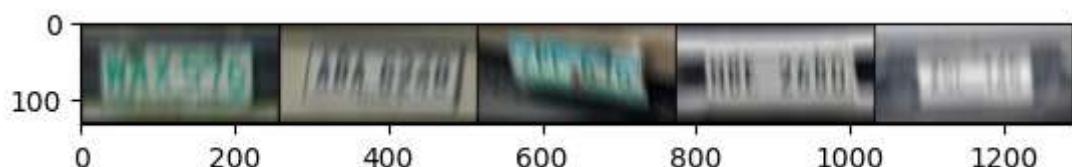
Generated Test Set



100% |

1250/1250 [19:47<00:00, 1.05it/s]

Training Data

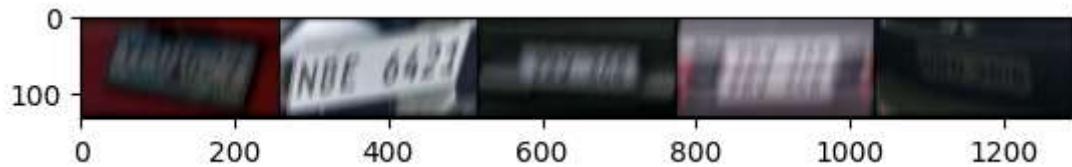


Epoch [33/200], loss_g: 8.2800, loss_d: 0.3370, real_score: 0.6413, fake_score: 0.17

17

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [33/200] - Mean PSNR: 17.7610, Mean SSIM: 0.7363, FID: 50.74321746826172



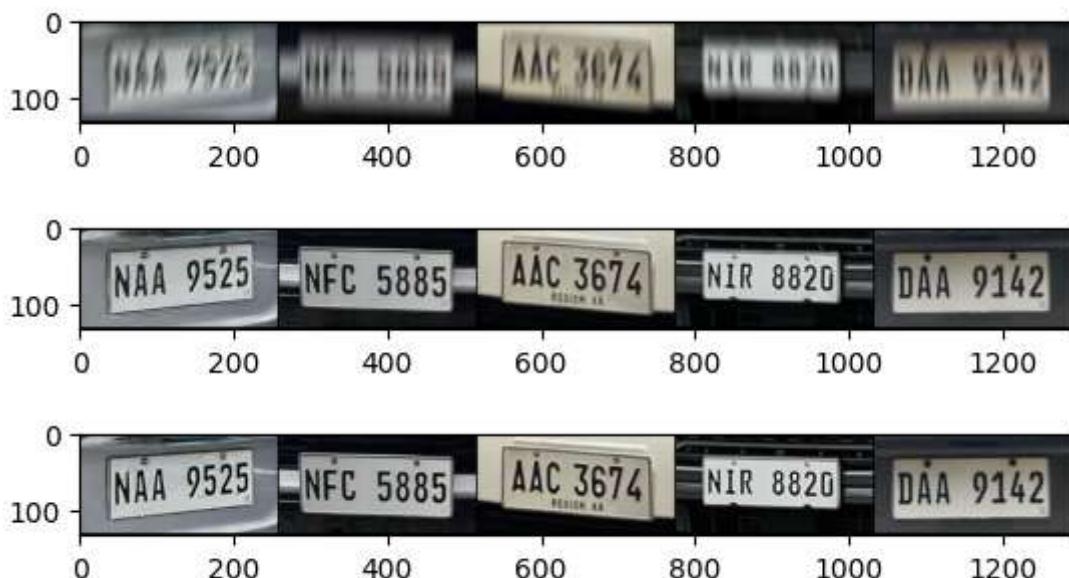
100% |
1250/1250 [19:42<00:00, 1.06it/s]
Training Data



Epoch [34/200], loss_g: 7.7352, loss_d: 0.8898, real_score: 0.2138, fake_score: 0.05
24

Current LRs – Generator: 0.000200, Discriminator: 0.000200

Validation Data



Epoch [34/200] - Mean PSNR: 17.8010, Mean SSIM: 0.7341, FID: 50.371219635009766



Epoch 00034: reducing learning rate of group 0 to 1.0000e-04.

100% |
1250/1250 [19:34<00:00, 1.06it/s]

Training Data



Epoch [35/200], loss_g: 7.7428, loss_d: 0.3664, real_score: 0.6519, fake_score: 0.23
62

Current LRs – Generator: 0.000200, Discriminator: 0.000100

Validation Data



Epoch [35/200] - Mean PSNR: 17.7283, Mean SSIM: 0.7344, FID: 50.69340515136719

Raw Test Set

Generated Test Set



5% |
| 66/1250 [01:09<18:57, 1.04it/s]

Post Training

In [34]: `torch.cuda.empty_cache()`

Save Model

```
In [34]: torch.save(generator.state_dict(), "[RAU-NET]generator-Datasetv2-5epoch.pt")
torch.save(discriminator.state_dict(), "[RAU-NET]discriminator-Datasetv2-5epoch.pt"

In [35]: torch.save({
    'generator_state_dict': generator.state_dict(),
    'discriminator_state_dict': discriminator.state_dict(),
    'opt_g_state_dict': g_opt.state_dict(),
    'opt_d_state_dict': d_opt.state_dict(),
    'history': history
}, f'[RAU-NET]Checkpoint-Datasetv2-30epoch.pt')
```

Model Evaluation

```
In [56]: psnr_vals = []
ssim_vals = []

fid_metric = FrechetInceptionDistance(feature=2048).to(device)
fid_metric.reset()

generator.eval()
for val_input, val_target in valid_dl:
    val_input = val_input.to(device)
    val_target = val_target.to(device)

    with torch.no_grad():
        gen_output = generator(val_input)

    # Convert to uint8 for FID
    gen_uint8 = to_uint8(gen_output)
    target_uint8 = to_uint8(val_target)

    fid_metric.update(gen_uint8, real=False)
    fid_metric.update(target_uint8, real=True)

    #show_generated_images(gen_output)

    # Compute PSNR/SSIM
    psnr = psnr_metric(gen_output, val_target).item()
    ssim = ssim_metric(gen_output, val_target).item()
    psnr_vals.append(psnr)
    ssim_vals.append(ssim)

print("Mean PSNR:", sum(psnr_vals) / len(psnr_vals))
print("Mean SSIM:", sum(ssim_vals) / len(ssim_vals))
print("FID:", fid_metric.compute().item())
```

Mean PSNR: 15.653532288291238
Mean SSIM: 0.6187068169767206
FID: 67.66695404052734

```
In [33]: def deblur_image(image_path, generator, device):
    image = Image.open(image_path).convert('RGB')
    input_tensor = valid_transform(image).unsqueeze(0).to(device)

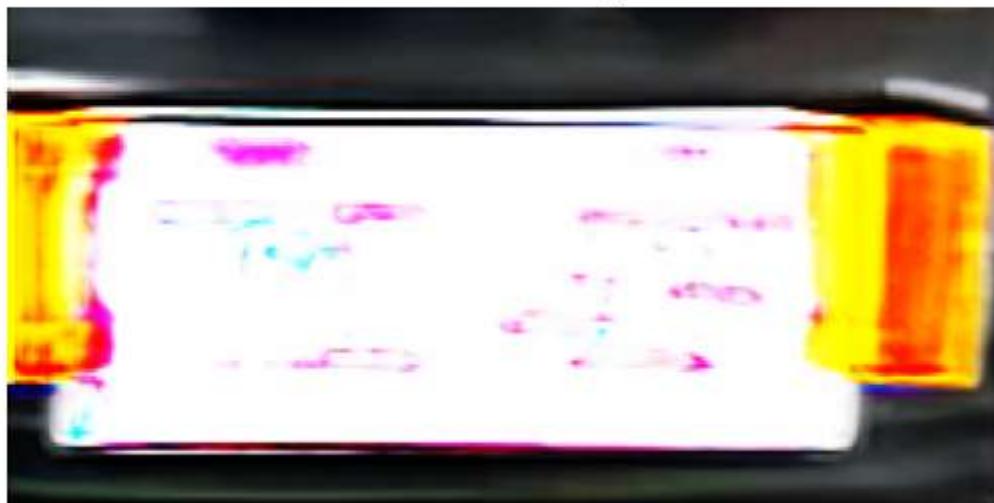
    generator.eval()
    with torch.no_grad():
        gen_image = generator(input_tensor)

    output = denorm(gen_image.squeeze(0).cpu())

    plt.imshow(output.permute(1, 2, 0)) # CHW → HWC
    plt.axis('off')
    plt.title("Generated Image")
    plt.show()
```

```
In [31]: deblur_image(os.path.join(DATA_DIR, 'h.blur', 'v2', 'test', 'h.blur', '129.jpg'), g
```

Generated Image



```
In [71]: raw_grid_img, gen_grid_img = create_image_grid_from_loader(generator, test_dl, device)

fig, axes = plt.subplots(1, 2, figsize=(15, 15)) # 1 row, 2 columns

# Raw Test Set
axes[0].imshow(raw_grid_img)
axes[0].axis("off")
axes[0].set_title("Raw Test Set")

# Generated Test Set
axes[1].imshow(gen_grid_img)
axes[1].axis("off")
axes[1].set_title("Generated Test Set")

plt.tight_layout()
plt.show()
```



```
In [33]: torch.cuda.empty_cache()
```

Load Model

```
In [50]: MODELS_DIR = os.path.join('..', '..', '..', 'models')
```

```
In [55]: generator = AttentionResUNetGenerator()
loaded_g_wts = torch.load("C:\\Thesis\\LiPAD with Paddle\\models\\gan_v.blur.pt")
generator.load_state_dict(loaded_g_wts)
generator = to_device(generator, device)

#discriminator = PatchDiscriminator()
#Loaded_d_wts = torch.load('[RAU-NET OCRDET LR]discriminator-Datasetv2-80 epoch.pt')
#discriminator.load_state_dict(loaded_d_wts)
#discriminator = to_device(discriminator, device)
```

```
In [34]: checkpoint = torch.load('[RAU-NET OCRDET LR]Checkpoint-Logs-Datasetv2-80 epoch.pt')
checkpoint2 = torch.load('[RAU-NET OCRDET LR 2]Checkpoint-Logs-Datasetv2-145 epoch.
#g_opt.load_state_dict(checkpoint['opt_g_state_dict'])
#d_opt.load_state_dict(checkpoint['opt_d_state_dict'])
#history = checkpoint['losses_g']
```

```
In [39]: len(checkpoint)
```

```
Out[39]: 14
```

```
In [39]: def are_models_identical(model1, model2):
    # Check if state dictionaries have the same keys
    if model1.state_dict().keys() != model2.state_dict().keys():
        return False

    # Compare each parameter's values
    for key in model1.state_dict():
        if not torch.equal(model1.state_dict()[key], model2.state_dict()[key]):
            return False
    return True
```

```
In [45]: if are_models_identical(generator1, generator2):
    print("Models have identical weights and biases")
else:
    print("Models have different weights or biases")
```

```
Models have different weights or biases
```

```
In [9]: len(history)
```

```
Out[9]: 79
```

Save Reconstructed Test Set

```
In [39]: def save_generated_tests(generator, test_dl, sample_dir='./results'):
    generator.eval()
    os.makedirs(sample_dir, exist_ok=True)

    with torch.no_grad():
        global_idx = 1 # running index for image IDs

        for inputs, _ in test_dl:
            inputs = inputs.to(device)

            fake_images = generator(inputs)
            fake_images = fake_images.clamp(-1, 1)

            batch_size = fake_images.size(0)

            for b in range(batch_size):
                img = fake_images[b]

                file_name = f"{global_idx}.jpg"
                save_path = os.path.join(sample_dir, file_name)

                save_image(denorm(img).unsqueeze(0), save_path, nrow=1)
                global_idx += 1
```

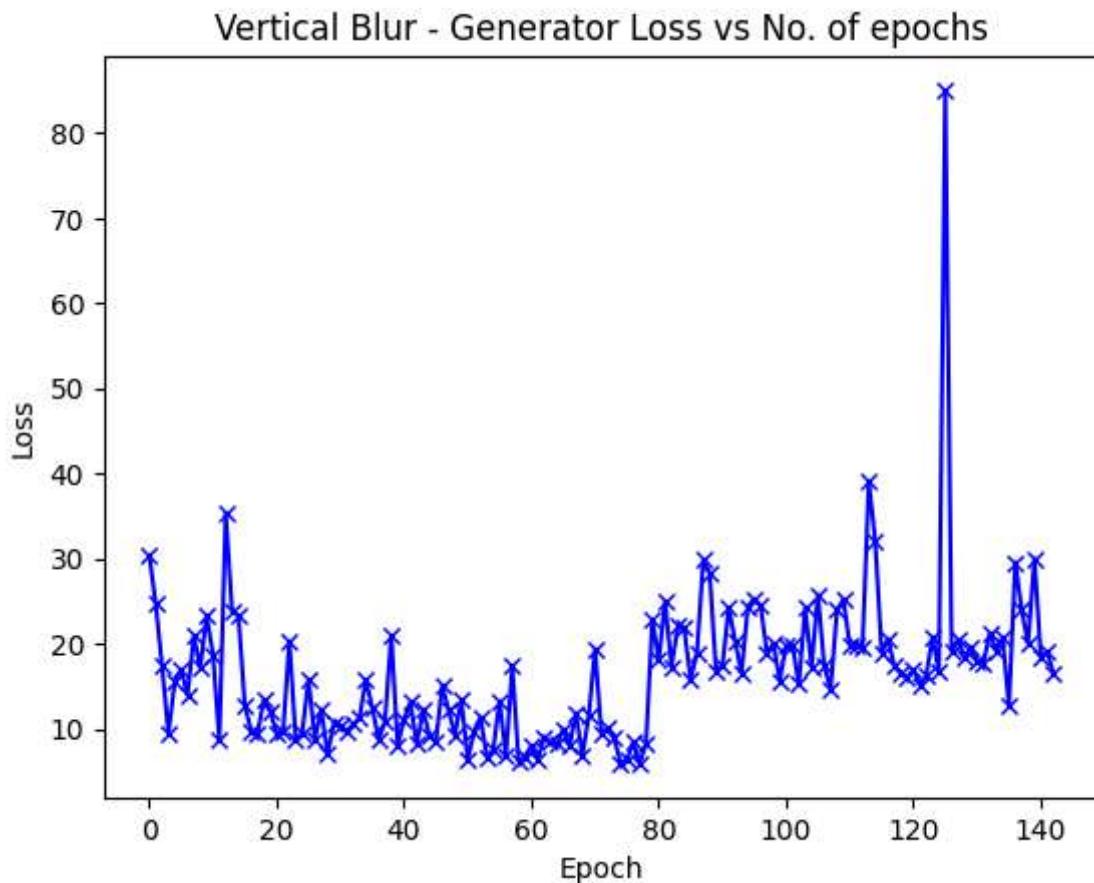
```
In [40]: save_generated_tests(generator, test_dl)
```

```
In [39]: from collections import OrderedDict

for key, value in checkpoint2.items():
    if key not in checkpoint:
        checkpoint[key] = value
    else:
        if isinstance(value, list):
            checkpoint[key].extend(value) # concatenate lists
        elif isinstance(value, OrderedDict):
            checkpoint[key].update(value) # merge OrderedDicts
        else:
            # fallback, keep checkpoint2 value
            checkpoint[key] = value
```

```
In [40]: def plot_g_losses(history):
    losses_g = [x for x in history['losses_g']]
    plt.plot(losses_g, '-bx')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Vertical Blur - Generator Loss vs No. of epochs')

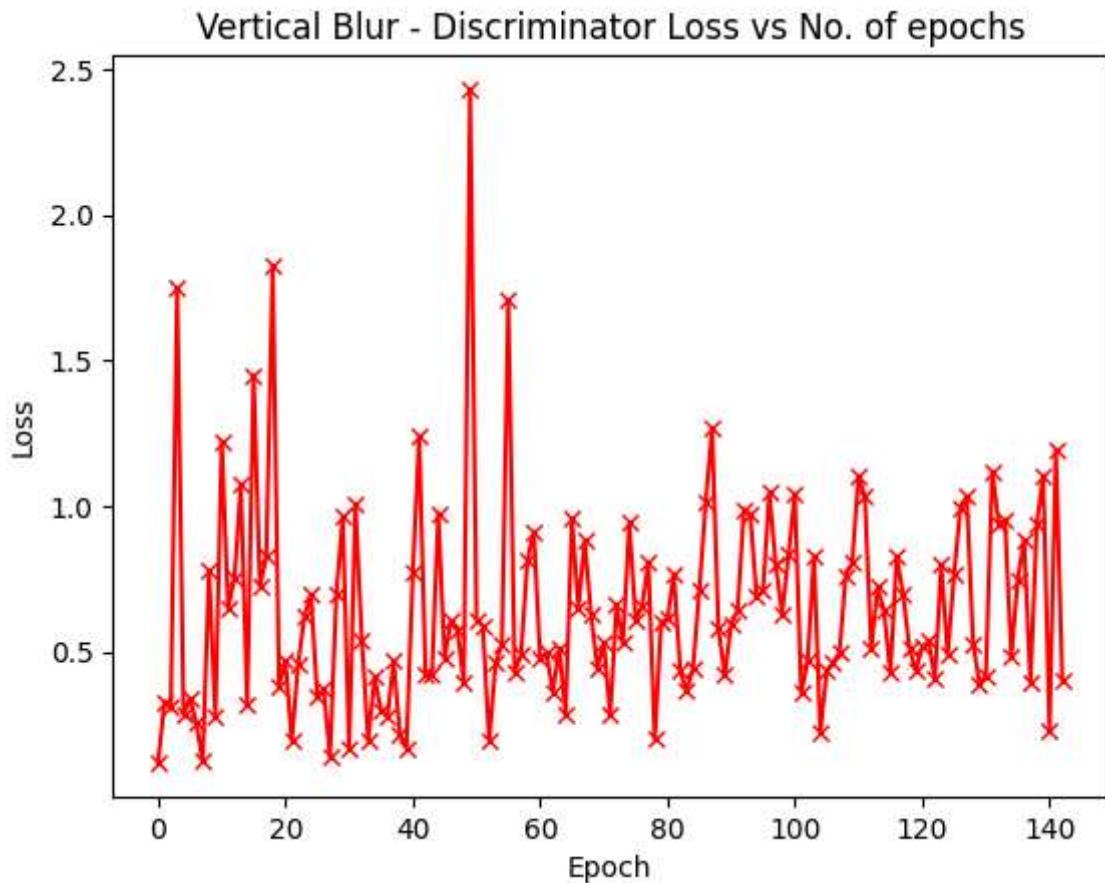
plot_g_losses(checkpoint)
```



```
In [45]: def plot_d_losses(history):
    losses_d = [x for x in history['losses_d']]
```

```
plt.plot(losses_d, '-rx')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Vertical Blur - Discriminator Loss vs No. of epochs')

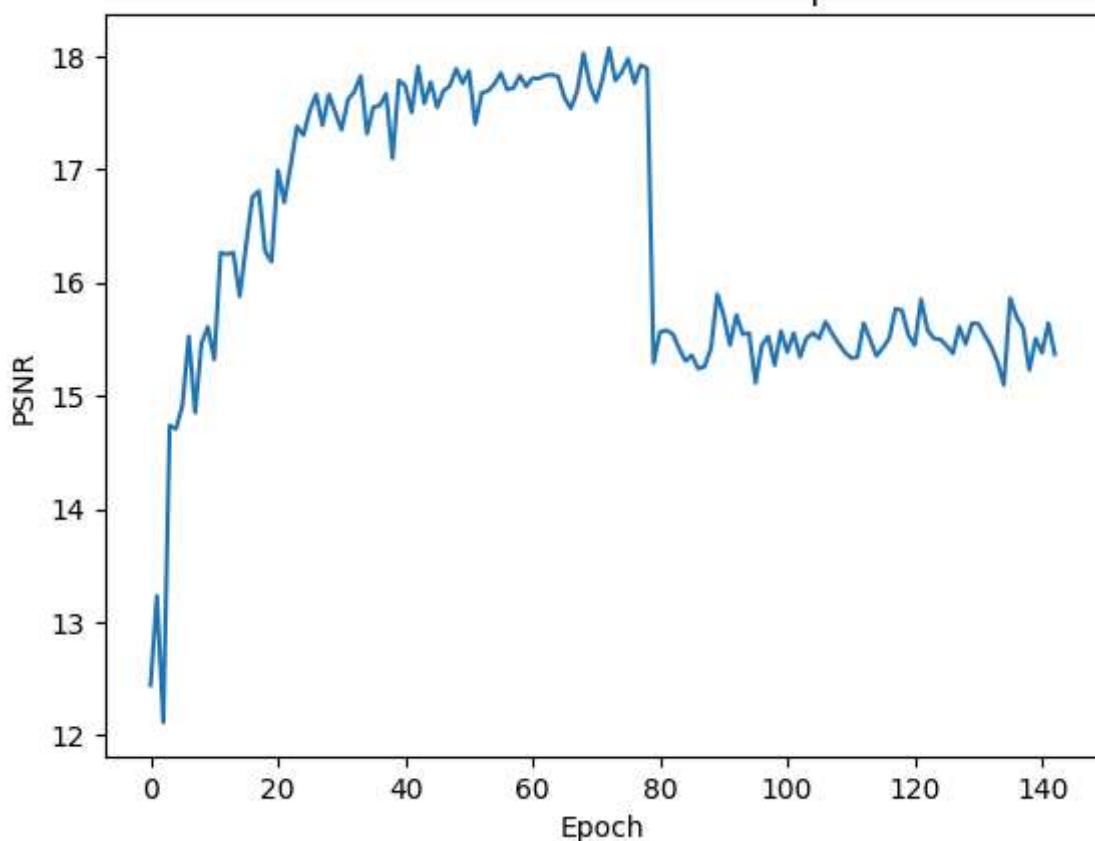
plot_d_losses(checkpoint)
```



```
In [46]: def plot_psnrs(history):
    psnrs = [x for x in history['psnrs']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('PSNR')
    plt.title('Vertical Blur - PSNR vs No. of epochs')

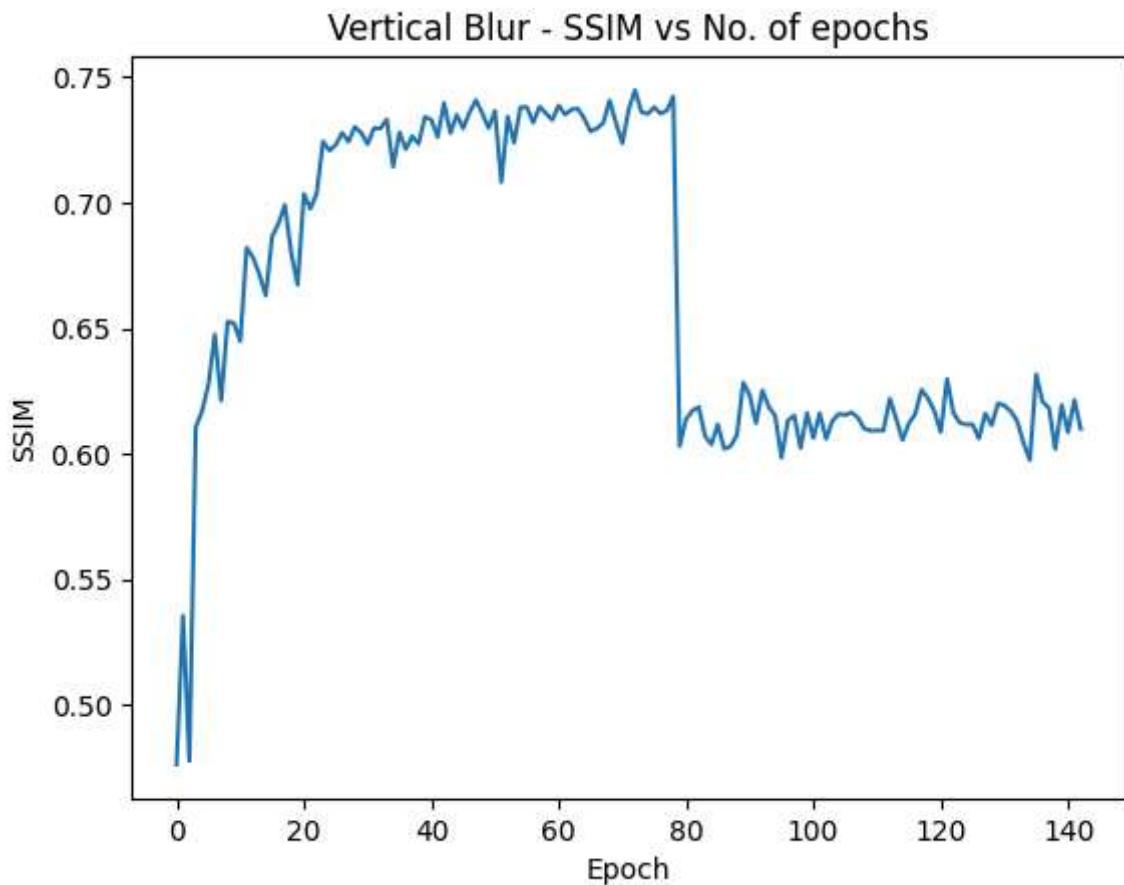
plot_psnrs(checkpoint)
```

Vertical Blur - PSNR vs No. of epochs



```
In [47]: def plot_ssims(history):
    psnrs = [x for x in history['ssims']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('SSIM')
    plt.title('Vertical Blur - SSIM vs No. of epochs')

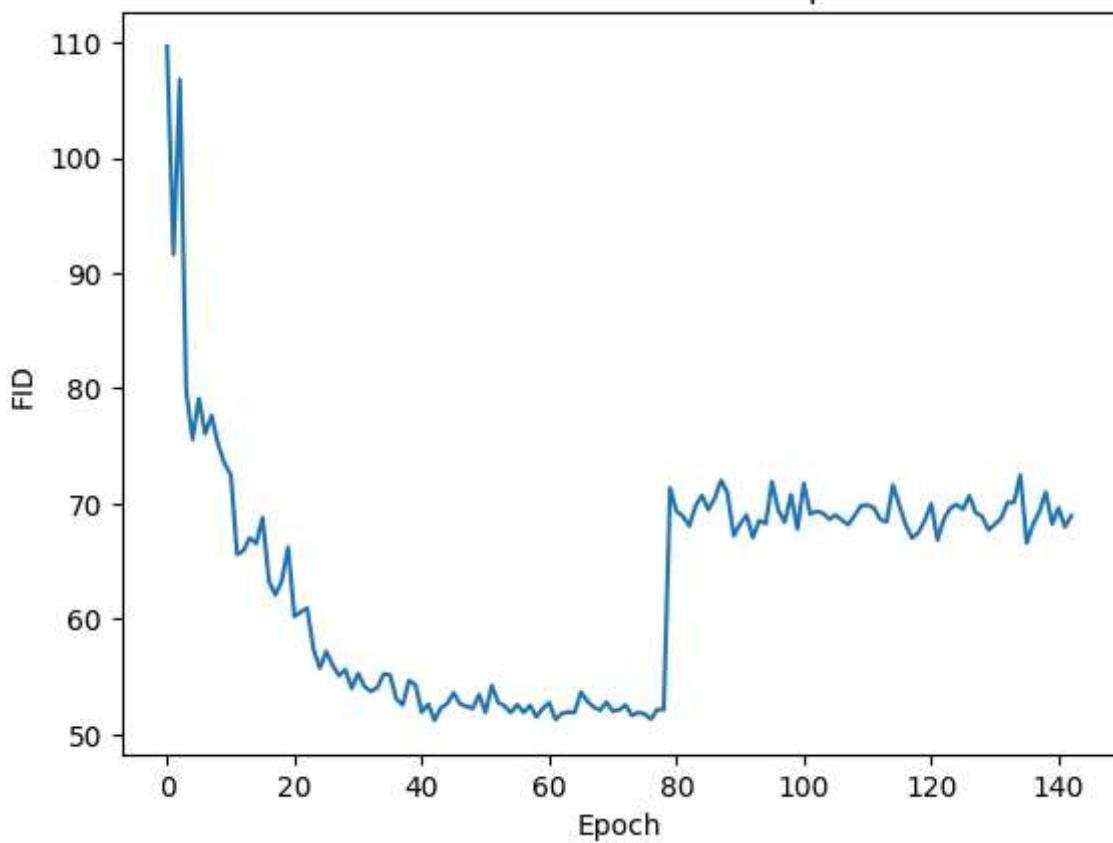
plot_ssims(checkpoint)
```



```
In [48]: def plot_fids(history):
    psnrs = [x for x in history['fids']]
    plt.plot(psnrs)
    plt.xlabel('Epoch')
    plt.ylabel('FID')
    plt.title('Vertical Blur - FID vs No. of epochs')

plot_fids(checkpoint)
```

Vertical Blur - FID vs No. of epochs



In []: